

Ford GoBike system

Regression project

YEHUDA FINKELSHTIEN

DANIEL BITON

Introduction

This data set includes information about individual rides made in a bike-sharing system covering the greater San Francisco Bay area.

The dataset will require some data wrangling to make it tidy for analysis.

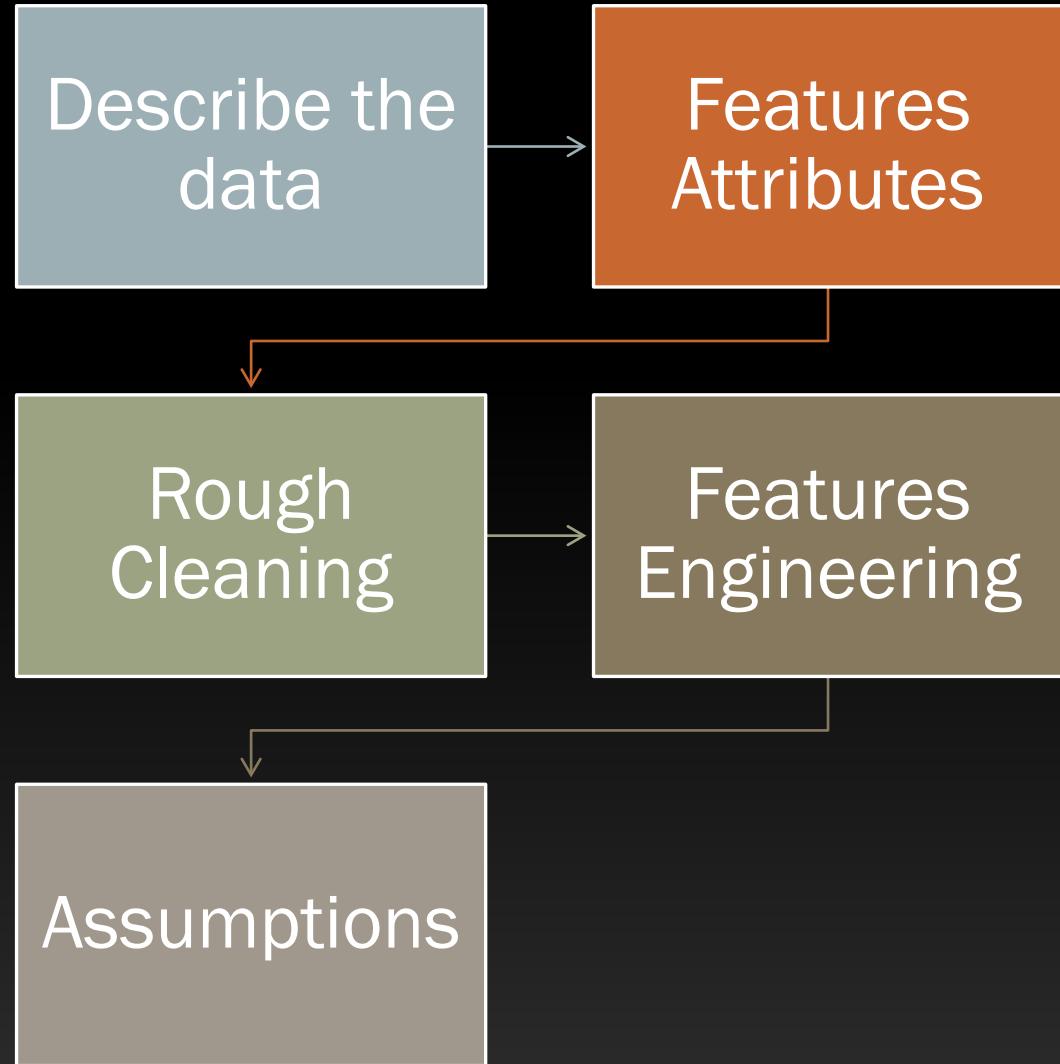
Data source: [Kaggle](#)

Prediction target definition

We would like to predict a **duration** of bike-sharing system usage based on the data which collect in the moment that ride is starting according to the location, gender, age and so forth.

This data should contribute the guidelines for the marketing division where to focus on the next campaign based on the attributes were examine.

Exploratory Data Analysis



Data .Info()

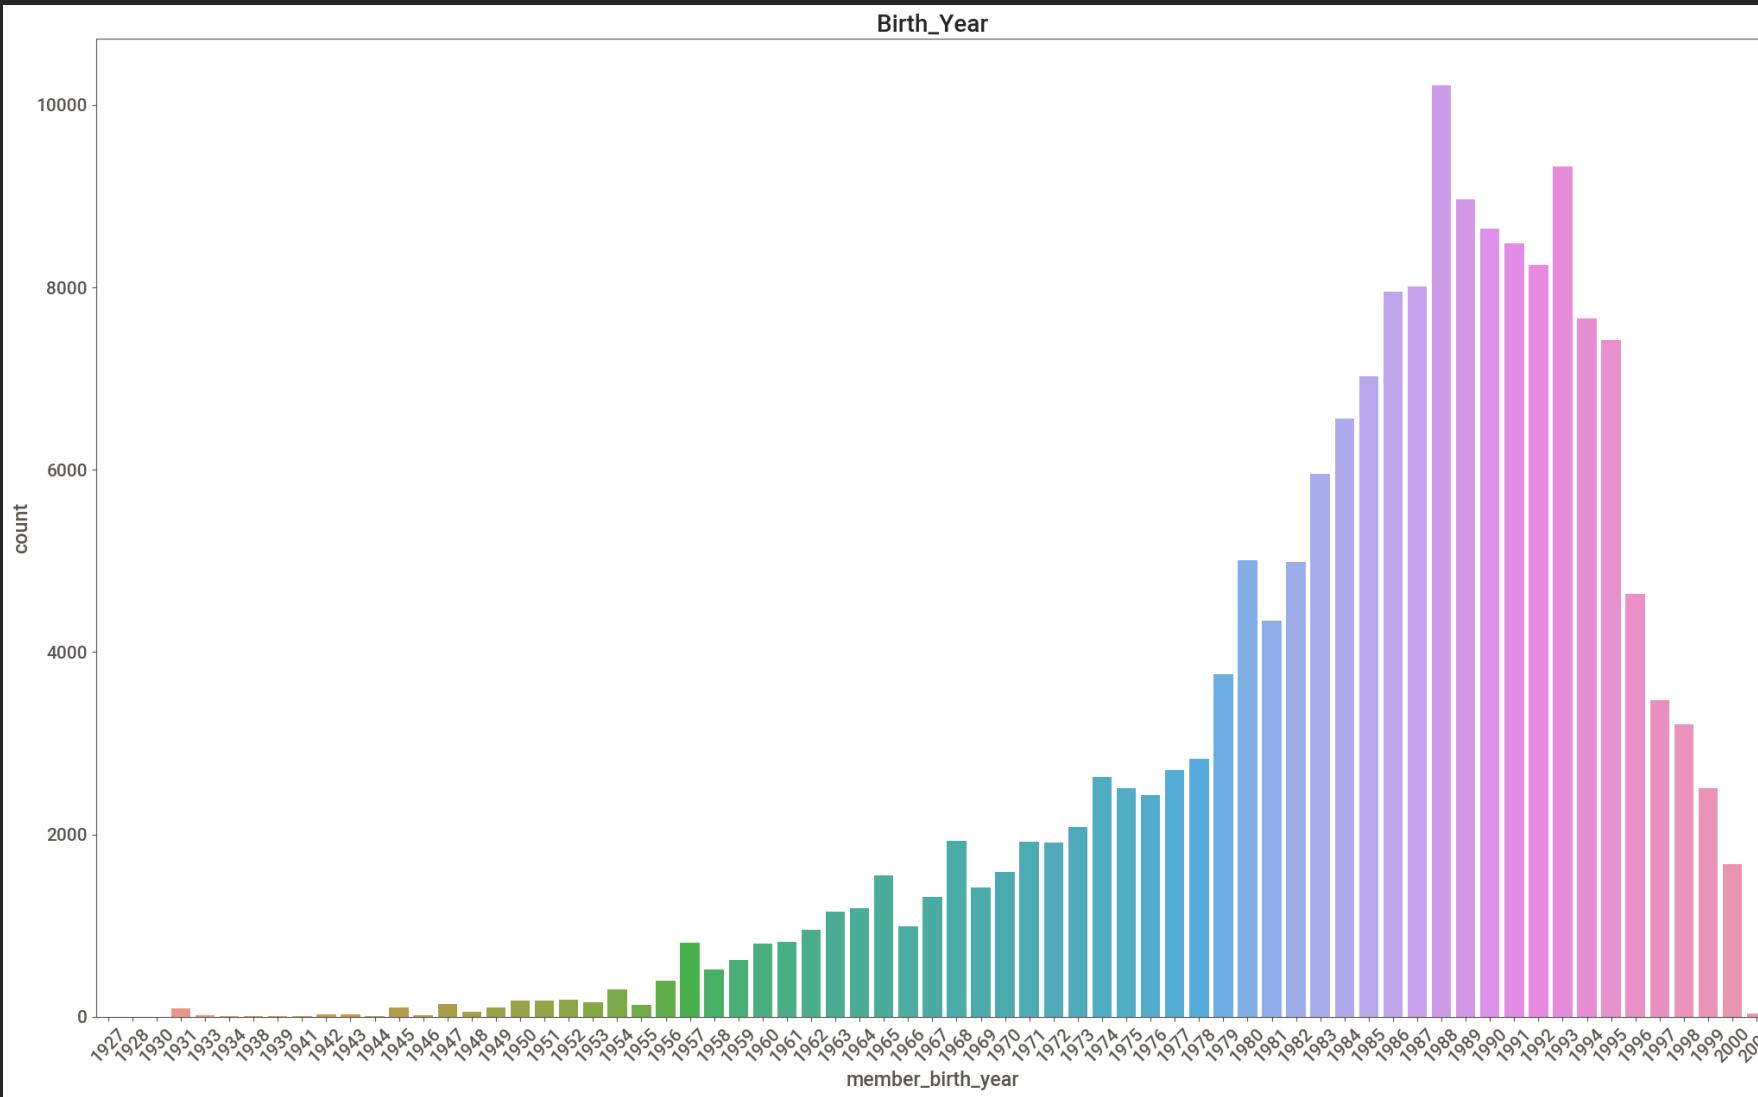
NULLS

```
duration_sec          0
start_time            0
end_time              0
start_station_id      197
start_station_name    197
start_station_latitude 0
start_station_longitude 0
end_station_id        197
end_station_name      197
end_station_latitude  0
end_station_longitude 0
bike_id               0
user_type              0
member_birth_year     8265
member_gender          8265
bike_share_for_all_trip 0
dtype: int64
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 183416 entries, 0 to 183415
Data columns (total 16 columns):
 #   Column           Non-Null Count   Dtype  
--- 
 0   duration_sec    183416 non-null    int64  
 1   start_time       183416 non-null    object  
 2   end_time         183416 non-null    object  
 3   start_station_id 183219 non-null    float64 
 4   start_station_name 183219 non-null    object  
 5   start_station_latitude 183416 non-null    float64 
 6   start_station_longitude 183416 non-null    float64 
 7   end_station_id   183219 non-null    float64 
 8   end_station_name 183219 non-null    object  
 9   end_station_latitude 183416 non-null    float64 
 10  end_station_longitude 183416 non-null    float64 
 11  bike_id          183416 non-null    int64  
 12  user_type        183416 non-null    object  
 13  member_birth_year 175151 non-null    float64 
 14  member_gender    175151 non-null    object  
 15  bike_share_for_all_trip 183416 non-null    object  
dtypes: float64(7), int64(2), object(7)
memory usage: 22.4+ MB
```

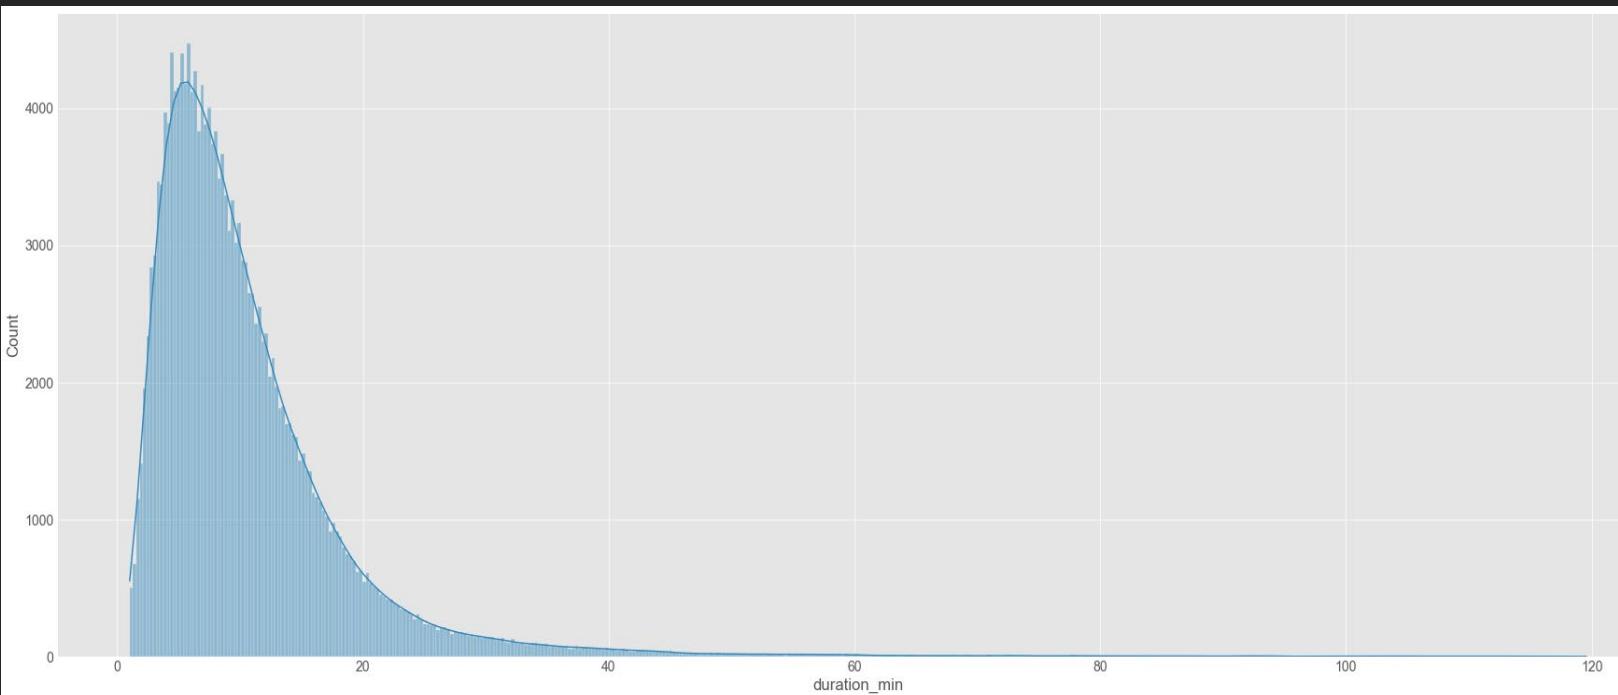
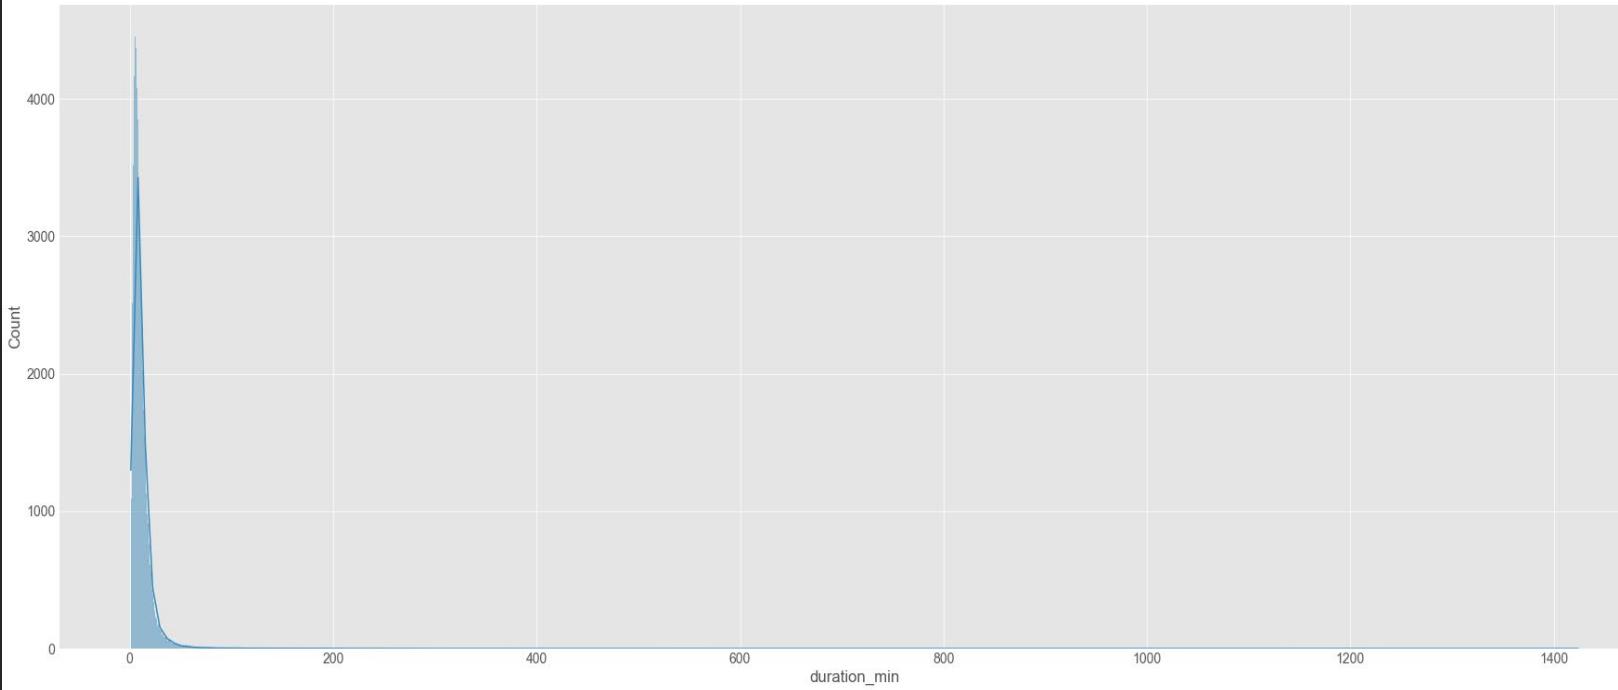

Age Issues

- Cleaning to range of 20-100 years old
- Imputing missing values by average



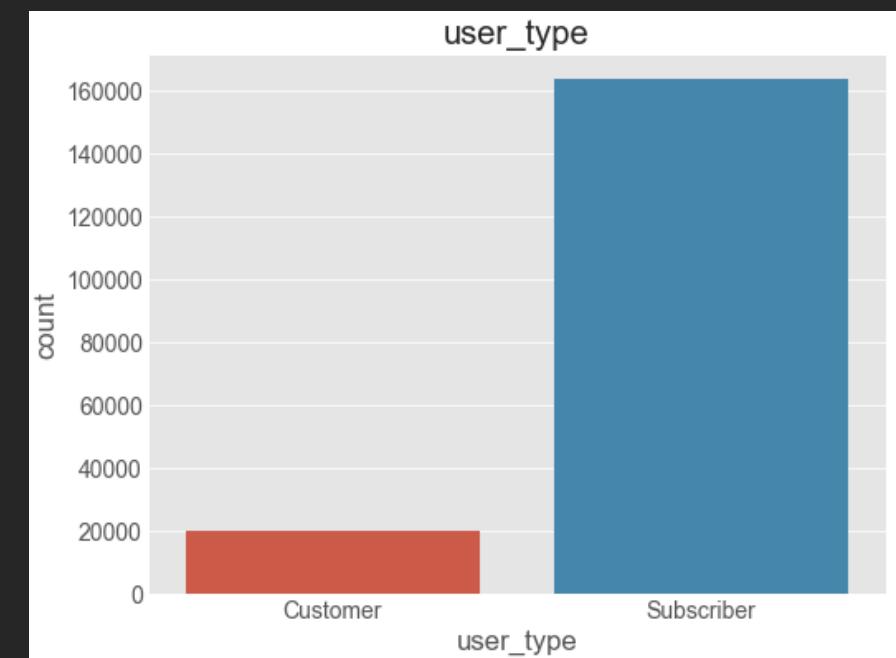
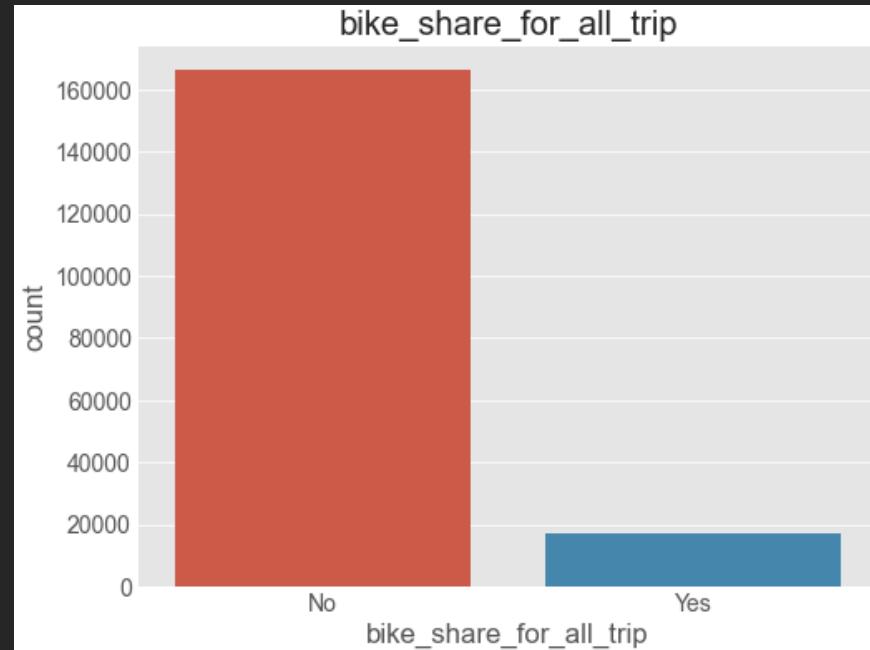
Duration distribution

- Cleaning out of range of 120 min



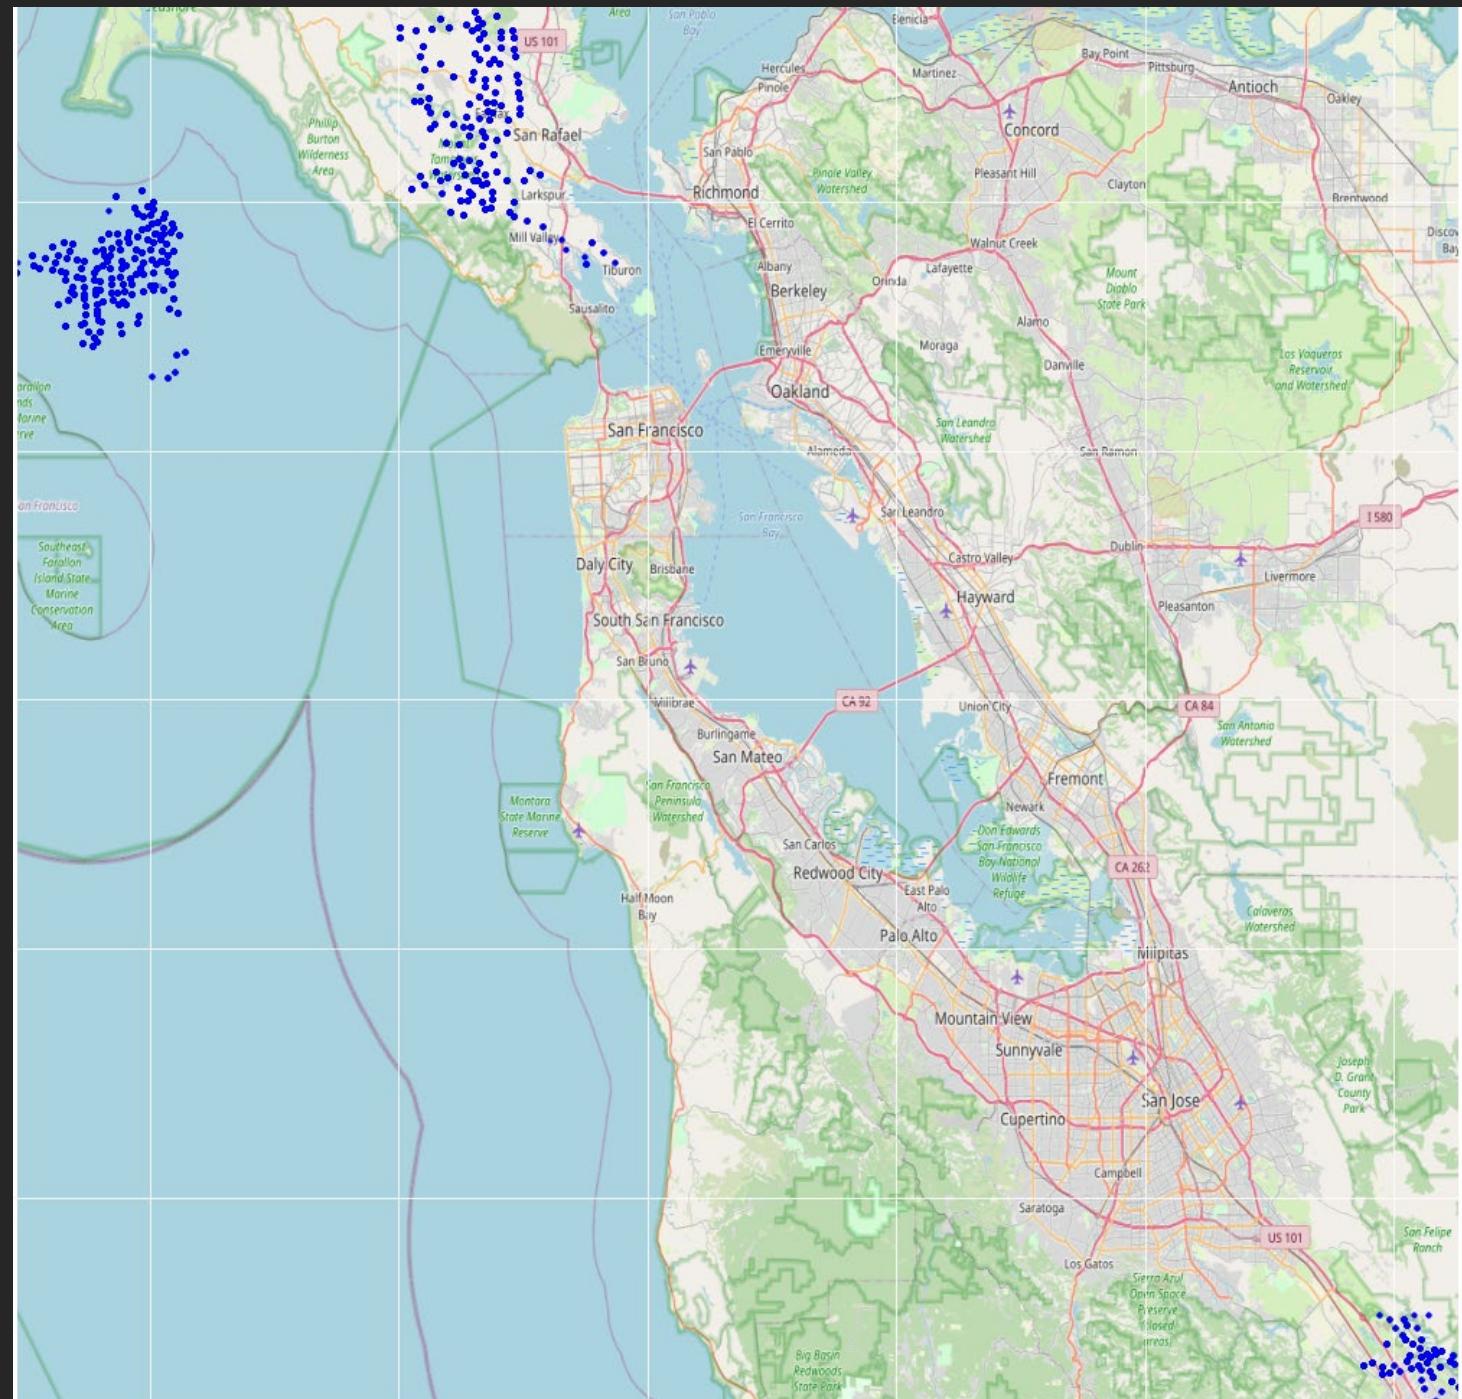
Categorical Features Distribution

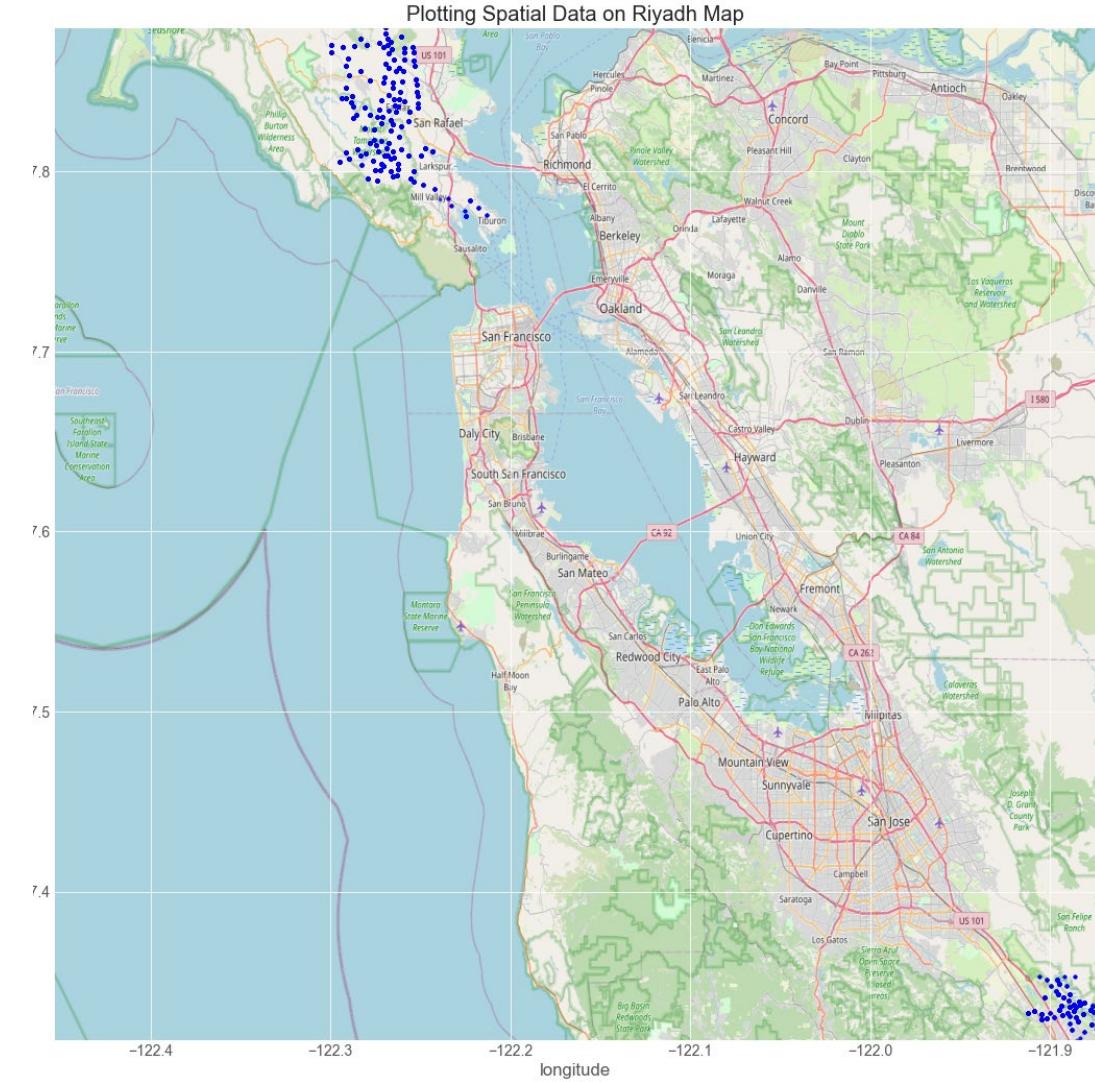
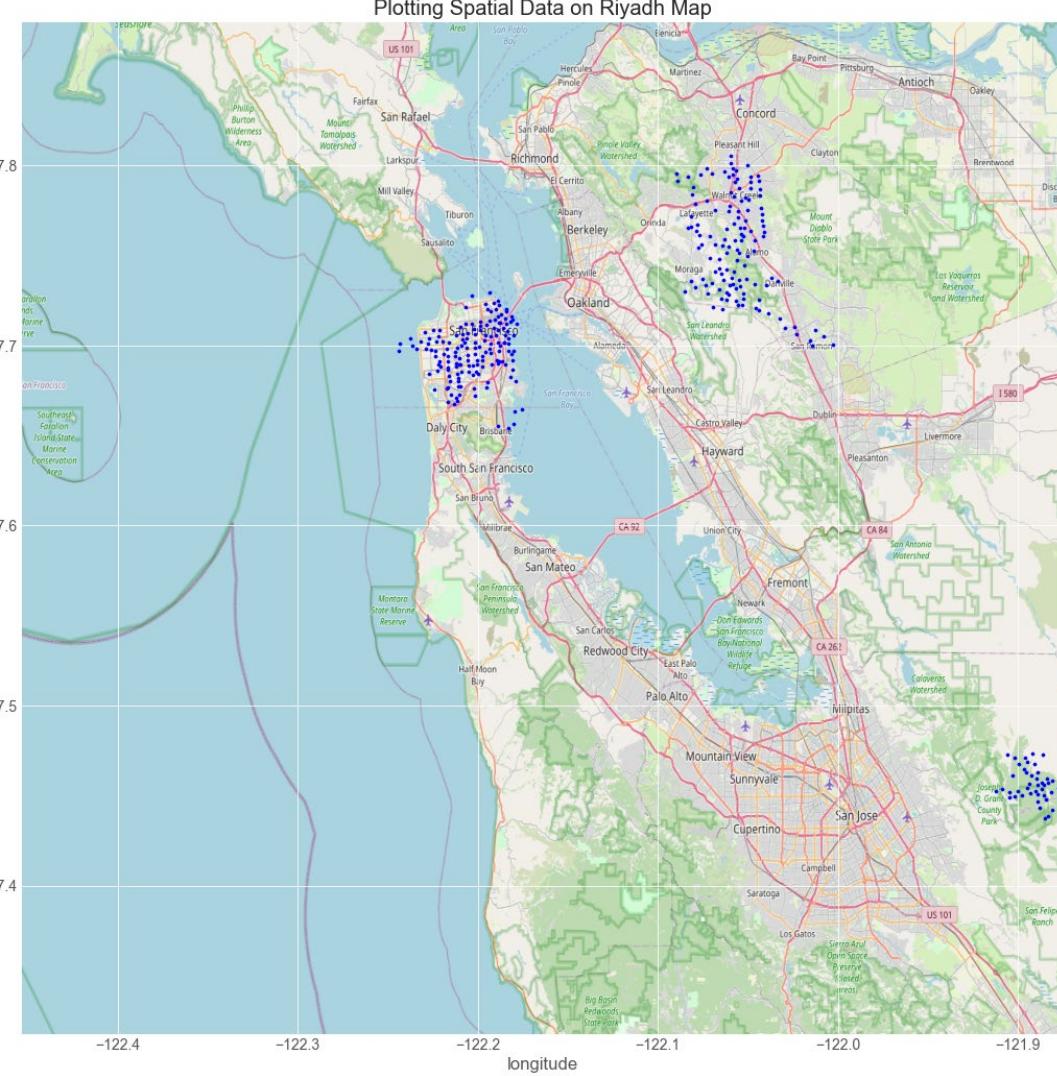
“All_Trip” organizing tours in San Francisco Bay area.



Mapping stations

THERE ARE SOME STATIONS WHERE ARE LOCATED ON THE SEA



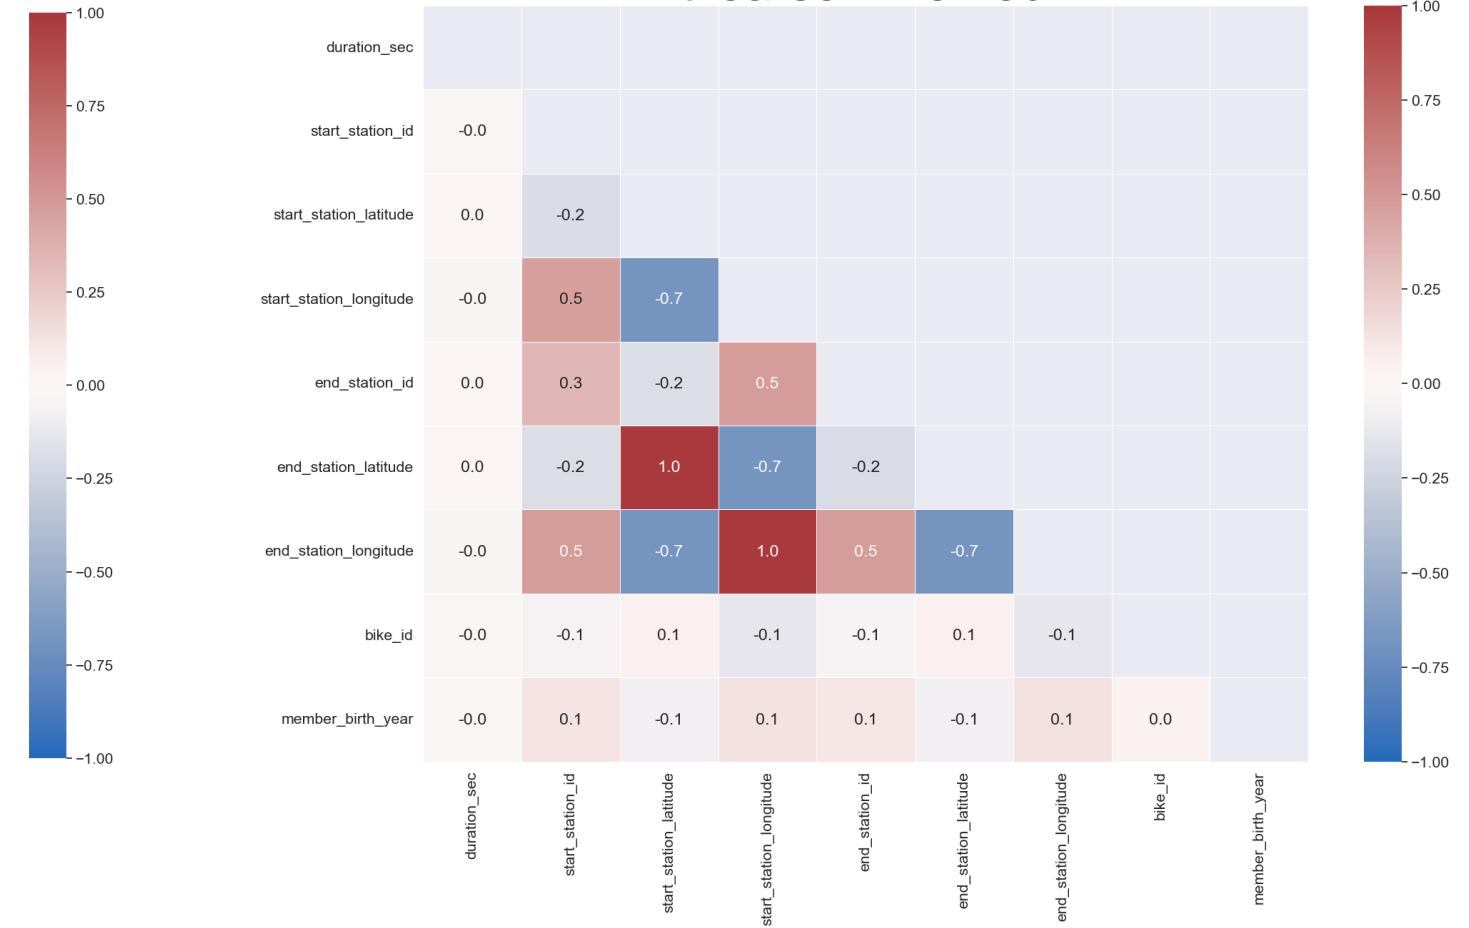




Spearman method

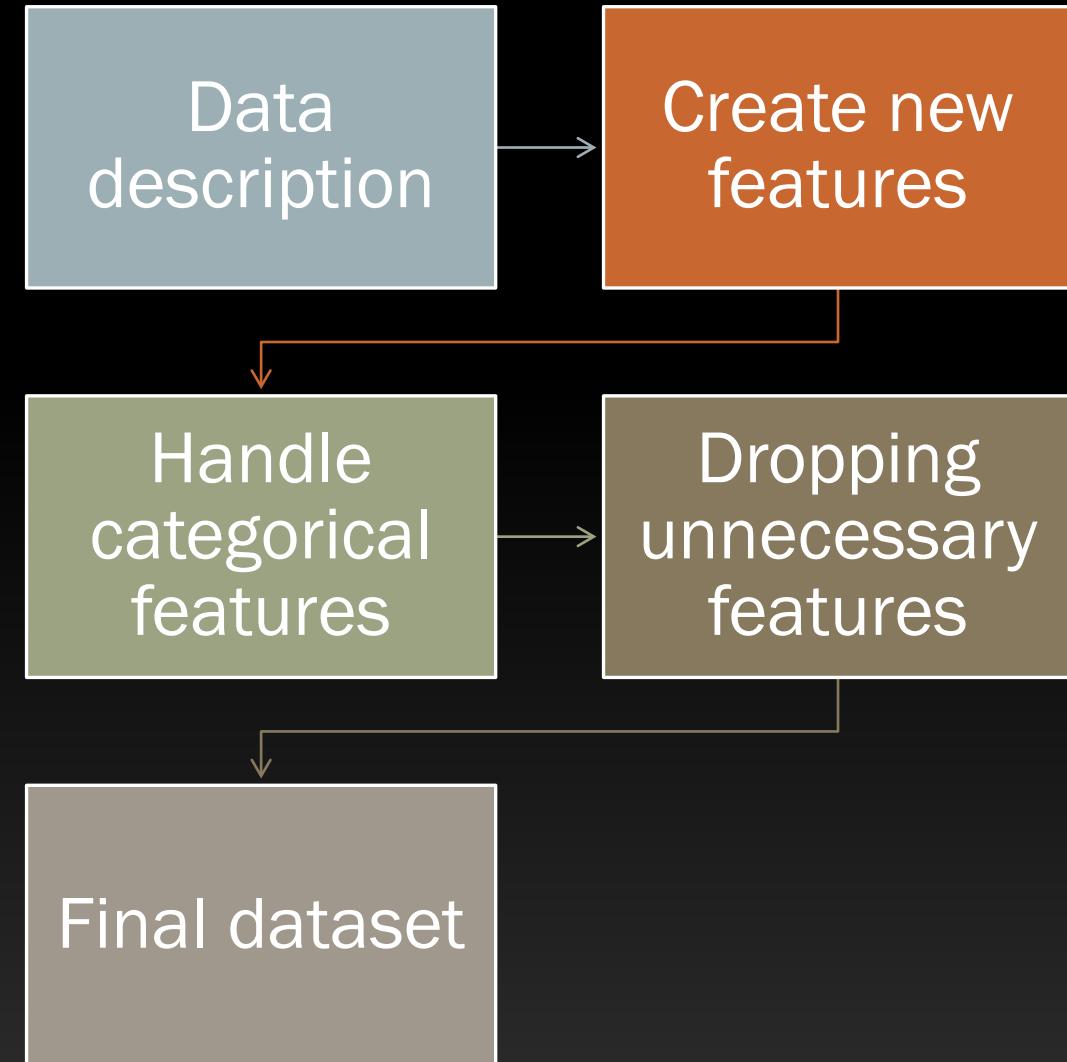


Pearson method



Correlations matrix

Preprocessing



Create new features

Calculate the distance between two locations based on **Haversine formula**.

$$d = 2r \arcsin \left(\sqrt{\sin^2 \left(\frac{\phi_2 - \phi_1}{2} \right) + \cos(\phi_1) \cos(\phi_2) \sin^2 \left(\frac{\lambda_2 - \lambda_1}{2} \right)} \right)$$

Calculate average speed

$$v = \frac{\Delta x}{\Delta t} = \frac{x_2 - x_1}{t_2 - t_1}$$

Calculate altitude

$$E = \arctan \left[\frac{\cos(G) \cos(L) - .1512}{\sqrt{1 - \cos^2(G) \cos^2(L)}} \right]$$

E = elevation of antenna in degrees

S = satellite longitude in degrees

N = site longitude in degrees

L = site latitude in degrees

G = S - N

Handle categorical Features

```
>>> Data info:  
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 182429 entries, 4 to 183411  
Data columns (total 26 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   duration_sec    182429 non-null   int64    
 1   start_time      182429 non-null   object    
 2   end_time        182429 non-null   object    
 3   start_station_id 182429 non-null   float64  
 4   start_station_name 182429 non-null   object    
 5   start_station_latitude 182429 non-null   float64  
 6   start_station_longitude 182429 non-null   float64  
 7   end_station_id   182429 non-null   float64  
 8   end_station_name 182429 non-null   object    
 9   end_station_latitude 182429 non-null   float64  
 10  end_station_longitude 182429 non-null   float64  
 11  bike_id         182429 non-null   int64    
 12  user_type       182429 non-null   category  
 13  member_birth_year 182429 non-null   float64  
 14  member_gender    182429 non-null   category  
 15  bike_share_for_all_trip 182429 non-null   category  
 16  age              182429 non-null   float64  
 17  duration_min    182429 non-null   float64  
 18  start_station_latitude_radians 182429 non-null   float64  
 19  start_station_longitude_radians 182429 non-null   float64  
 20  end_station_latitude_radians 182429 non-null   float64  
 21  end_station_longitude_radians 182429 non-null   float64  
 22  latitude_diff   182429 non-null   float64  
 23  longitude_diff  182429 non-null   float64  
 24  trip_distance   182429 non-null   float64  
 25  avg_trip_speed  182429 non-null   float64  
dtypes: category(3), float64(17), int64(2), object(4)  
memory usage: 38.0+ MB  
None
```

Dropping unnecessary features & Order columns



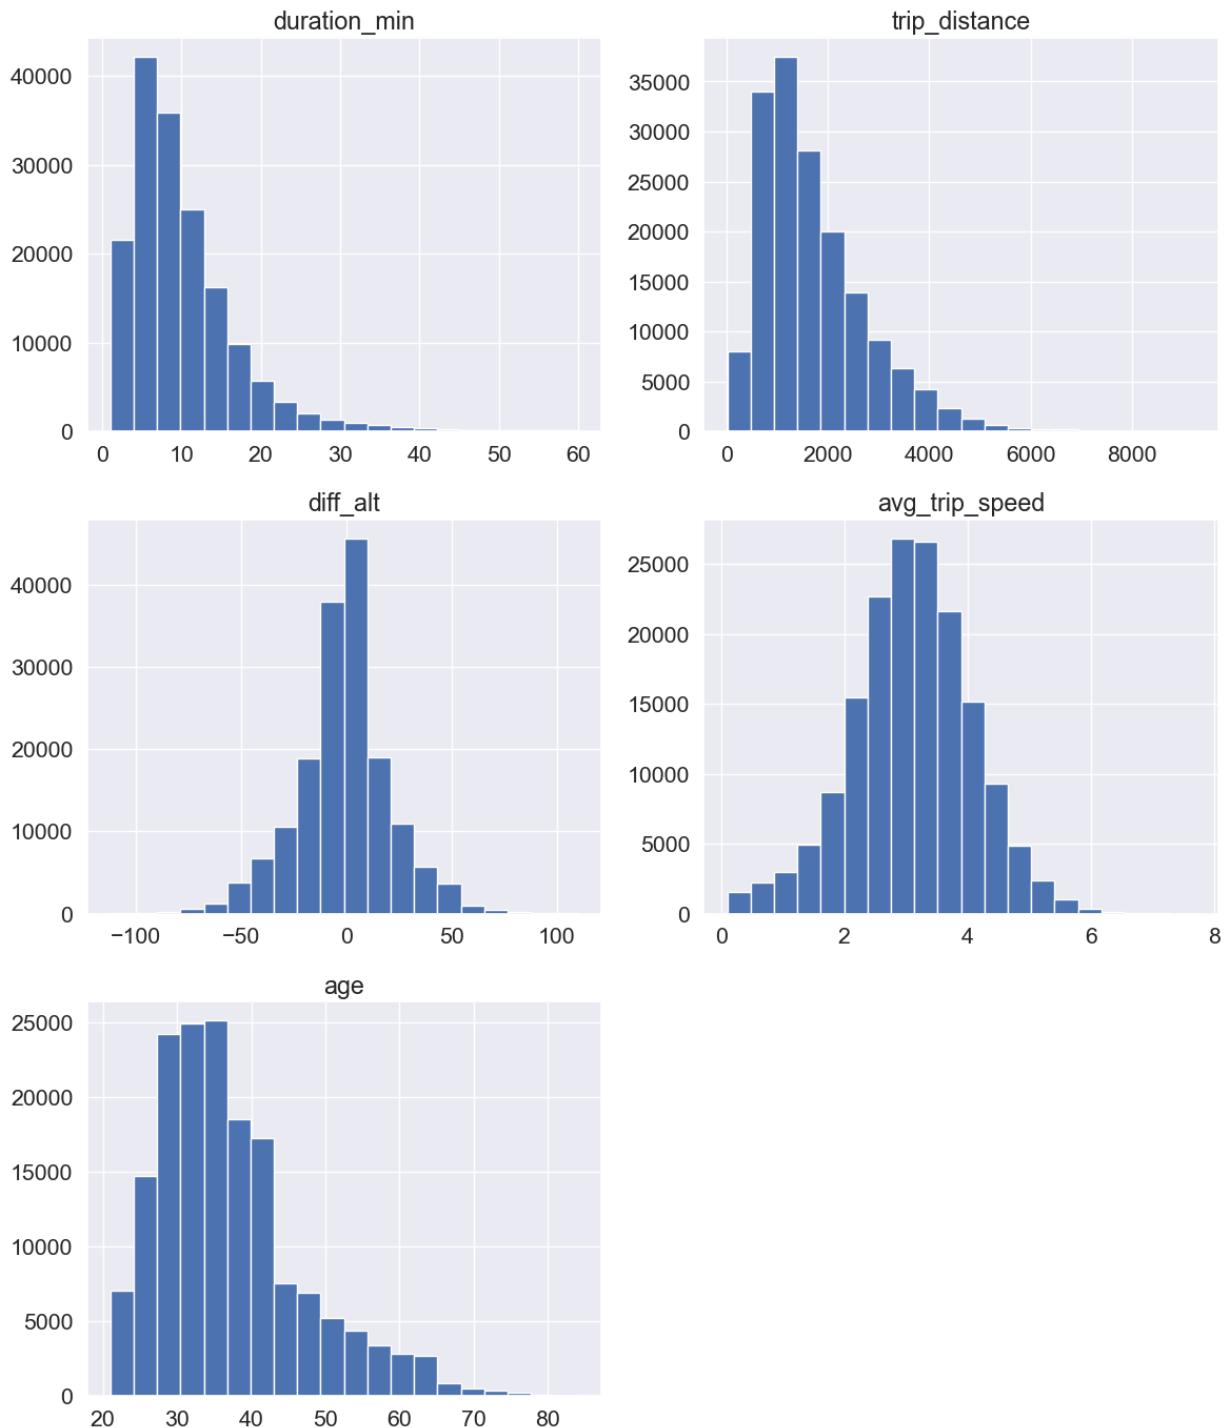
```
1 order_columns = [
2     'bike_id',
3     'start_station_id',
4     'start_station_name',
5     'end_station_id',
6     'end_station_name',
7     'trip_distance',
8     'diff_alt',
9     'duration_min',
10    'avg_trip_speed',
11    'counts_start',
12    'counts_end',
13    'user_type',
14    'member_gender',
15    'age',
16    'bike_share_for_all_trip'
17 ]
18
19 fdf = fdf[order_columns]
20 fdf.to_csv(r'fdf.csv', index=False)
21 fdf.columns
```



```
1 drop_columns_list = [
2     'duration_sec',
3     'start_time',
4     'end_time',
5     '# 'start_station_id',
6     '# 'start_station_name',
7     'start_station_latitude',
8     'start_station_longitude',
9     '# 'end_station_id',
10    '# 'end_station_name',
11    'end_station_latitude',
12    'end_station_longitude',
13    '# 'bike_id',
14    '# 'user_type',
15    'member_birth_year',
16    '# 'member_gender',
17    '# 'bike_share_for_all_trip',
18    '# 'age',
19    '# 'duration_min',
20    'start_station_latitude_radians',
21    'start_station_longitude_radians',
22    'end_station_latitude_radians',
23    'end_station_longitude_radians',
24    'latitude_diff',
25    'longitude_diff',
26    '# 'trip_distance',
27    '# 'avg_trip_speed',
28    'station_latitude_start',
29    'station_longitude_start',
30    '# 'counts_start',
31    'station_altitude_start',
32    'station_latitude_end',
33    'station_longitude_end',
34    '# 'counts_end',
35    'station_altitude_end',
36    '# 'diff_alt'
37 ]
```

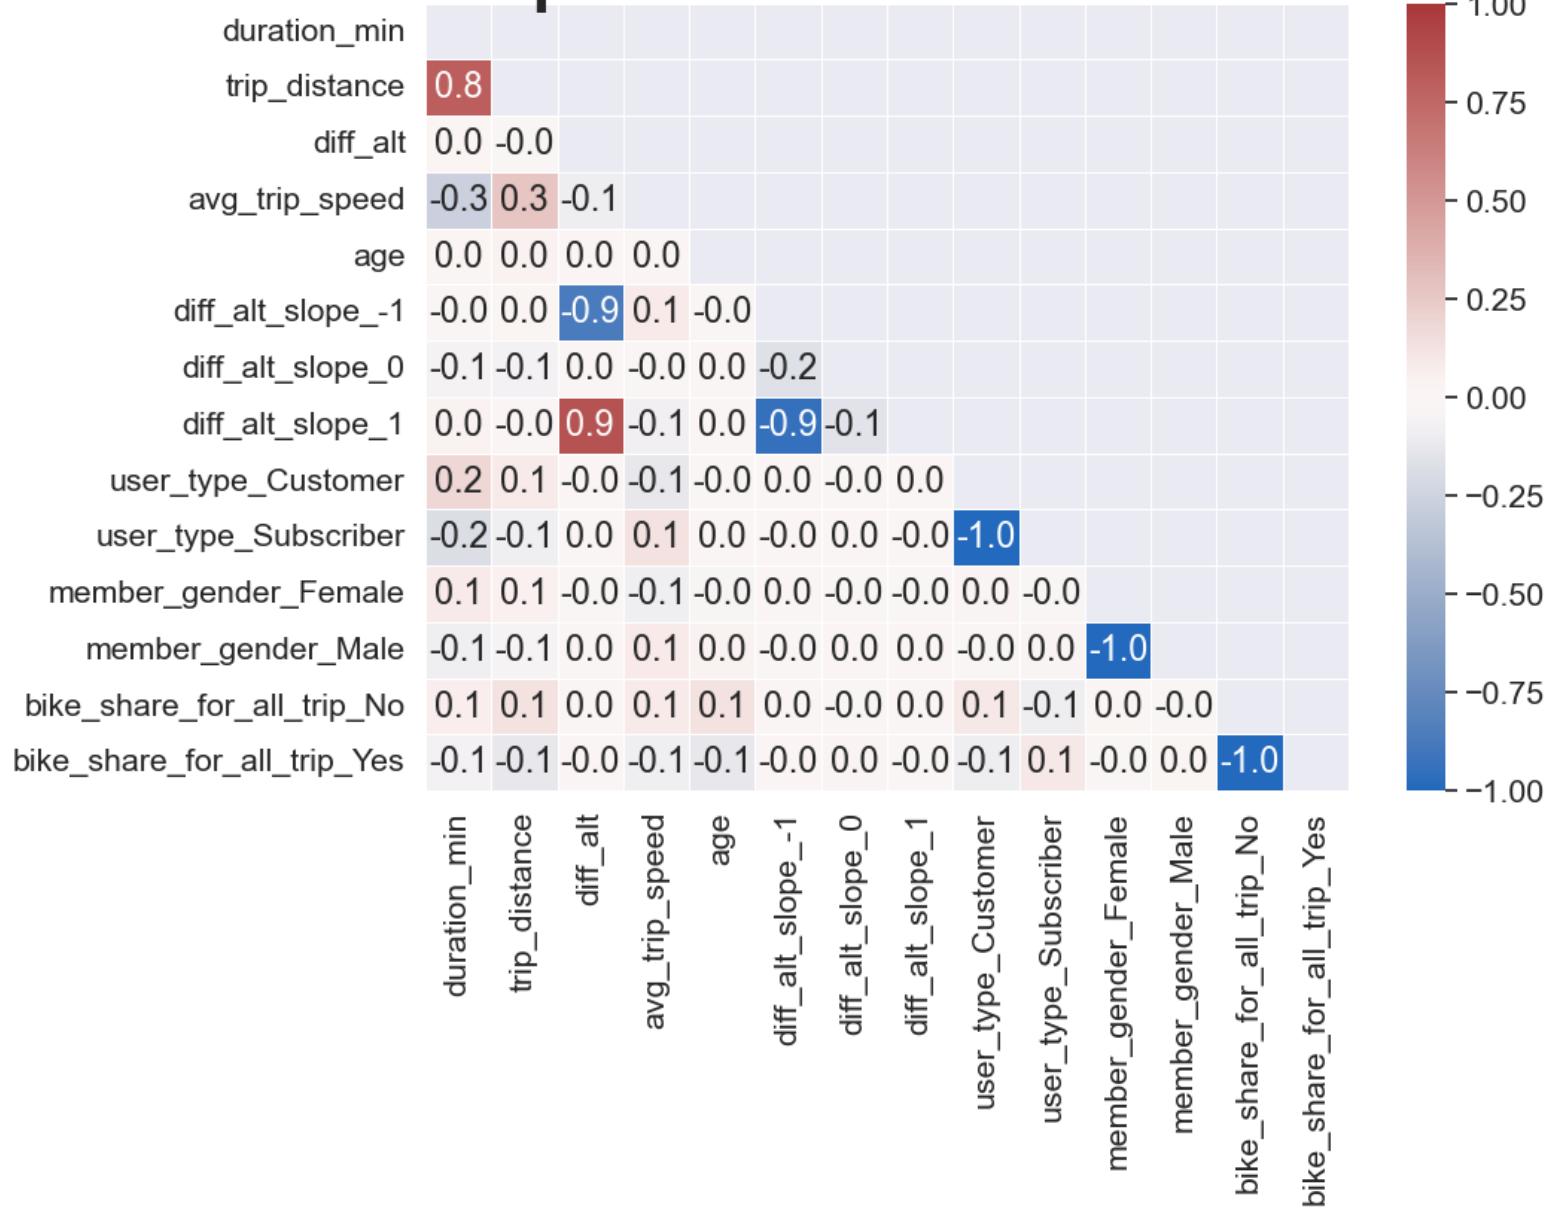
```
[115] 1 cdf = pdf.drop(drop_columns_list, axis=1)
2 cdf
```

Features Distributions

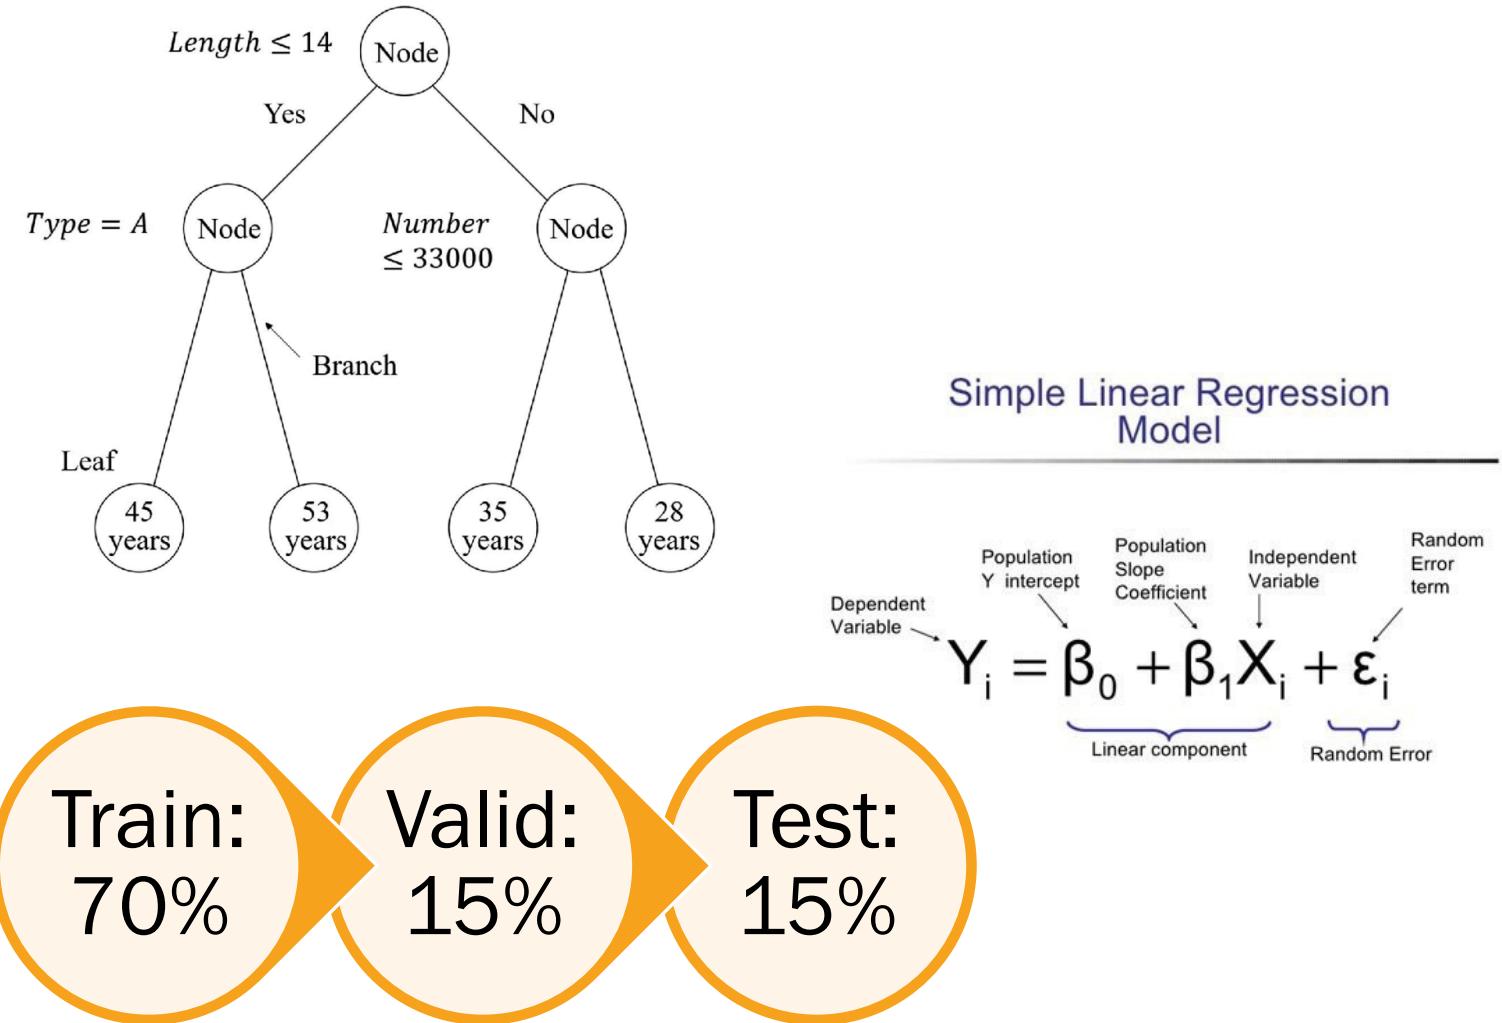




Spearman method

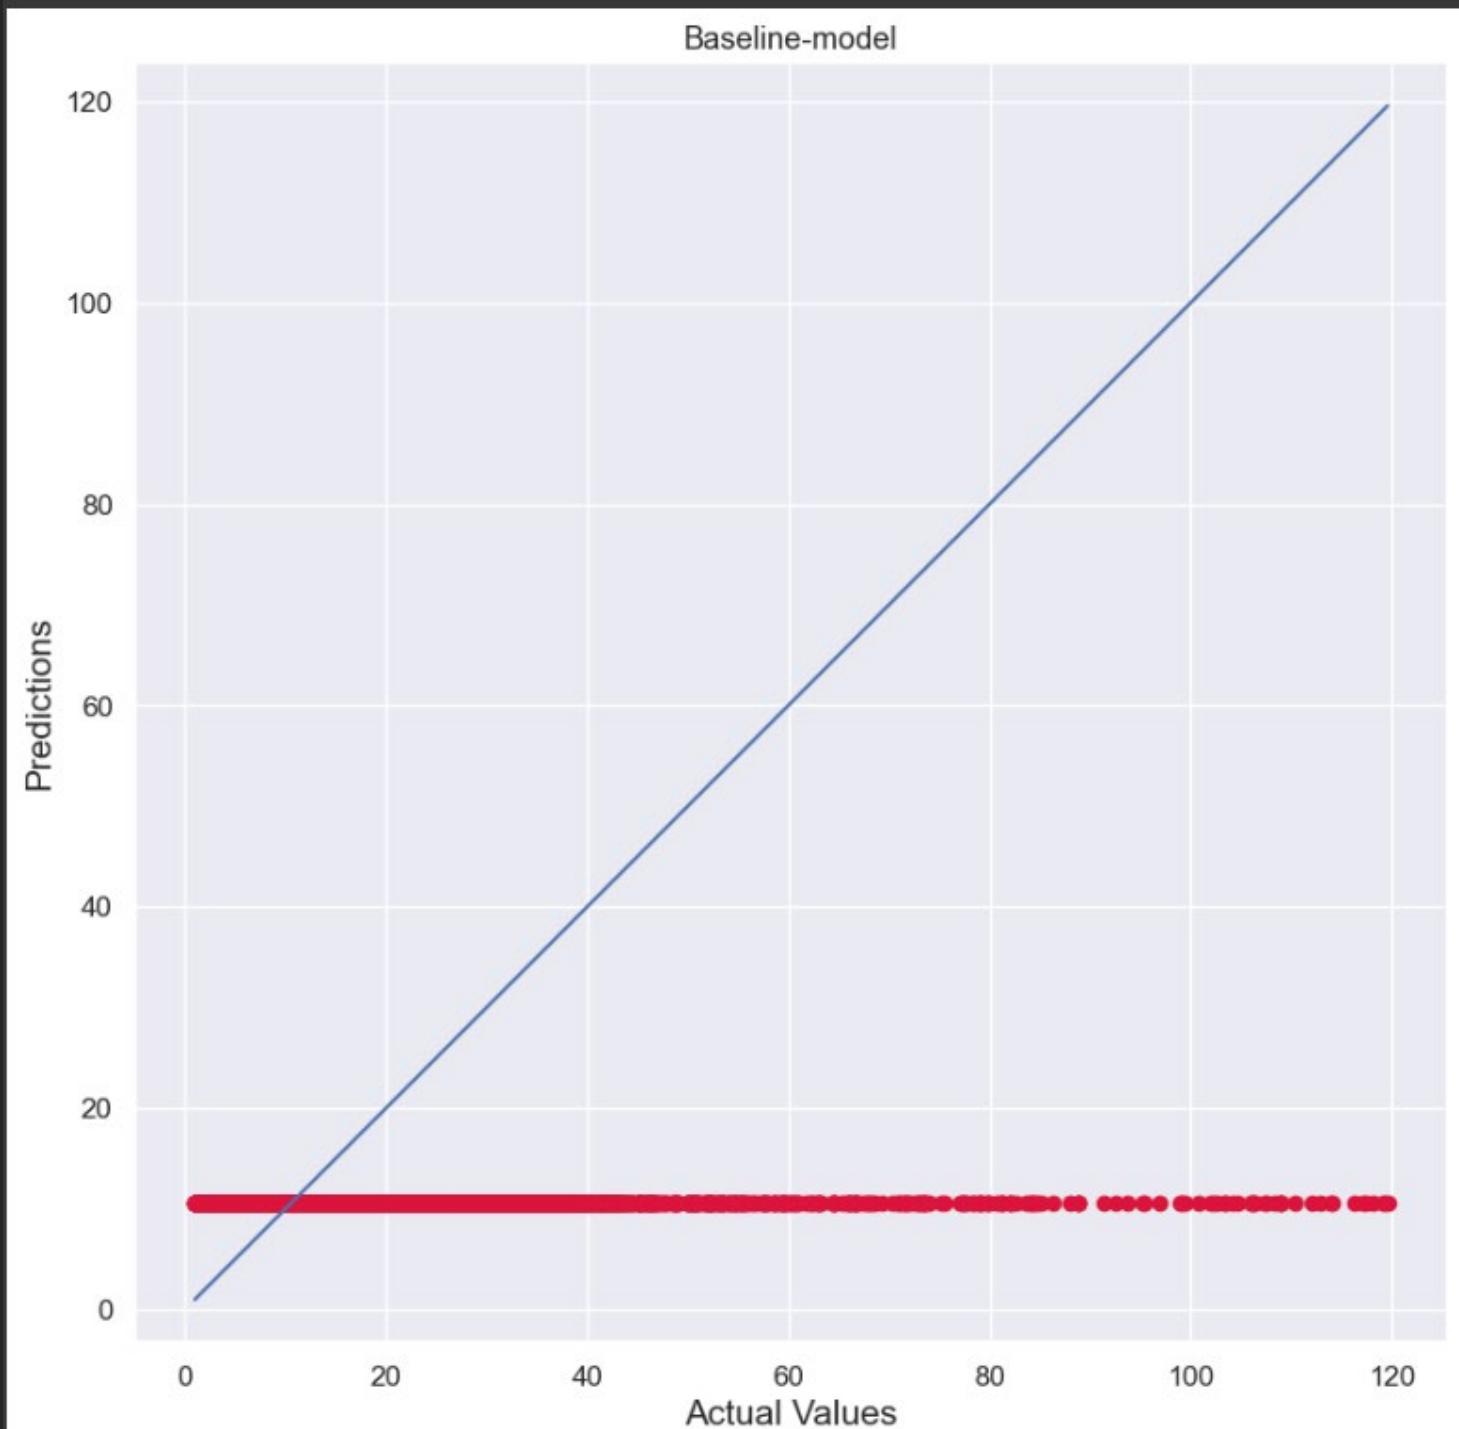


Regression models



Baseline model based on .mean()

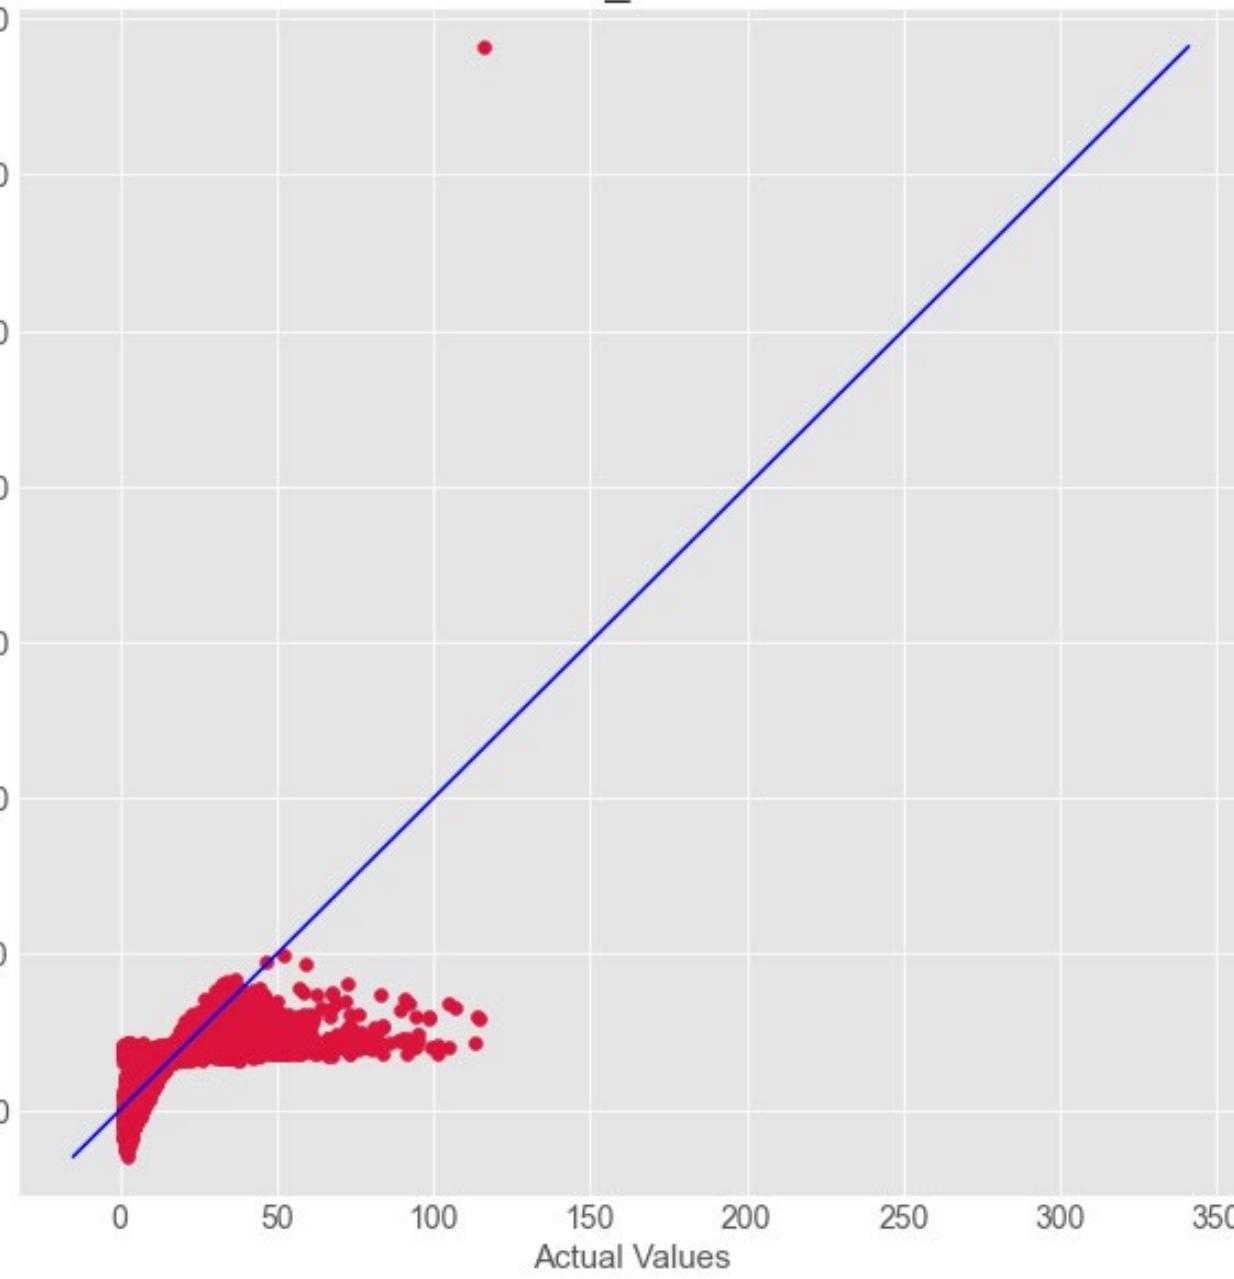
MAE: 5.654298977198291
MSE: 82.18859531479465
RMSE: 9.06579259164882
R2: 0



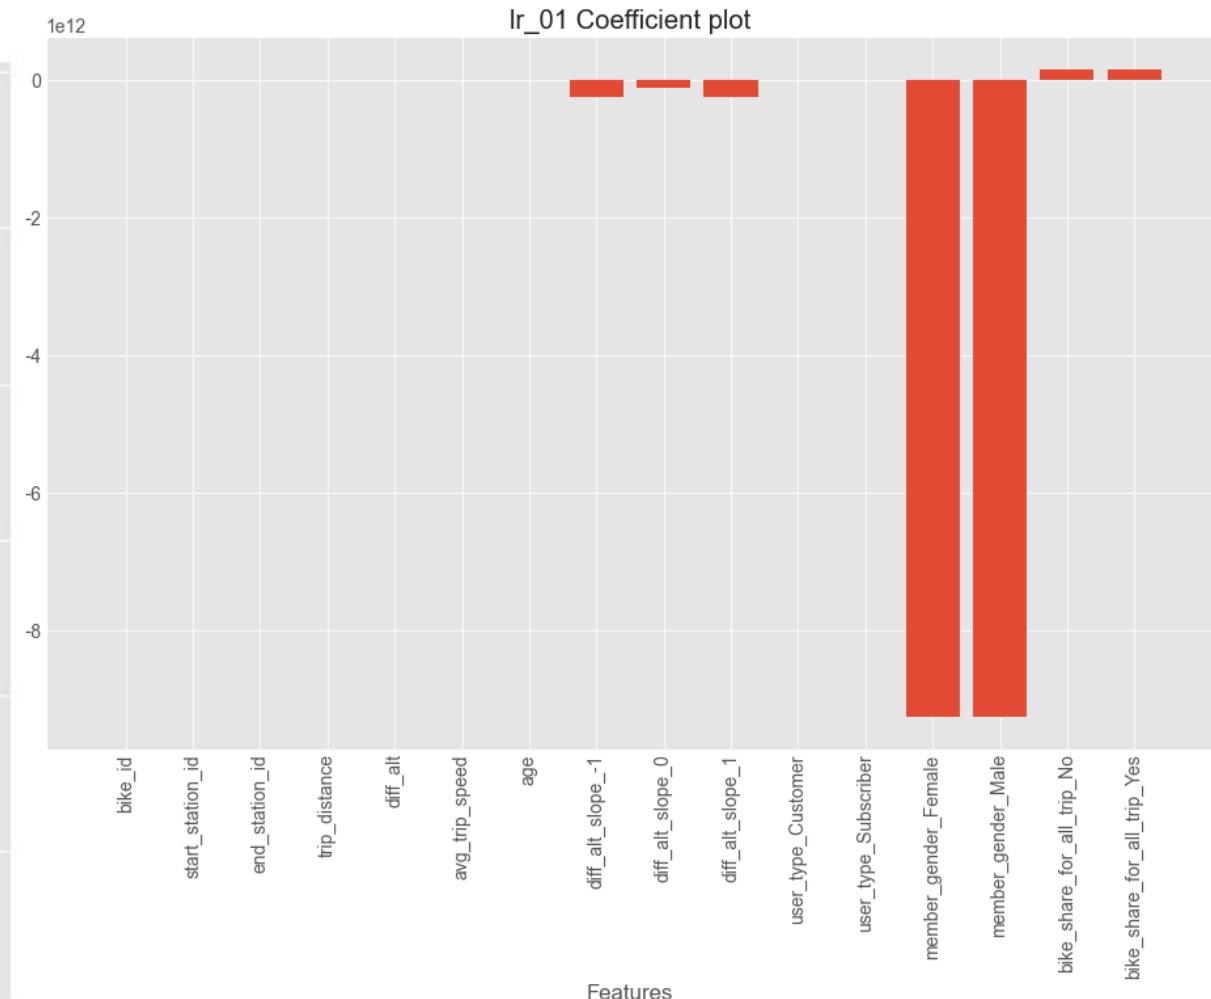


lr_01

Predictions

 $MAE = 2.70,$ $MSE = 31.59,$ $RMSE = 5.62,$ $R^2 = 0.573$

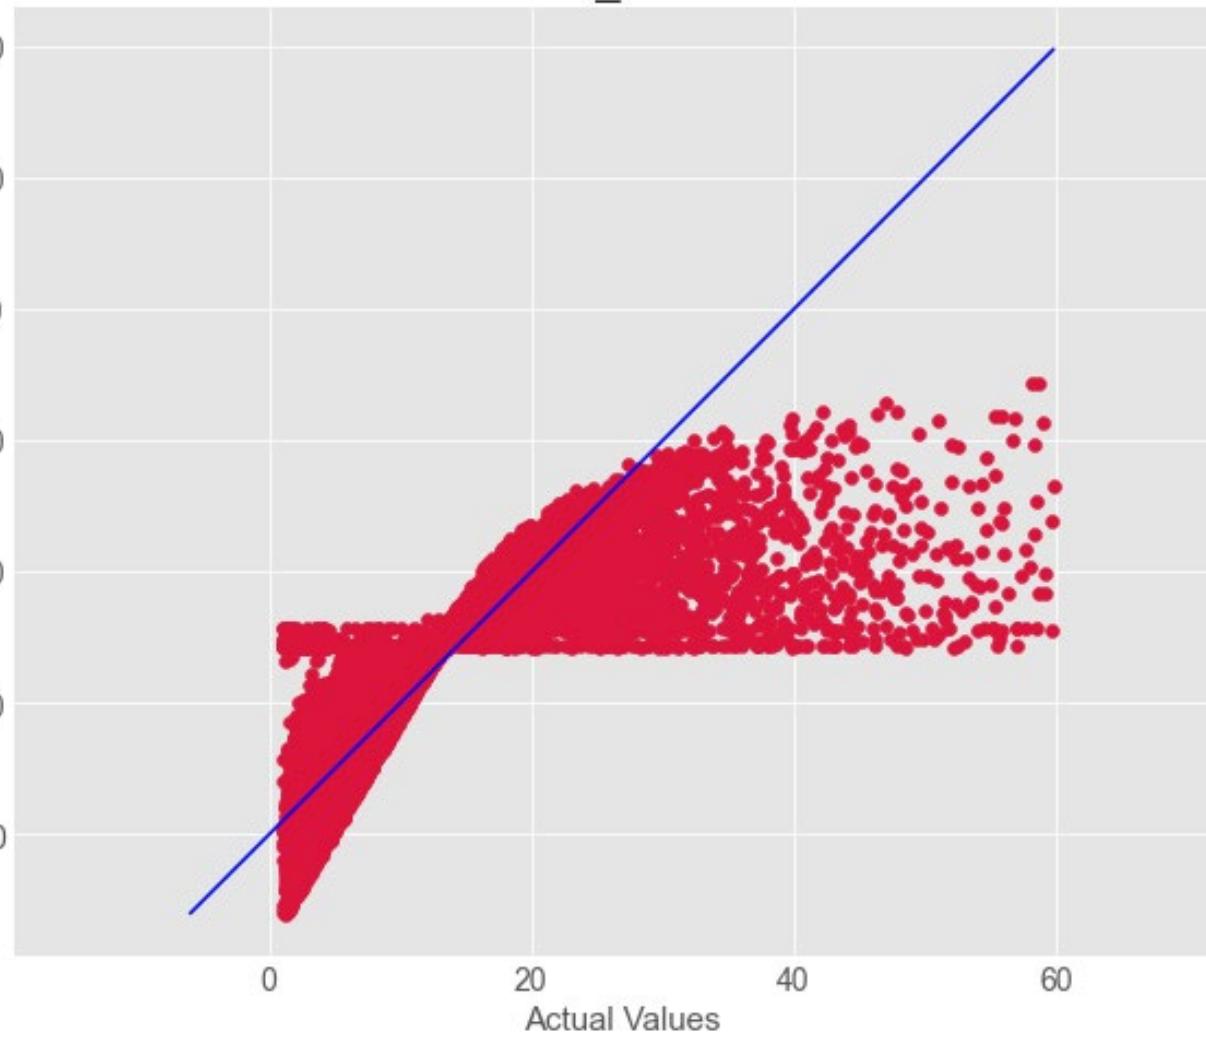
lr_01 Coefficient plot



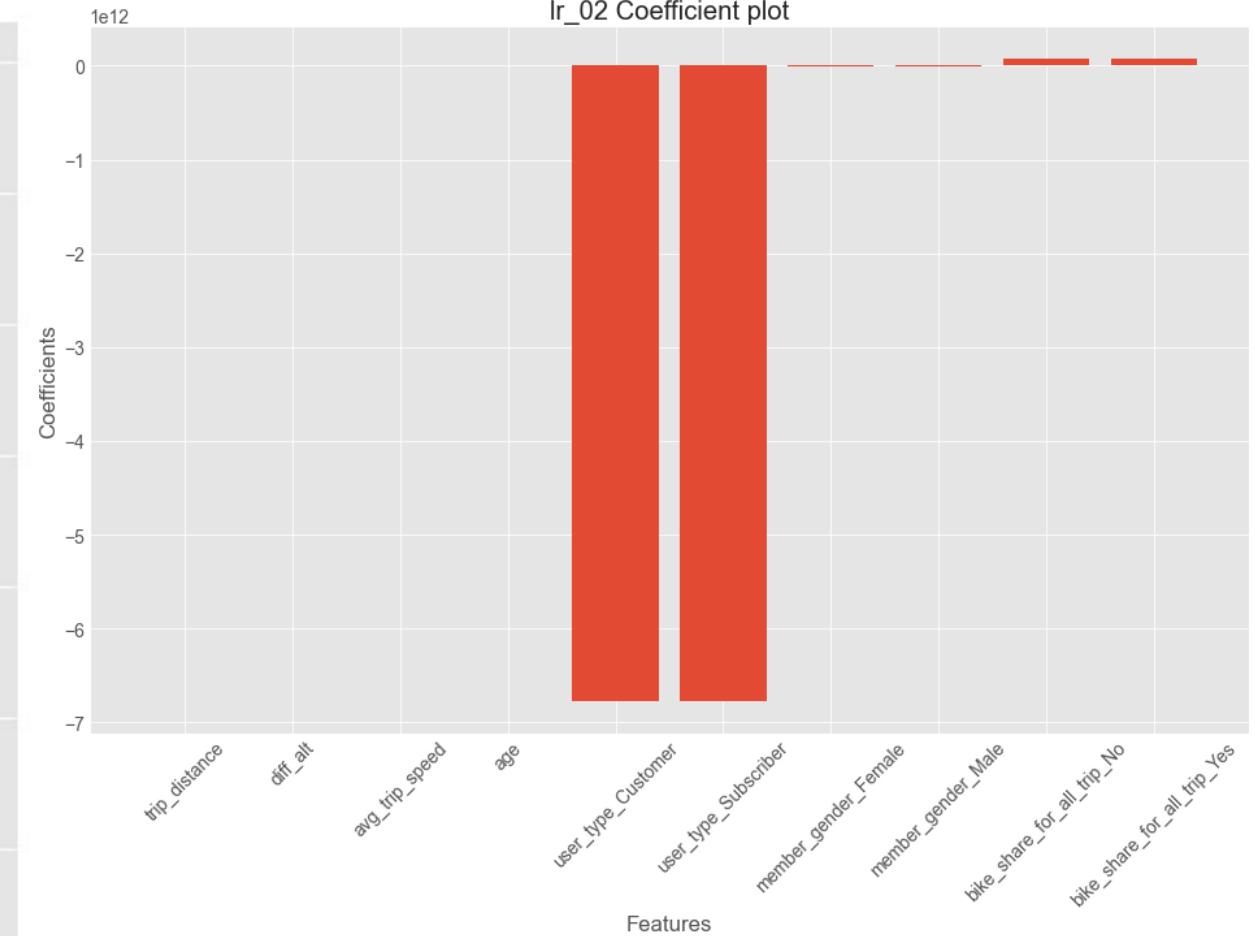


lr_02

Predictions



lr_02 Coefficient plot



$MAE = 2.27,$

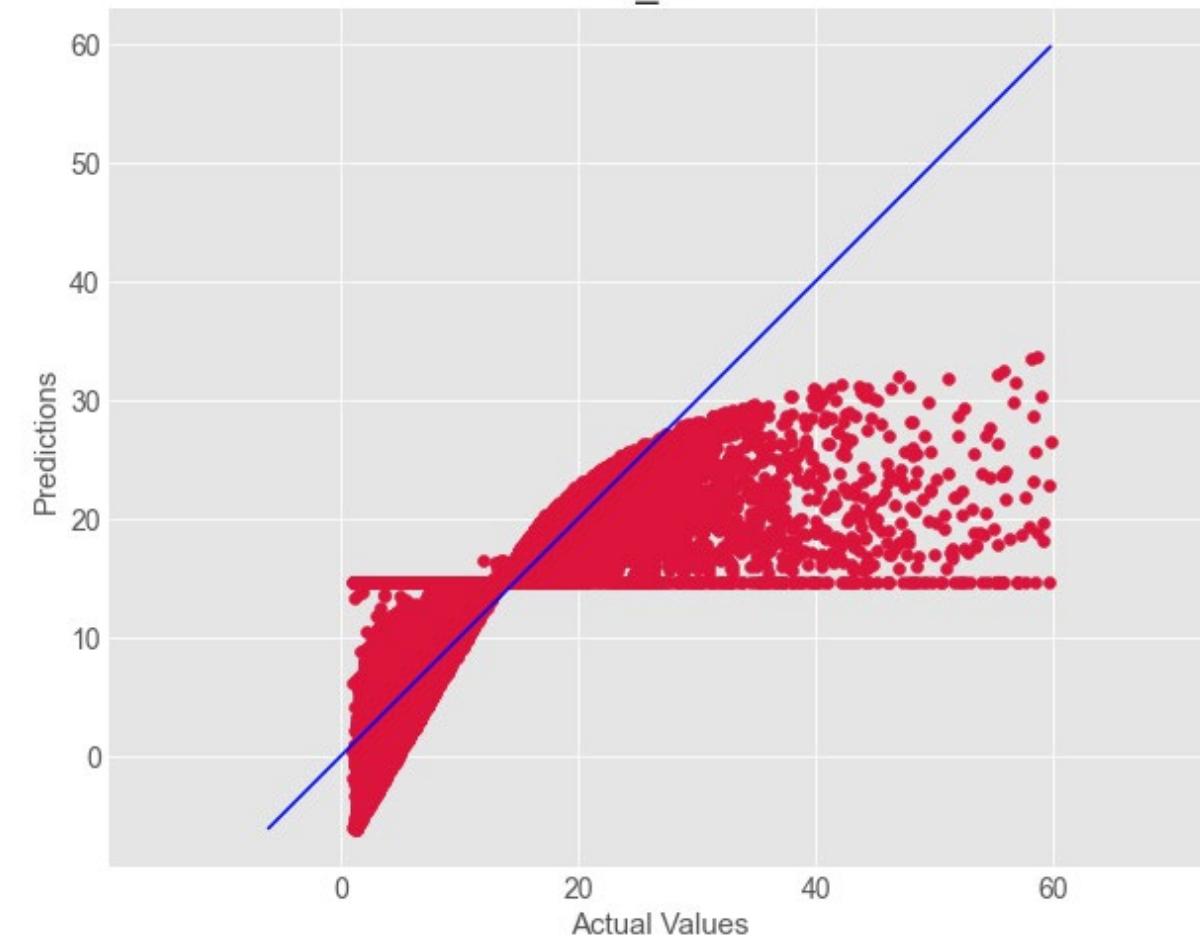
$MSE = 17.19,$

$RMSE = 4.14,$

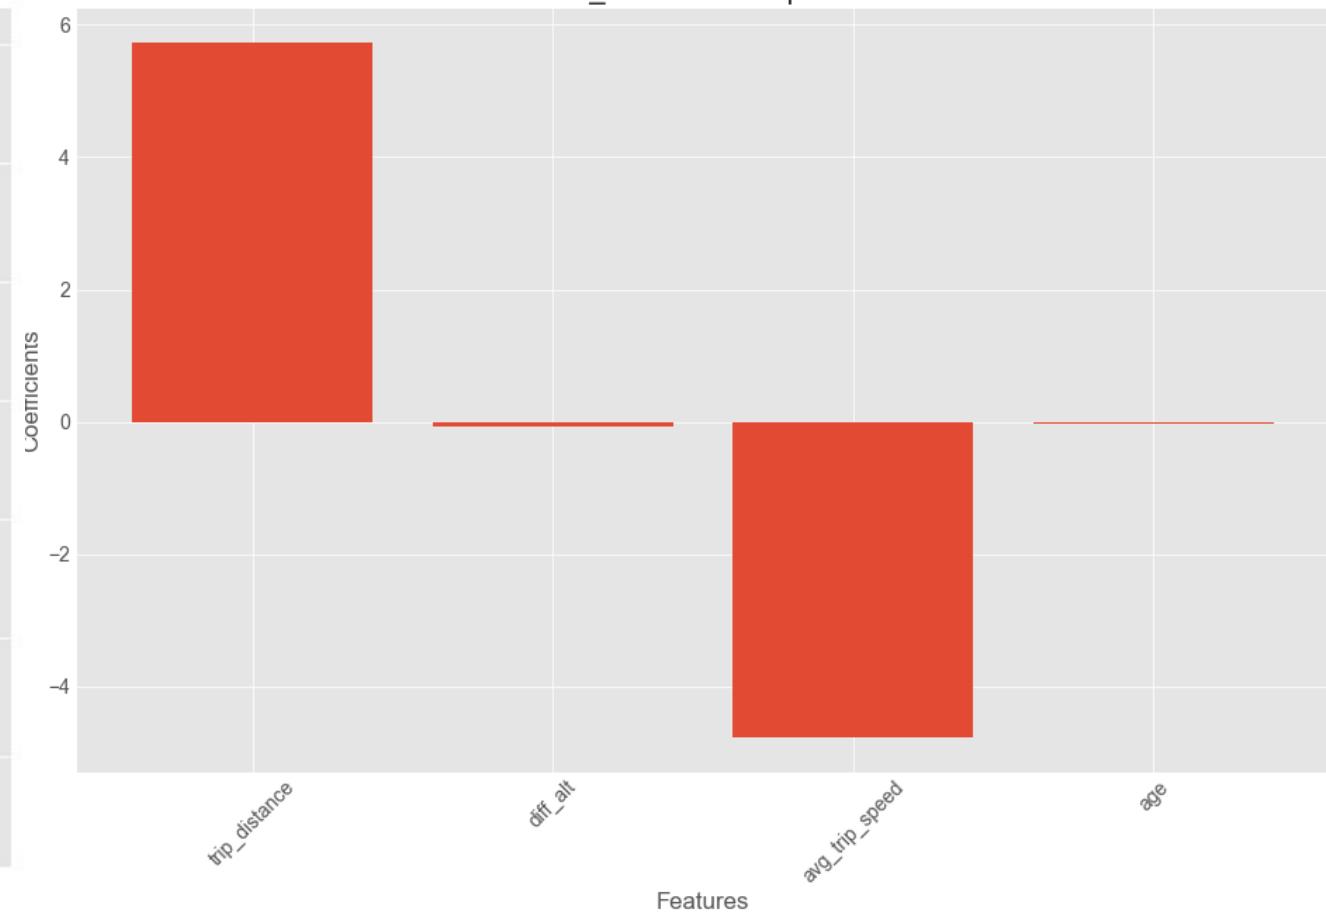
$R^2 = 0.69$



lr_03



lr_03 Coefficient plot



Validation: $MAE = 2.890$,
Test: $MAE = 2.256$,

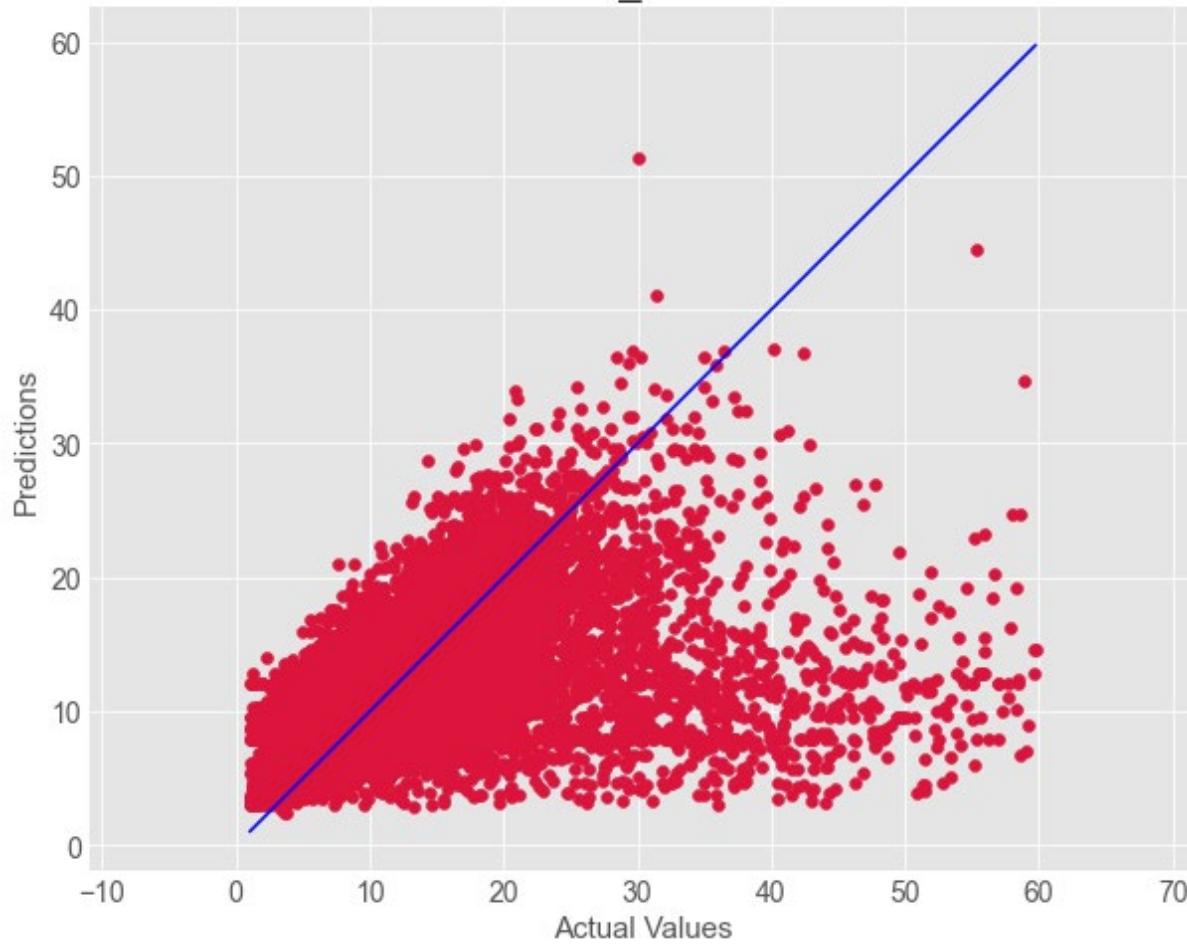
$MSE = 17.381$,
 $MSE = 16.877$,

$RMSE = 4.169$,
 $RMSE = 4.108$,

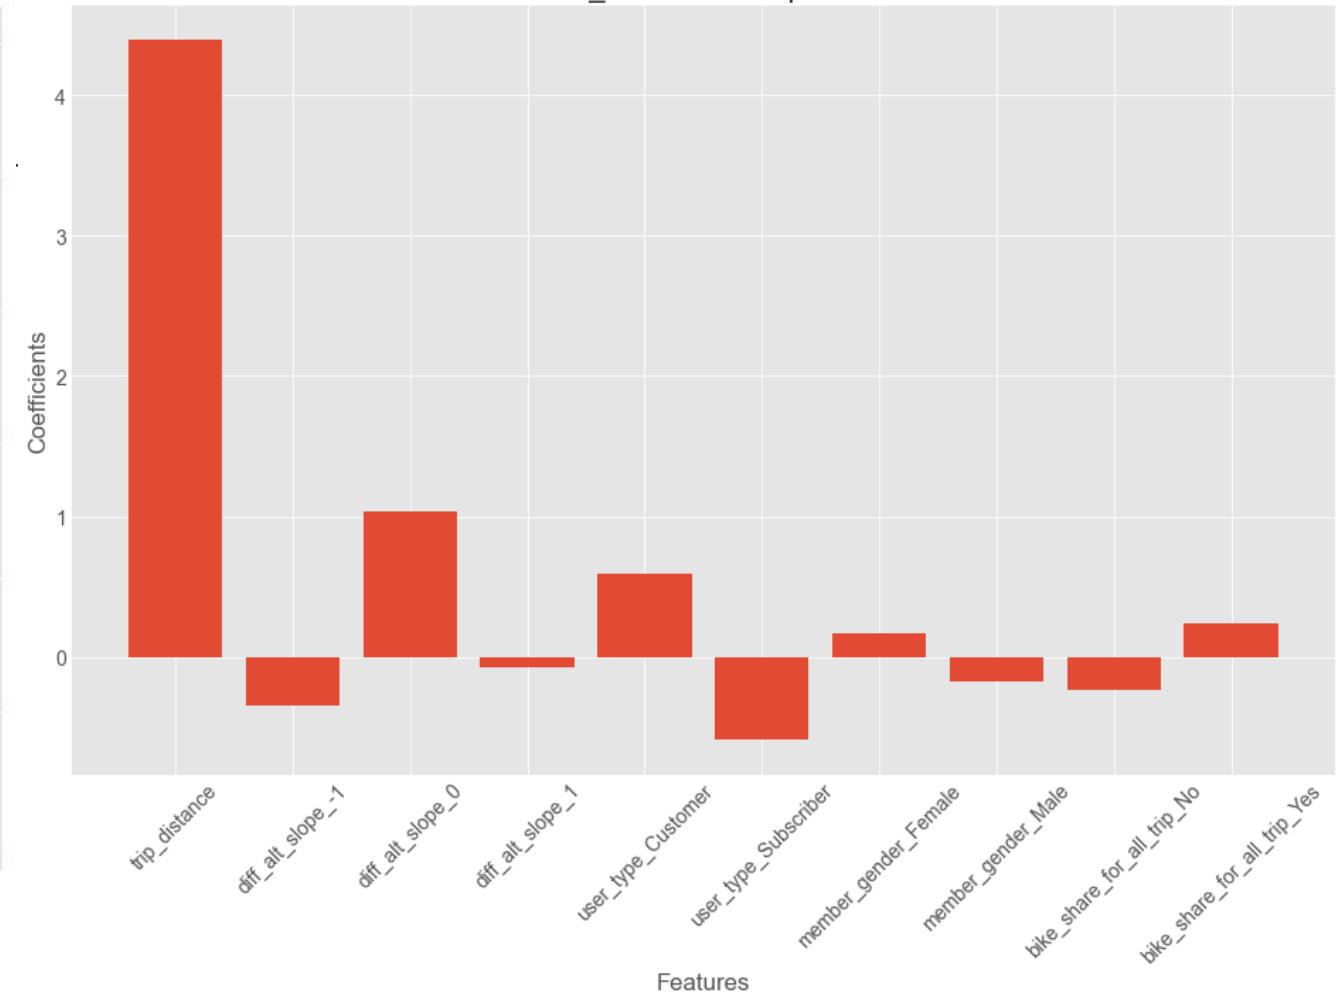
$R^2 = 0.688$
 $R^2 = 0.691$



lr_04



lr_04 Coefficient plot



Validation: MAE = 3.295,
Test: MAE = 3.266,

MSE = 34.803,
MSE = 33.864,

RMSE = 5.899,
RMSE = 5.819,

$R^2 = 0.377$
 $R^2 = 0.381$

Decision-Tree Modeling

This phase split up into two parts:

1. A dataset with clipping into the limits (clip_df)
2. A dataset without unruly data (cut_df)

Both were examined parallel with the same model attributes, same hyper-parameters and so forth.

Methods we had used:

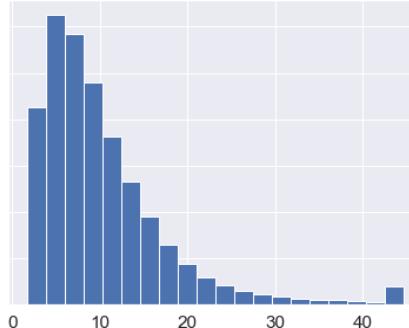
Handle with categorical features via `***get.dummies()***`

Control the attributes of decision regressor: [`random_state`, `tree_n_leaves`]

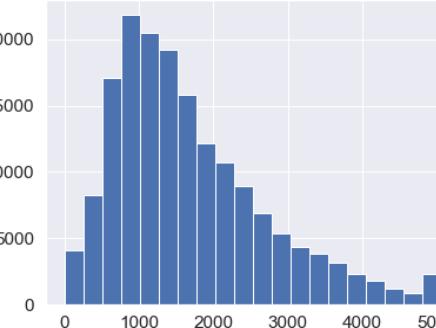
Tuning Hyper-parameters: [`max_depth`, `min_samples_split`, `min_samples_leaf`, `ccp_alpha`]

Clip_df

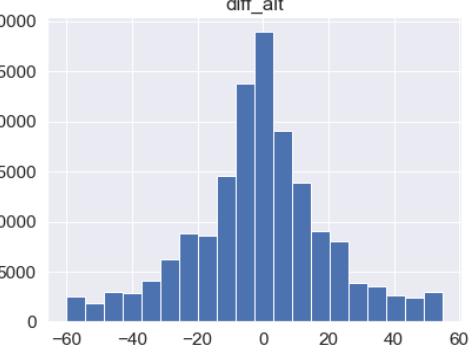
duration_min



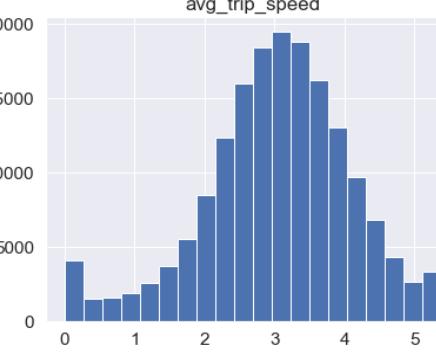
trip_distance



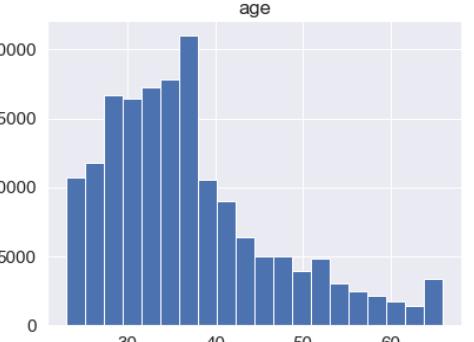
diff_alt



avg_trip_speed



age

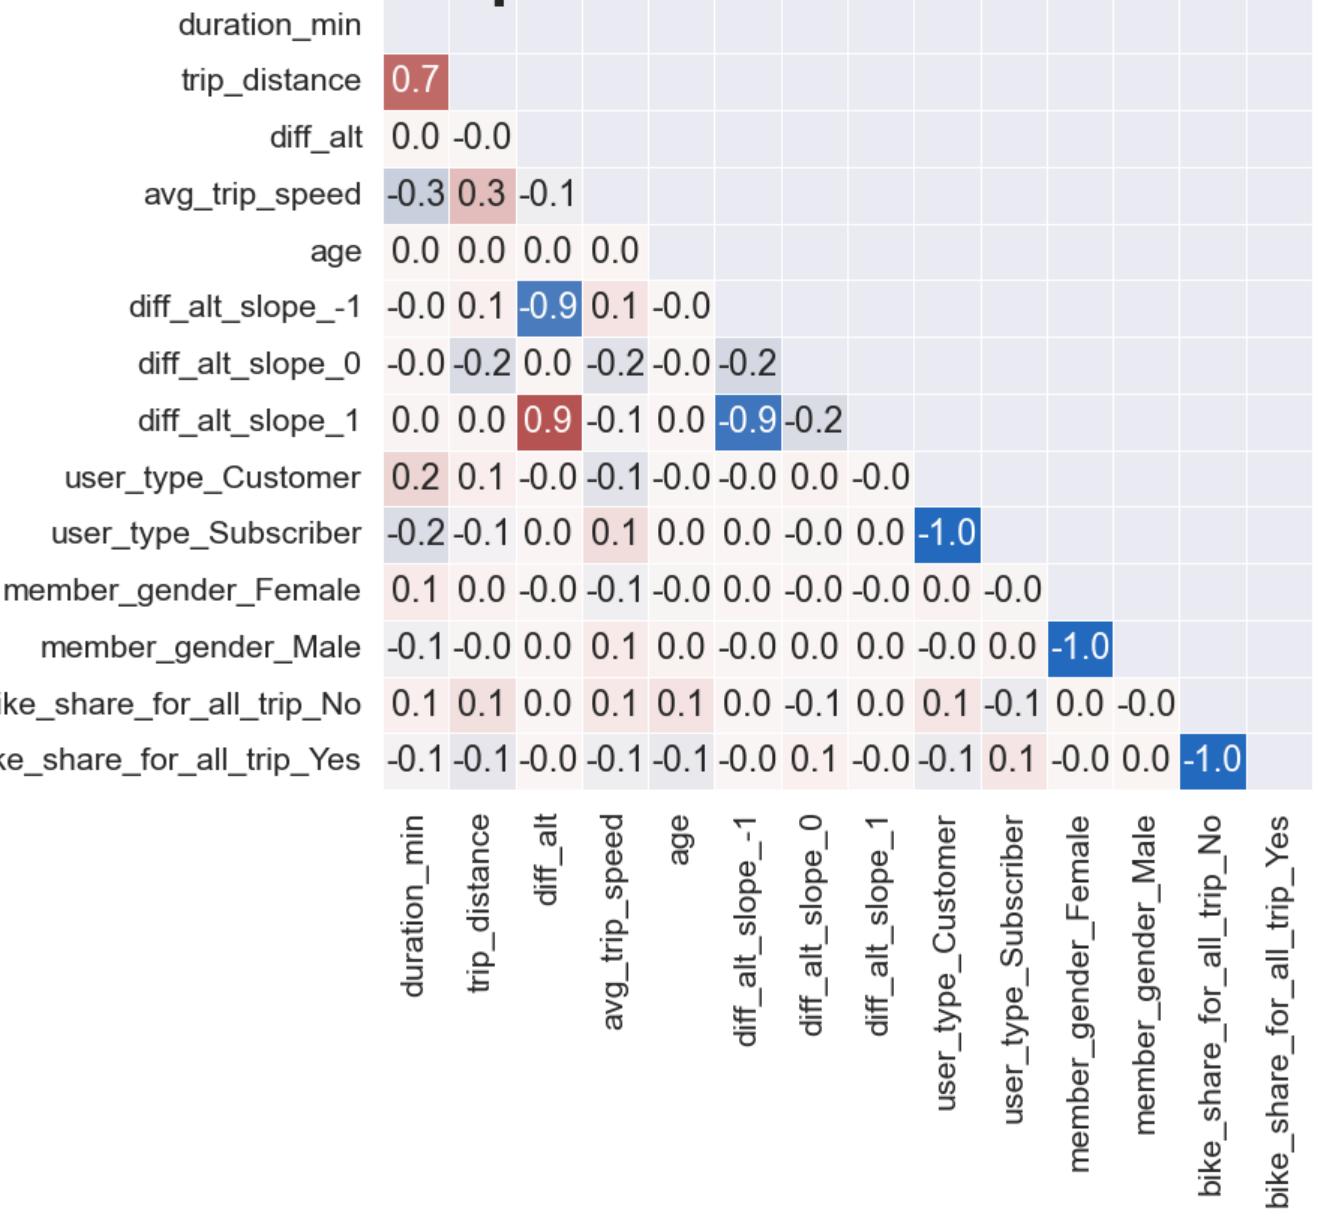


dtr_fdf_clip

(170723, 14)

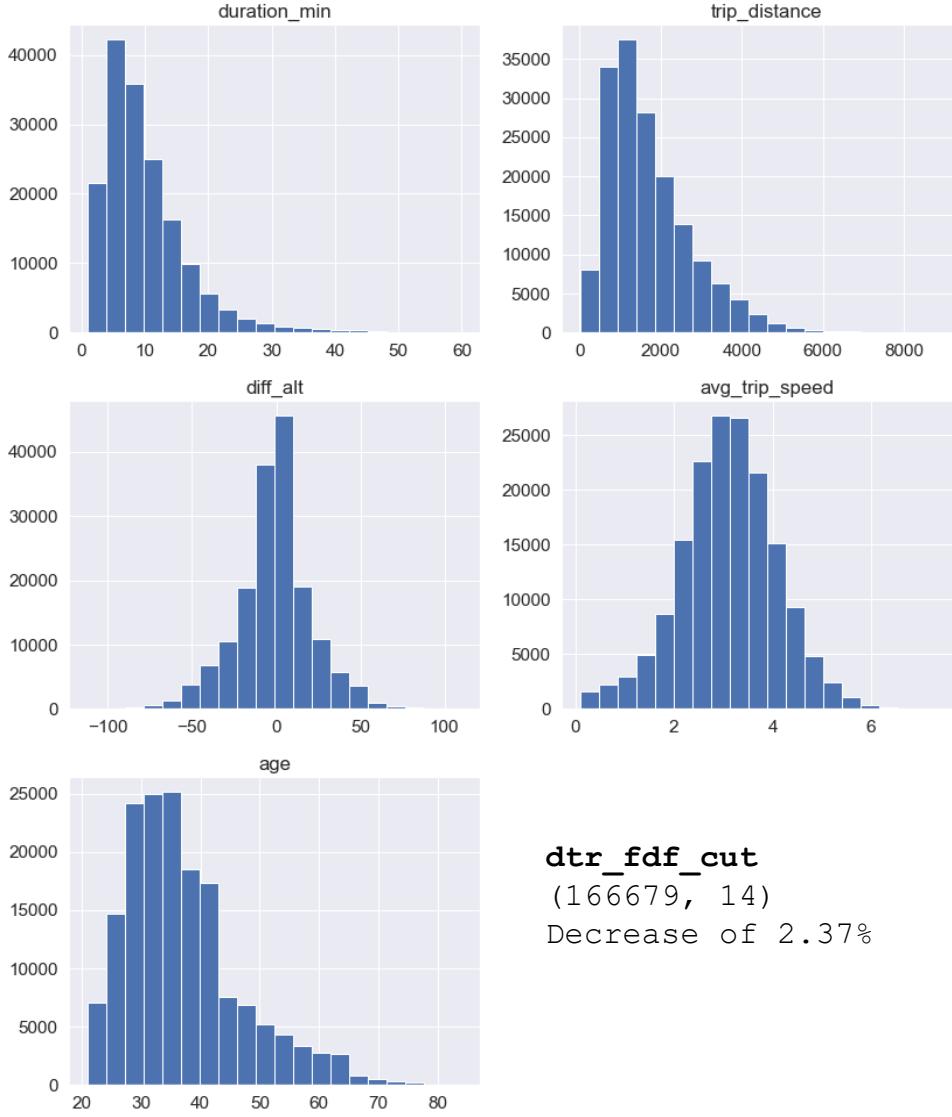
Decrease of 0%

Spearman method



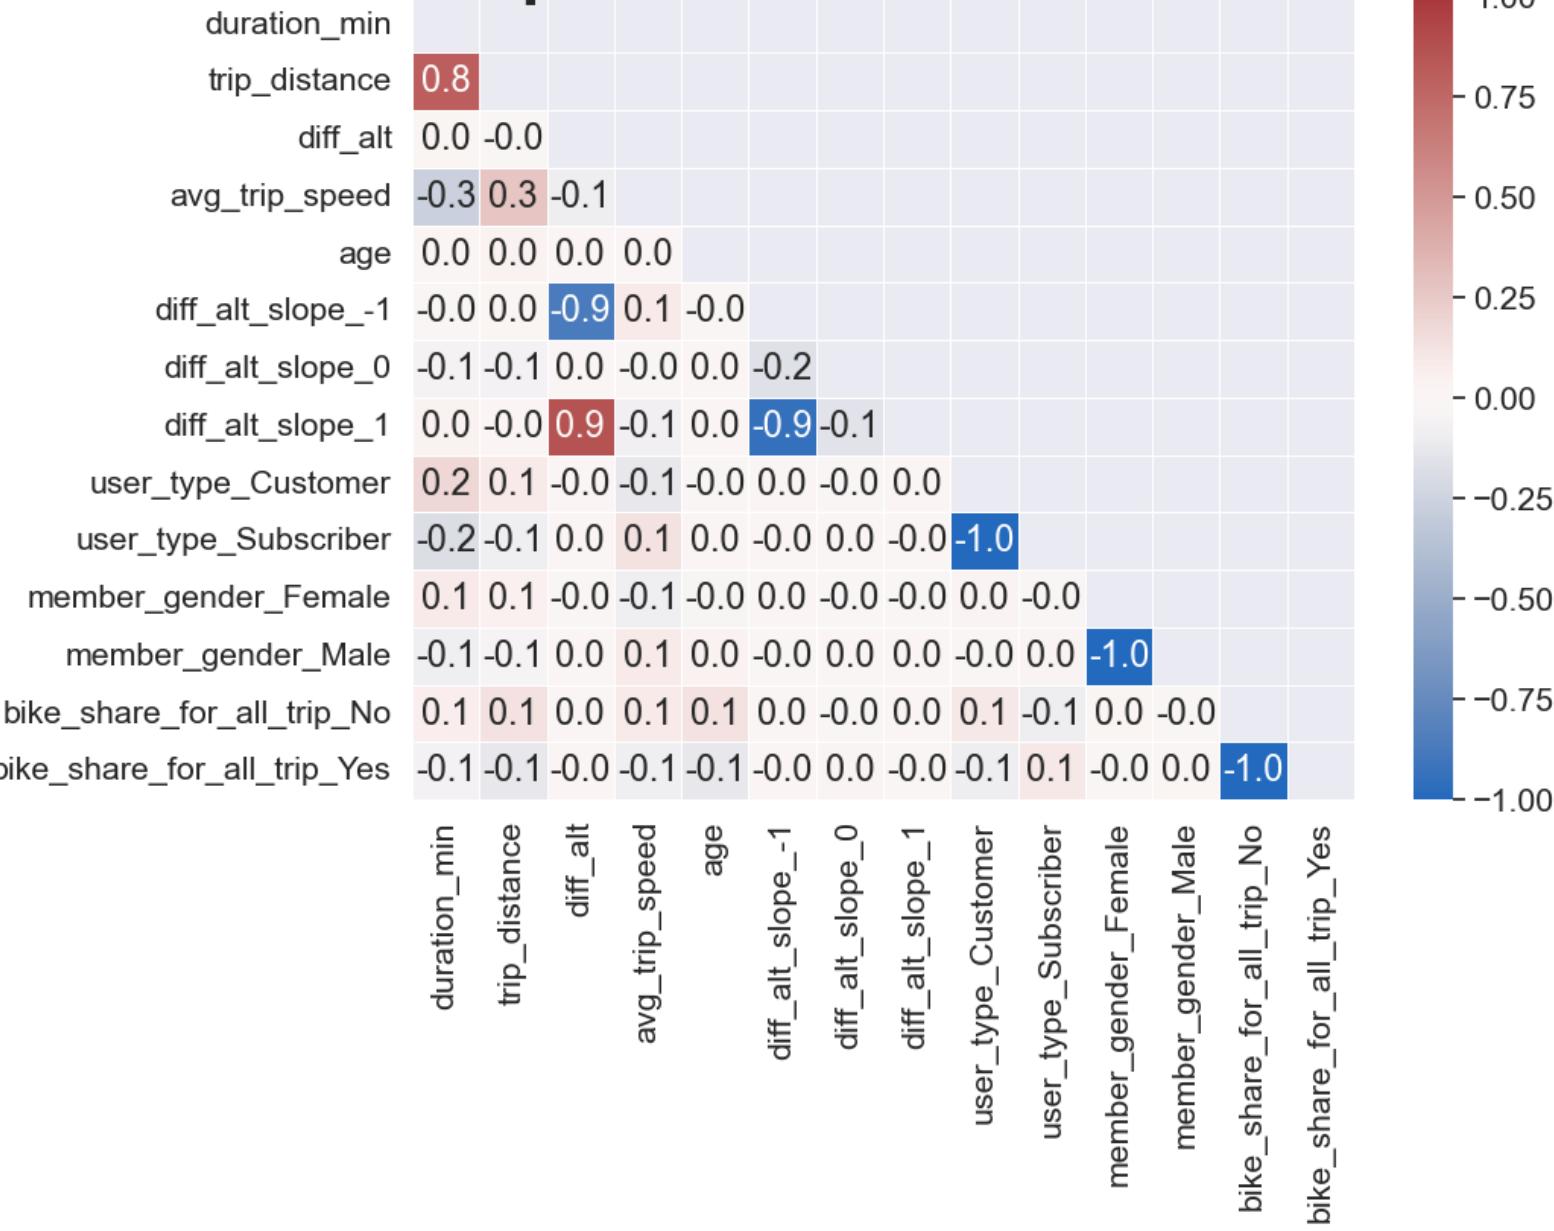
clip_train(119506, 14), clip_valid (25608 ,14), clip_test(25609 ,14)

Cut_df



cut_train(116675, 14), cut_valid (25002 ,14), cut_test(25002 ,14)

Spearman method

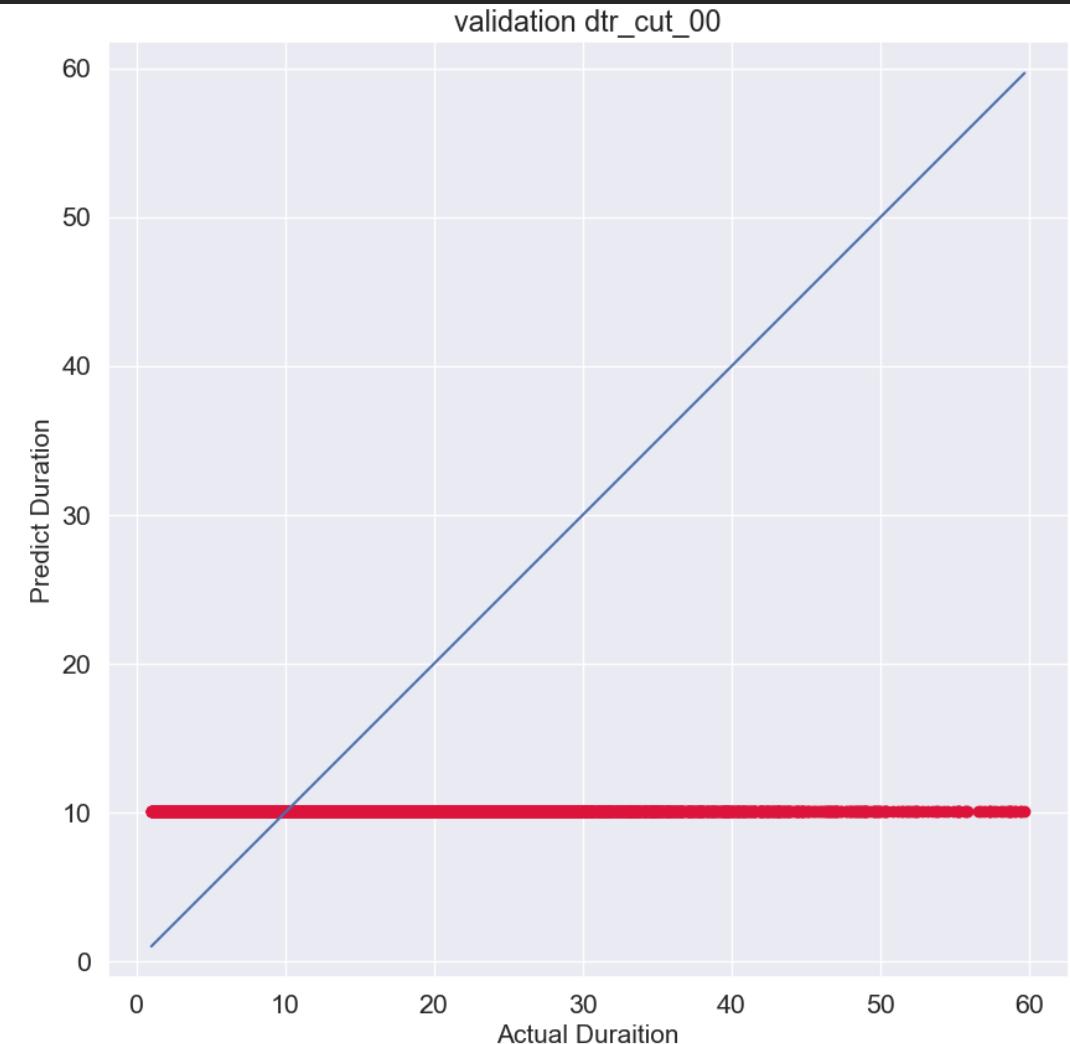
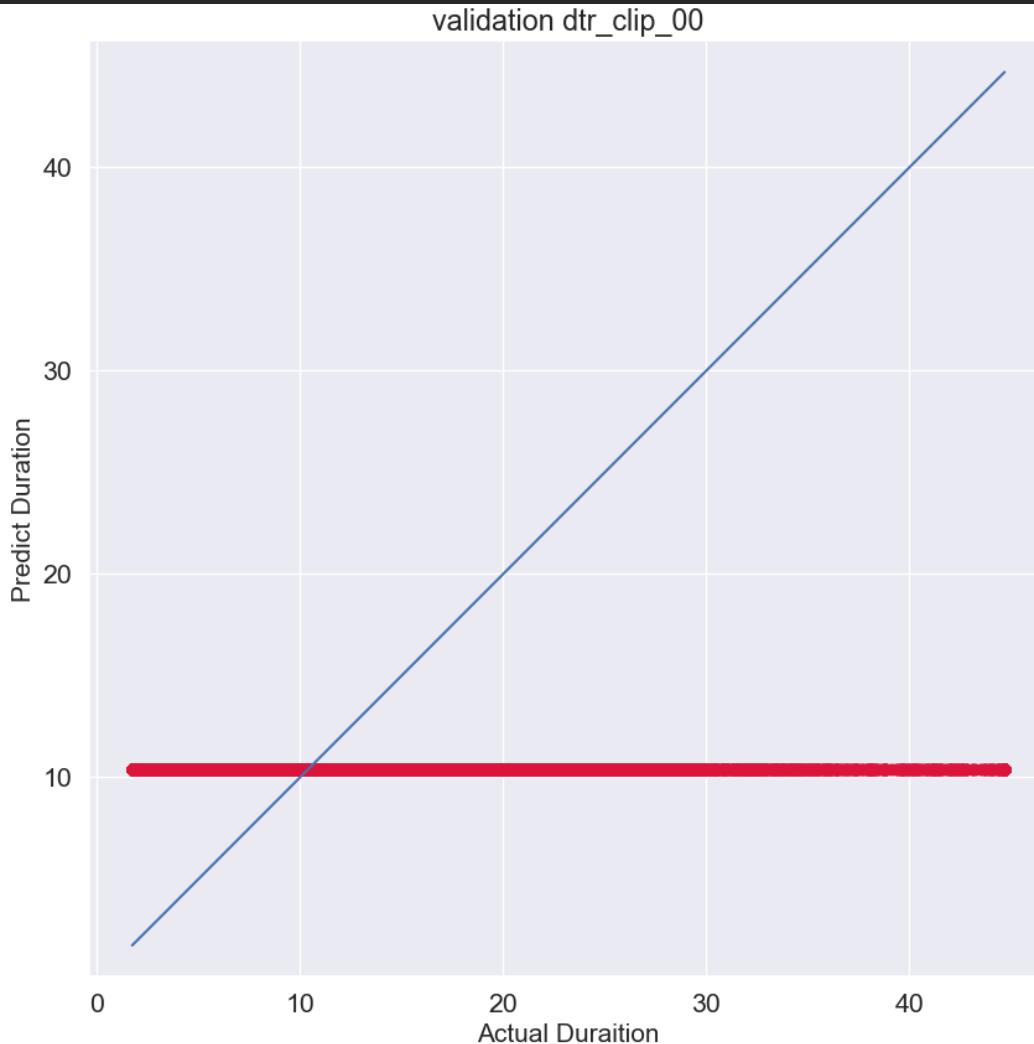


dtr_00

Baseline model based on `.mean()`

Train: $MAE = 5.341$, $MSE = 56.623$, $RMSE = 7.524$, $R^2 = 0.0$
Validation: $MAE = 5.358$, $MSE = 56.675$, $RMSE = 7.528$, $R^2 = 0.0$
Test: $MAE = 5.366$, $MSE = 56.893$, $RMSE = 7.542$, $R^2 = 0.0$

Train: $MAE = 5.026$, $MSE = 49.197$, $RMSE = 7.014$, $R^2 = 0.0$
Validation: $MAE = 5.084$, $MSE = 50.731$, $RMSE = 7.122$, $R^2 = 0.0$
Test: $MAE = 5.052$, $MSE = 50.804$, $RMSE = 7.127$, $R^2 = 0.0$



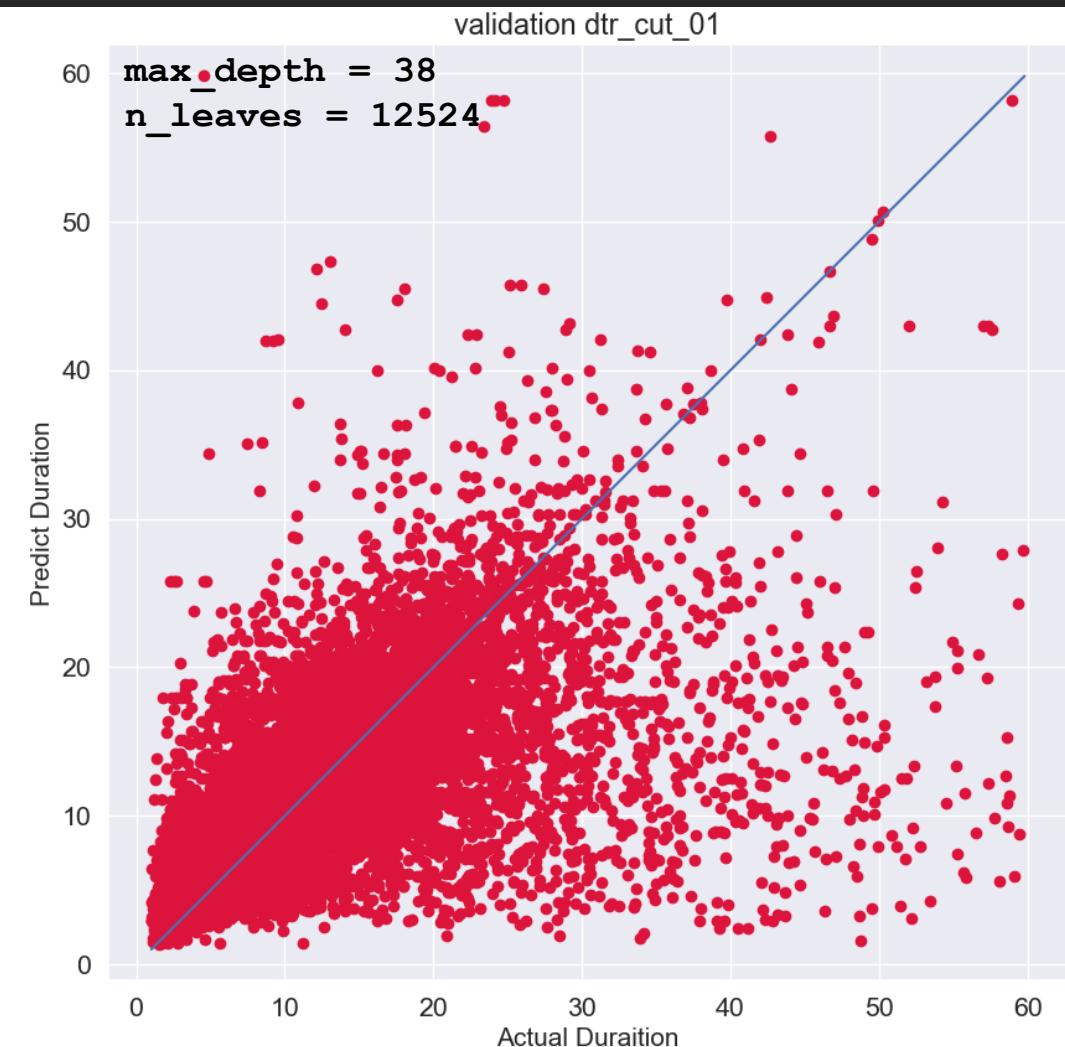
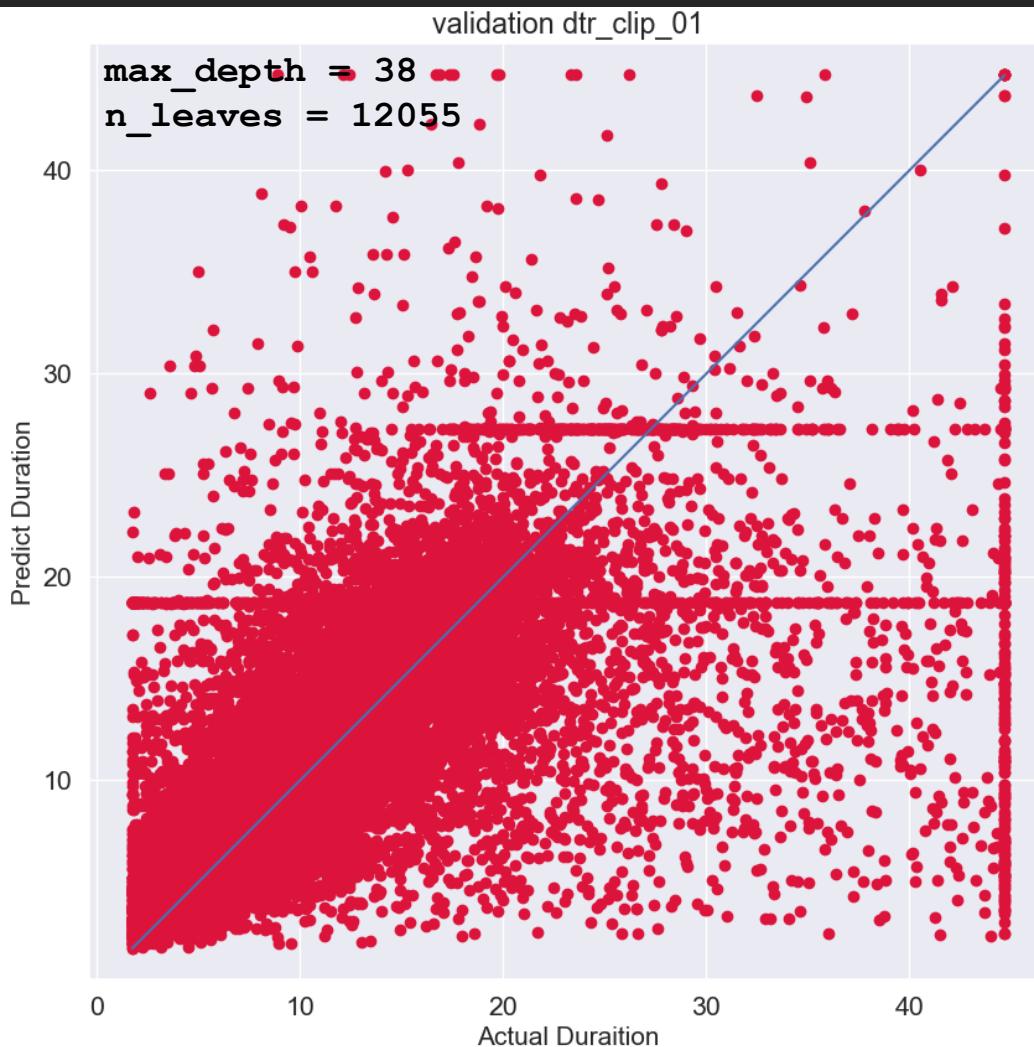
dtr_01

FEATURES:
['trip_distance']

HP:
max_depth = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0

Train: MAE = 2.774, MSE = 26.087, RMSE = 5.107, R^2 = 0.539
Validation: MAE = 3.179, MSE = 33.634, RMSE = 7.528, R^2 = 0.406

Train: MAE = 2.405, MSE = 20.088, RMSE = 4.482, R^2 = 0.597
Validation: MAE = 2.844, MSE = 27.532, RMSE = 5.247, R^2 = 0.457



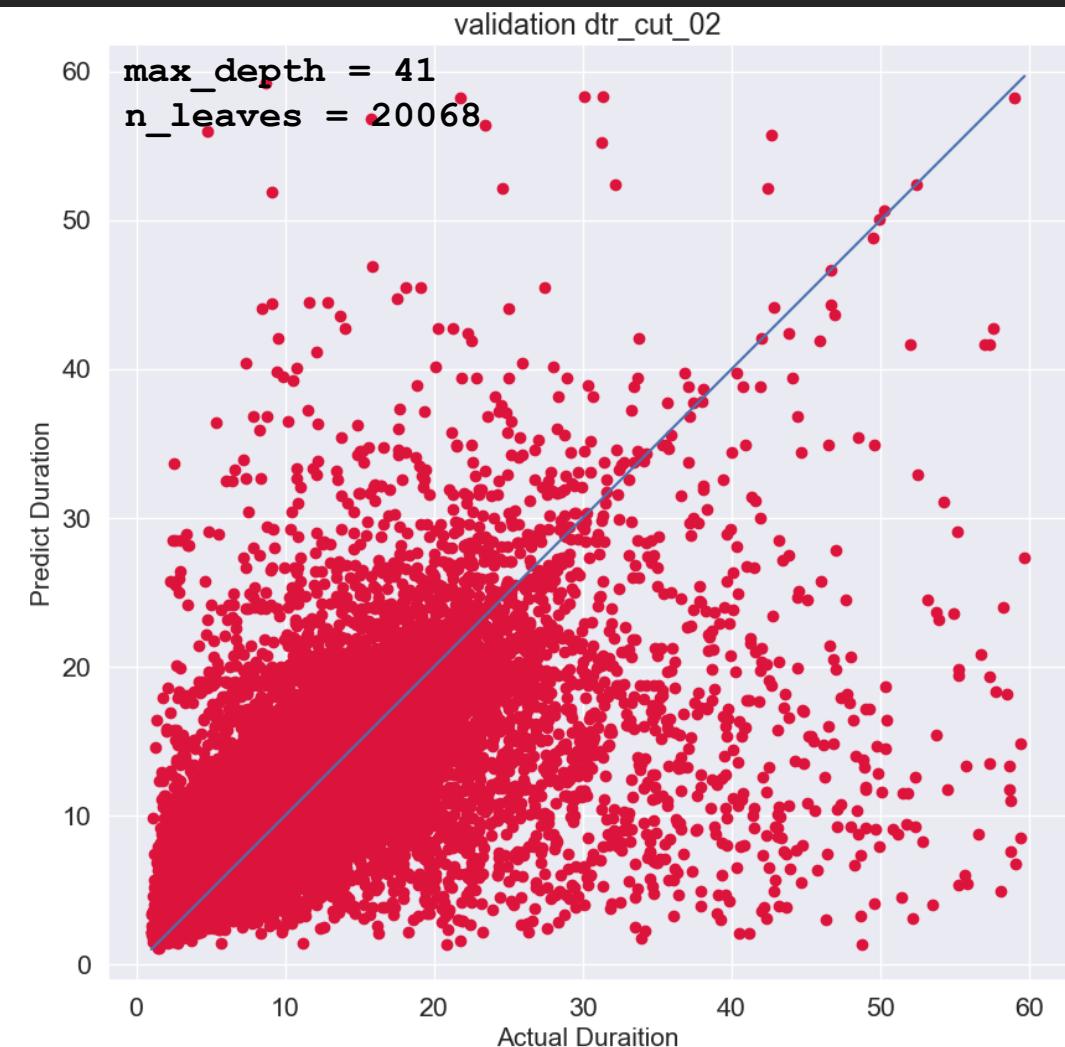
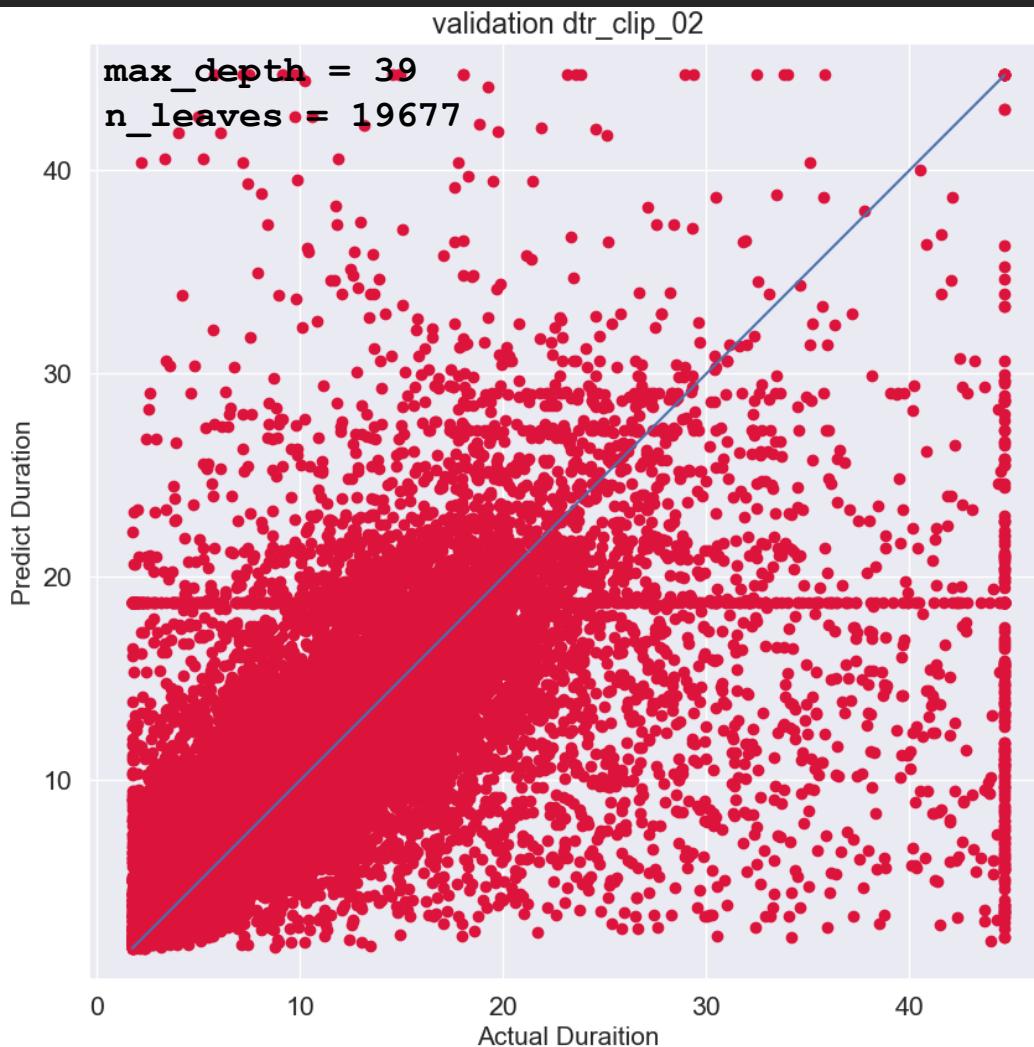
dtr_02

FEATURES:
['trip_distance', 'diff_alt']

HP:
max_dep = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0

Train: MAE = 2.517, MSE = 23.461, RMSE = 4.843, R^2 = 0.585
Validation: MAE = 3.205, MSE = 35.314, RMSE = 5.942, R^2 = 0.376

Train: MAE = 2.177, MSE = 17.583, RMSE = 4.193, R^2 = 0.642
Validation: MAE = 2.901, MSE = 29.413, RMSE = 5.423, R^2 = 0.421



dtr_03

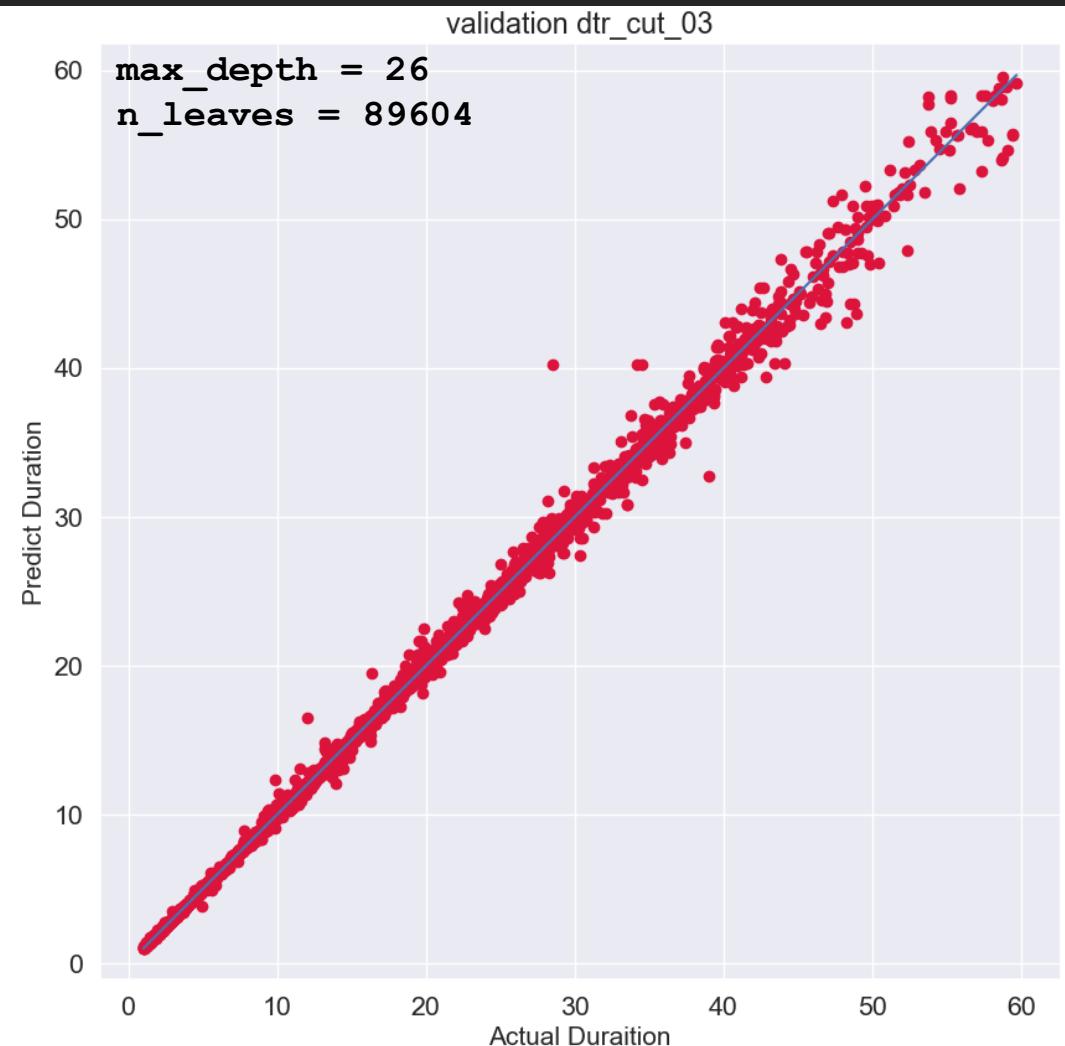
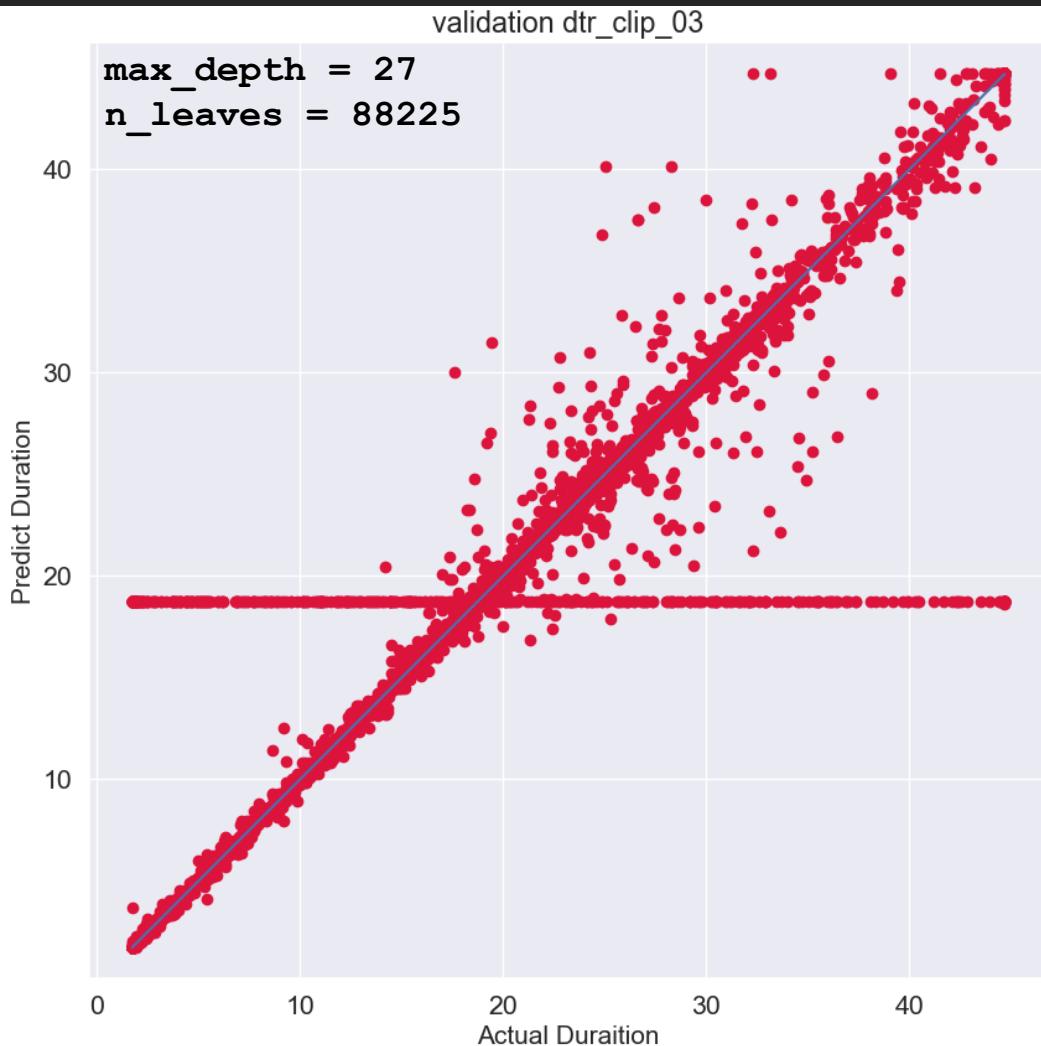
FEATURES:

['trip_distance', 'diff_alt', 'avg_trip_speed']

HP:
max_dep = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0

Train: MAE = 0.257, MSE = 4.412, RMSE = 2.101, R^2 = 0.922
Validation: MAE = 0.325, MSE = 4.223, RMSE = 2.055, R^2 = 0.925

Train: MAE = 1.387, MSE = 1.495, RMSE = 1.222, R^2 = 1
Validation: MAE = 0.071, MSE = 0.061, RMSE = 0.248, R^2 = 0.998



dtr_04

FEATURES:

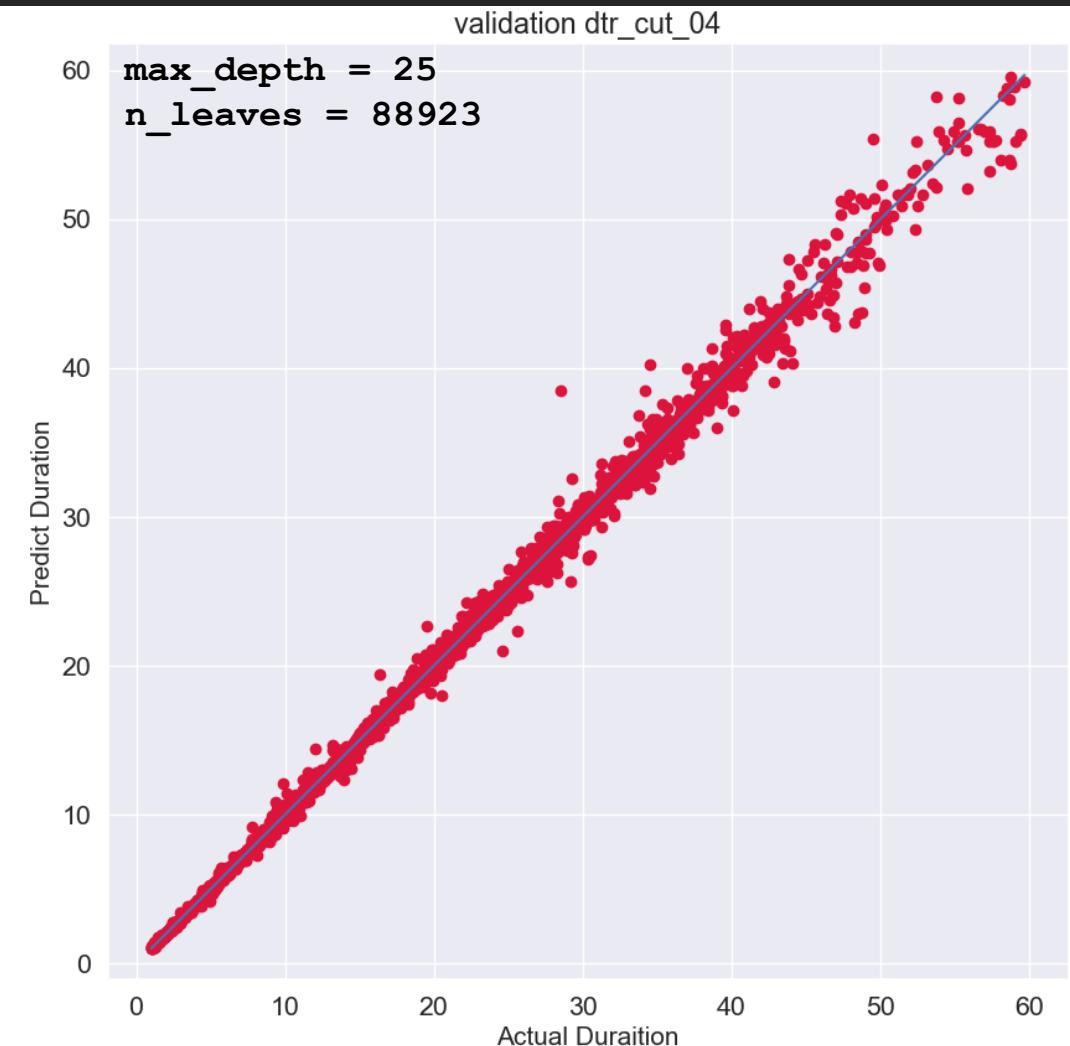
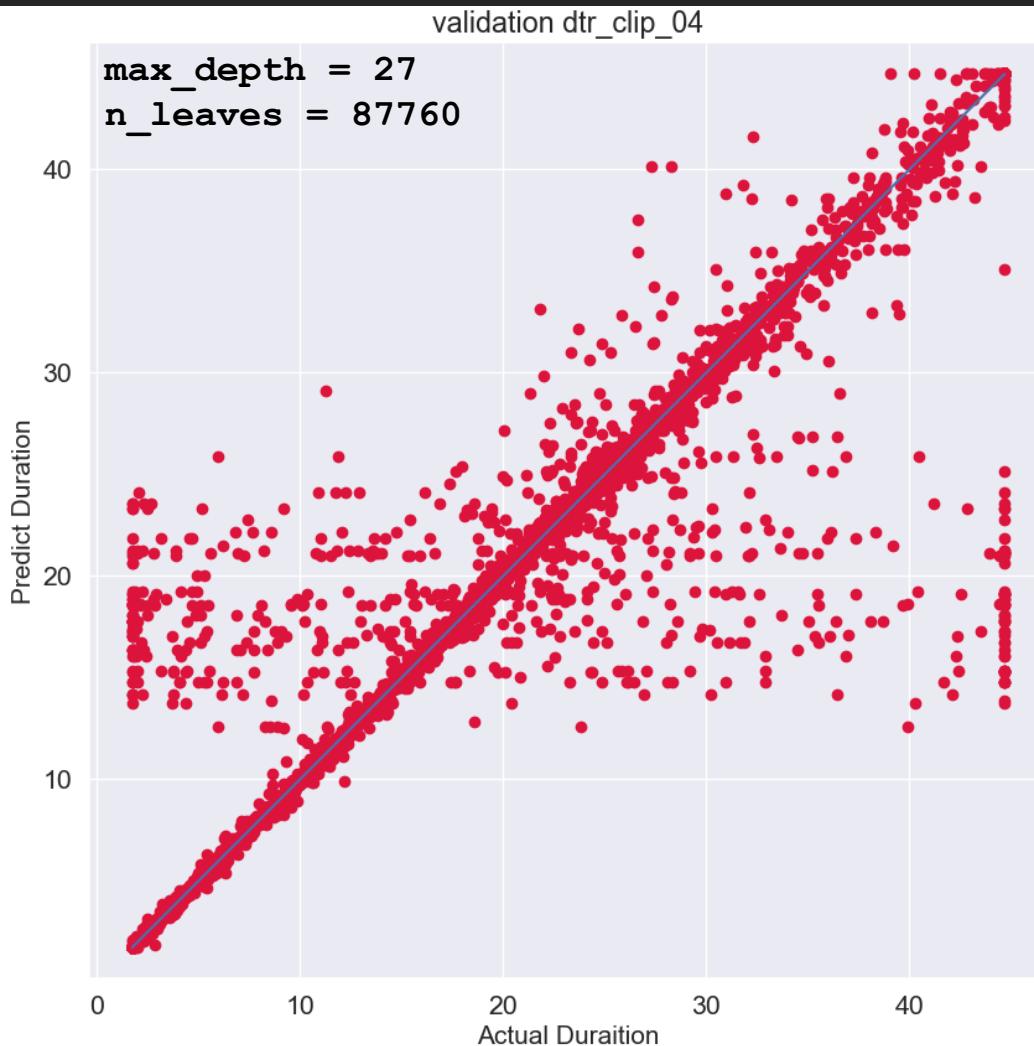
['trip_distance', 'diff_alt', 'avg_trip_speed', 'age']

HP:

max_dep = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0

Train: MAE = 0.248, MSE = 4.219, RMSE = 2.054, $R^2 = 0.925$
Validation: MAE = 0.328, MSE = 4.251, RMSE = 2.061, $R^2 = 0.924$

Train: MAE = 1.517, MSE = 1.691, RMSE = 1.301, $R^2 = 1$
Validation: MAE = 0.075, MSE = 0.065, RMSE = 0.255, $R^2 = 0.998$



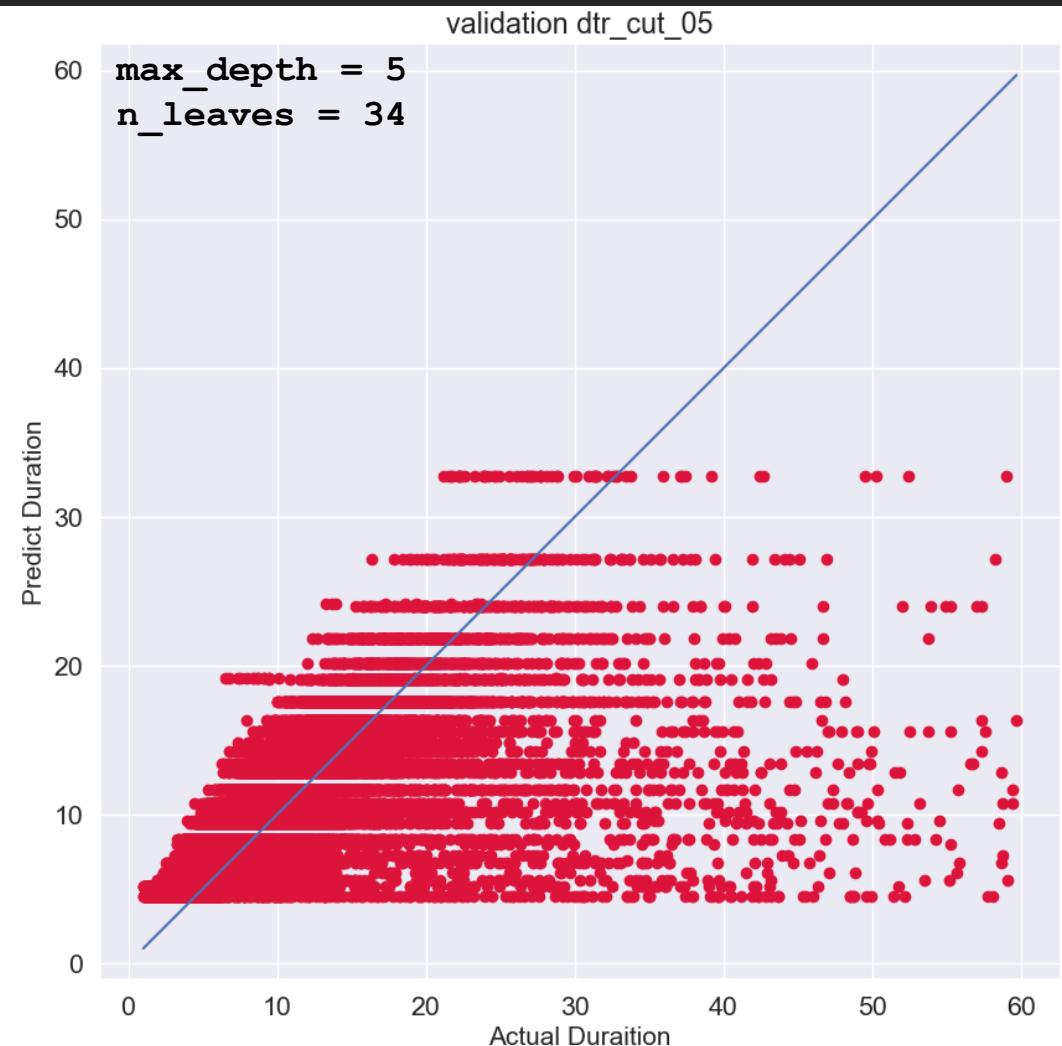
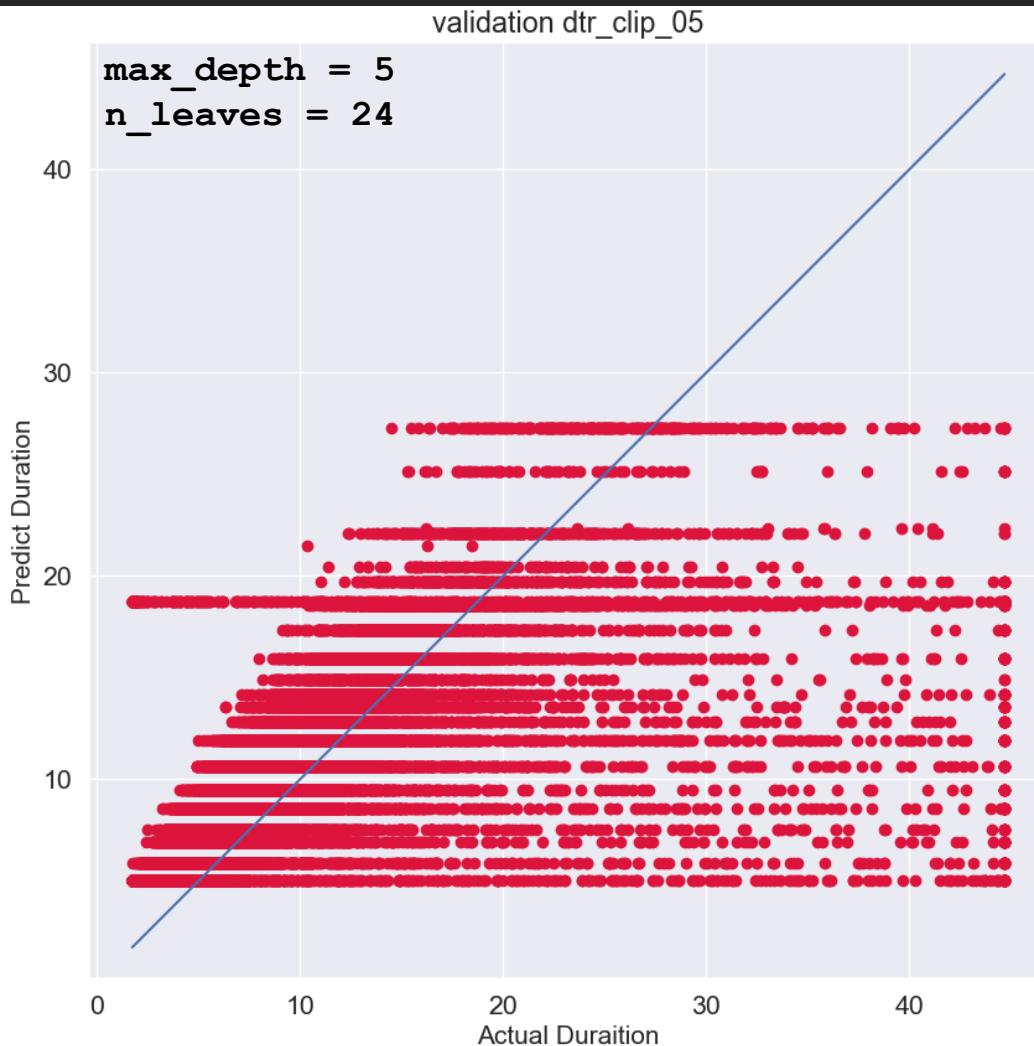
dtr_05

FEATURES:
['trip_distance']

HP: max_dep = 5
min_samp_split = 5
min_samp_leaf = 5
ccp = 0.0

Train: MAE = 3.322, MSE = 34.671, RMSE = 5.888, $R^2 = 0.387$
Validation: MAE = 3.306, MSE = 34.666, RMSE = 5.887, $R^2 = 0.388$

Train: MAE = 2.978, MSE = 28.166, RMSE = 5.307, $R^2 = 0.427$
Validation: MAE = 3.007, MSE = 28.867, RMSE = 5.372, $R^2 = 0.431$



dtr_06

FEATURES:

['trip_distance', 'diff_alt_slope_-1', 'diff_alt_slope_0', 'diff_alt_slope_1',
'user_type_Customer', 'user_type_Subscriber', 'member_gender_Female',
'member_gender_Male', 'bike_share_for_all_trip_No', 'bike_share_for_all_trip_Yes']

HP:

max_dep = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0

Train: MAE = 1.873,
Validation: MAE = 3.168,

MSE = 15.763,
MSE = 37.726,

RMSE = 3.971,
RMSE = 6.142,

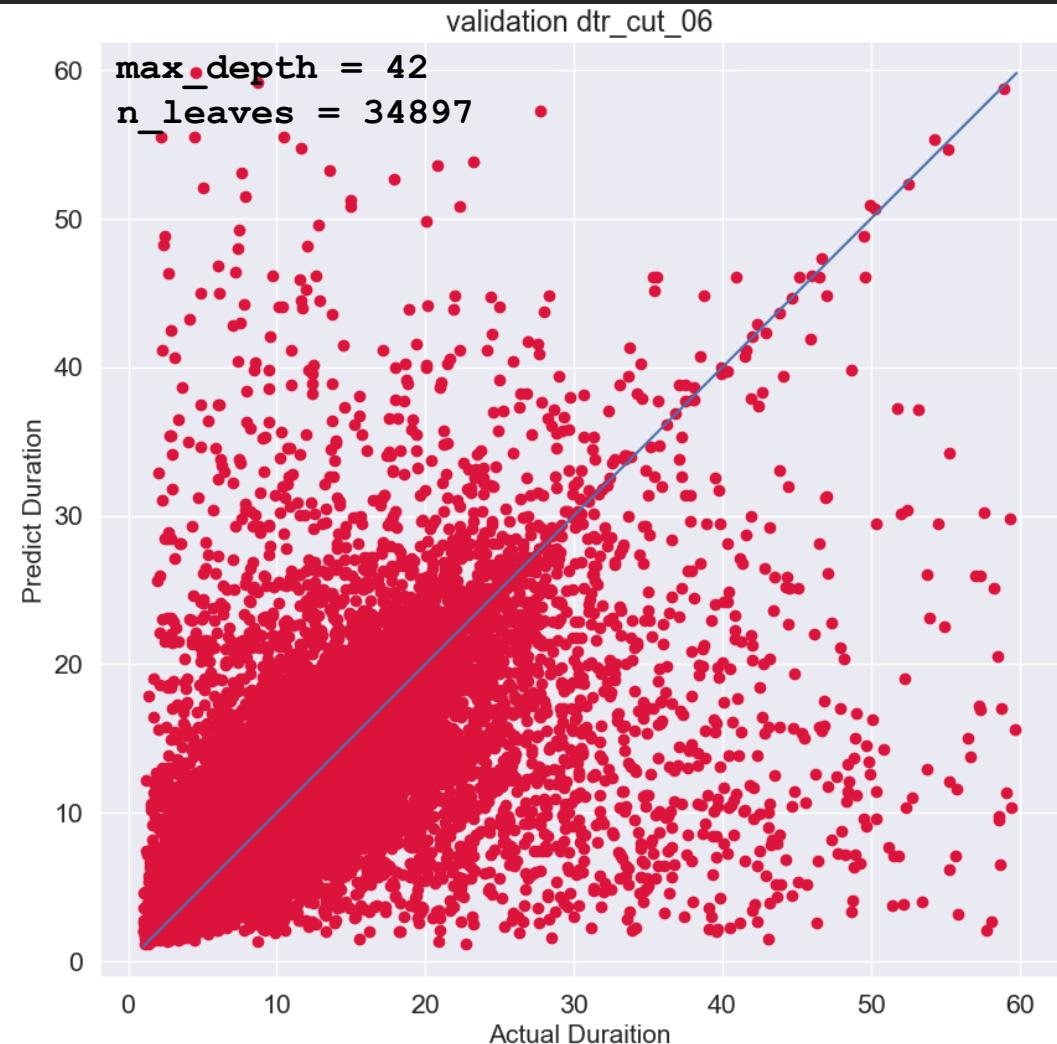
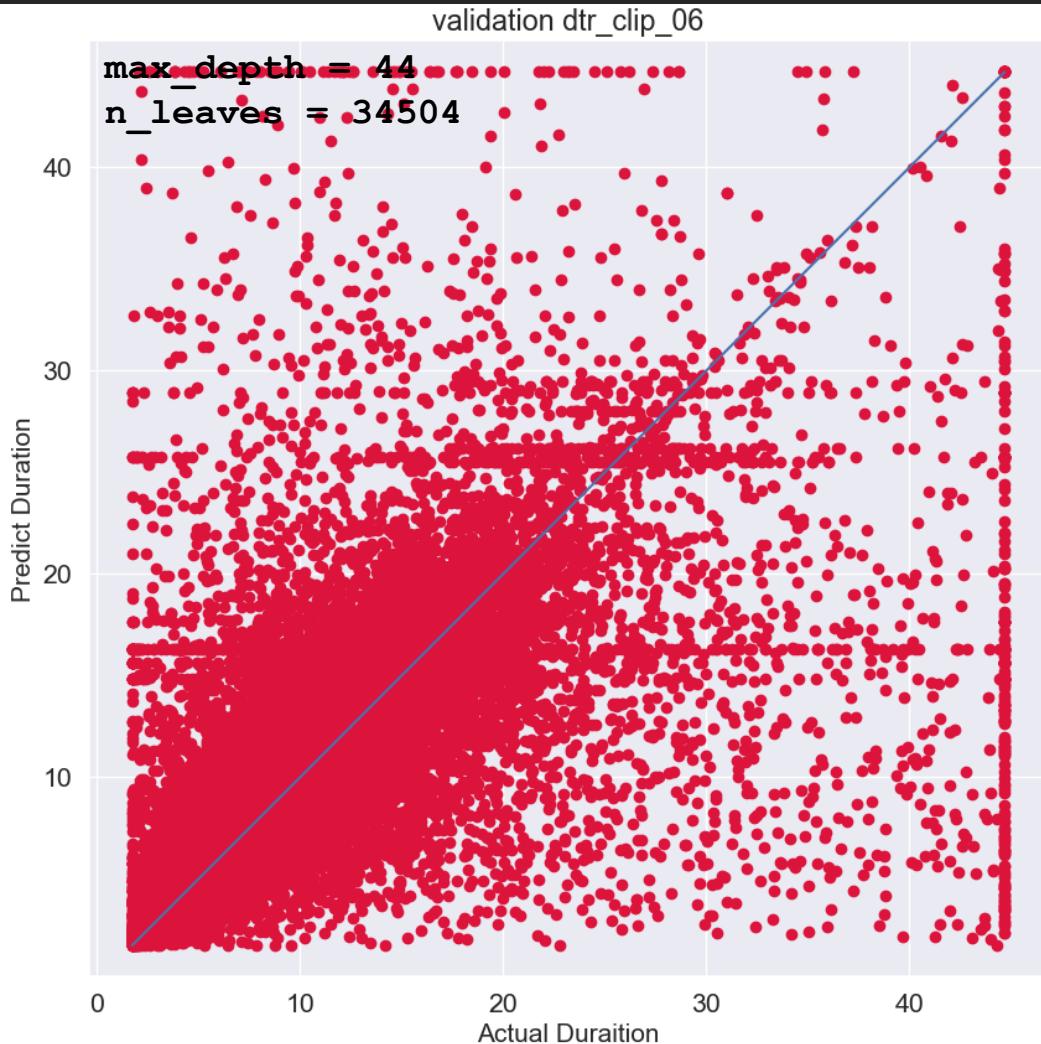
R^2 = 0.721
 R^2 = 0.334

Train: MAE = 1.584,
Validation: MAE = 2.902,

MSE = 10.907,
MSE = 32.634,

RMSE = 3.302,
RMSE = 5.712,

R^2 = 0.778
 R^2 = 0.356



dtr_07

FEATURES:

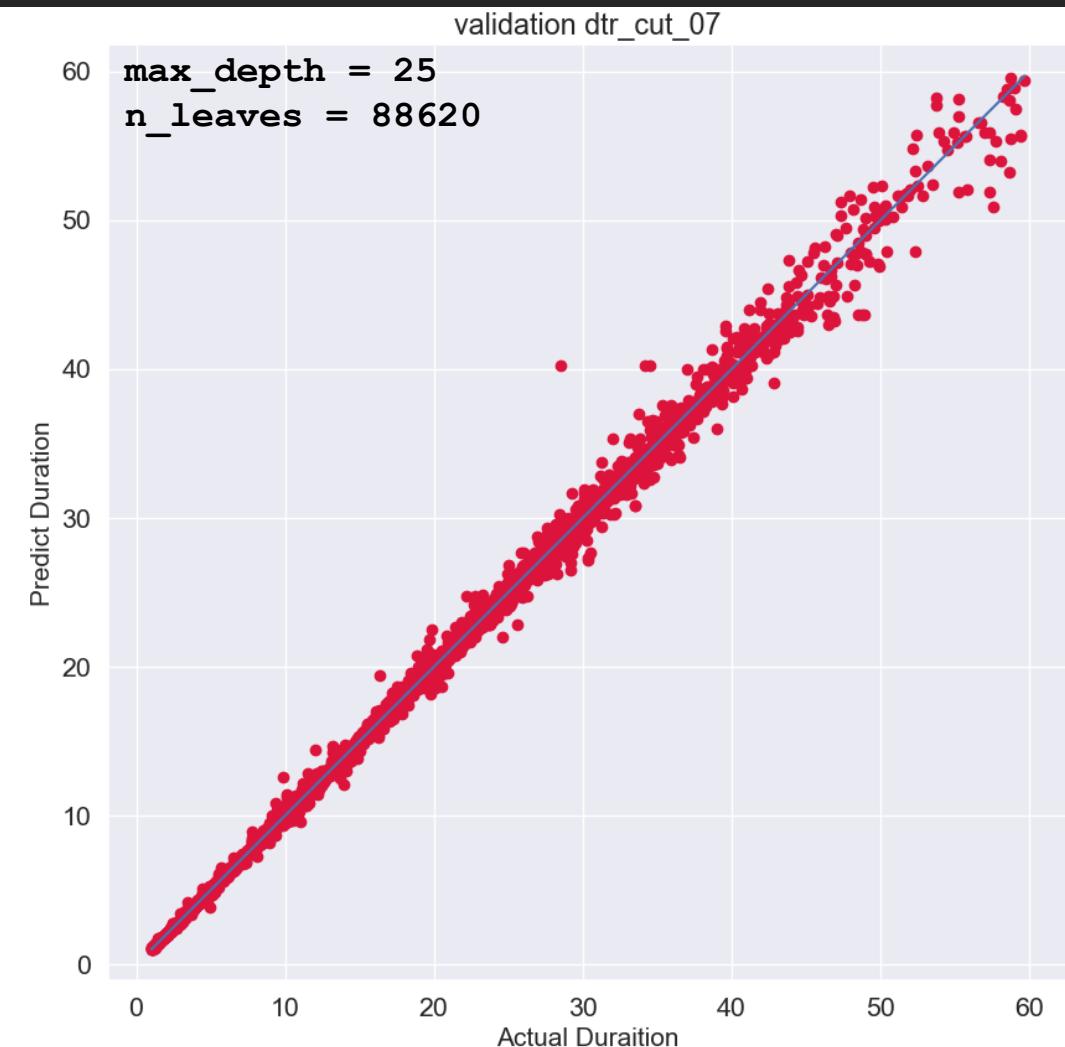
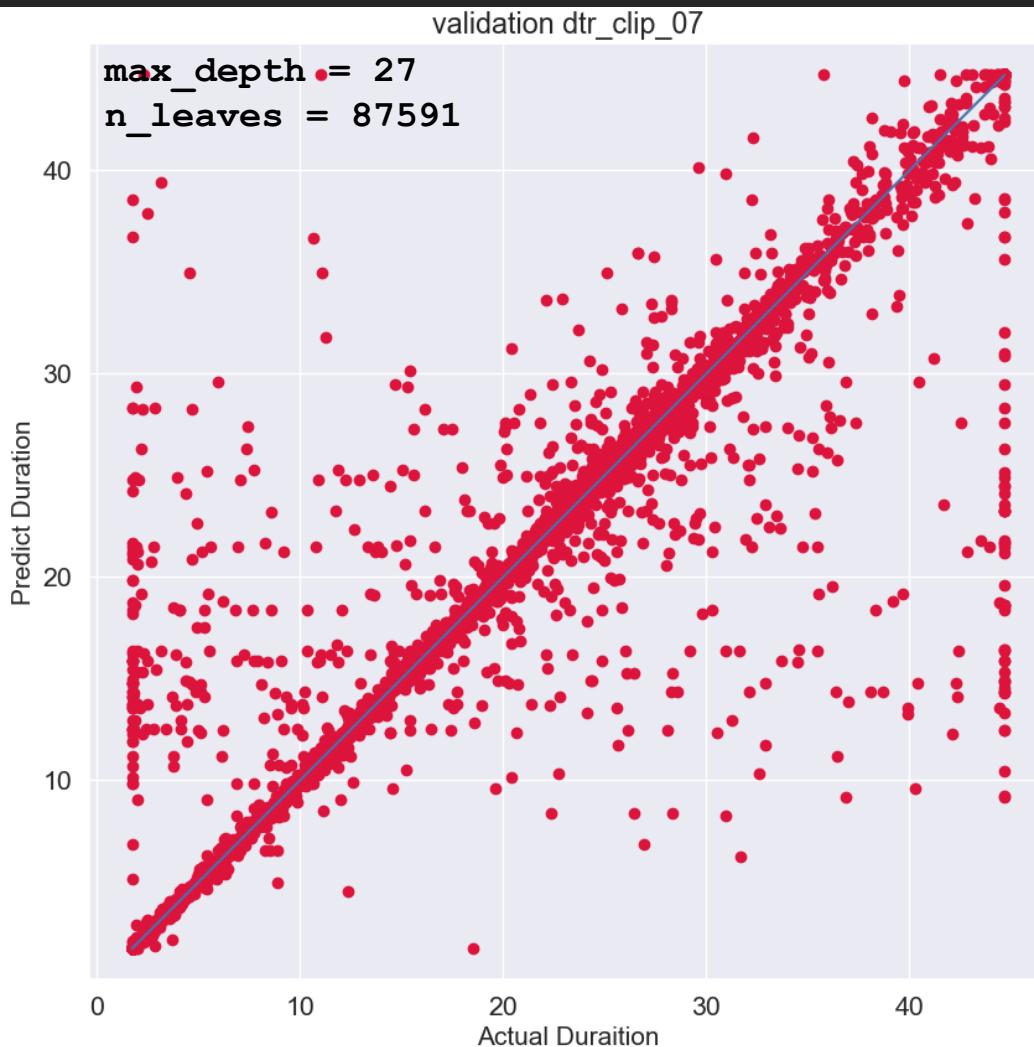
['trip_distance', 'diff_alt', 'avg_trip_speed', 'age', 'diff_alt_slope_-1', 'diff_alt_slope_0',
'diff_alt_slope_1', 'user_type_Customer', 'user_type_Subscriber', 'member_gender_Female',
'member_gender_Male', 'bike_share_for_all_trip_No', 'bike_share_for_all_trip_Yes']

HP:

max_dep = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0

Train: MAE = 0.211, MSE = 3.385, RMSE = 1.839, $R^2 = 0.941$
Validation: MAE = 0.318, MSE = 4.253, RMSE = 2.062, $R^2 = 0.924$

Train: MAE = 1.546, MSE = 1.771, RMSE = 1.331, $R^2 = 1$
Validation: MAE = 0.078, MSE = 0.0711, RMSE = 0.266, $R^2 = 0.998$



Models Comparison

Clipping_df

lr/dtr	#	tag	part	MAE	MSE	RMSE	R2
dtr	00	clip	train	5.341	56.623	7.524	0
dtr	00	clip	valid	5.358	56.675	7.528	0
dtr	01	clip	train	2.744	26.087	5.107	0.539
dtr	01	clip	valid	3.179	33.634	5.799	0.406
dtr	02	clip	train	2.517	23.461	4.843	0.585
dtr	02	clip	valid	3.205	35.314	5.942	0.376
dtr	03	clip	train	0.257	4.412	2.101	0.922
dtr	03	clip	valid	0.325	4.223	2.055	0.925
dtr	04	clip	train	0.248	4.219	2.054	0.925
dtr	04	clip	valid	0.328	4.251	2.061	0.924
dtr	05	clip	train	3.322	34.671	5.888	0.387
dtr	05	clip	valid	3.306	34.666	5.887	0.388
dtr	06	clip	train	1.873	15.763	3.971	0.721
dtr	06	clip	valid	3.168	37.726	6.142	0.334
dtr	07	clip	train	0.211	3.385	1.839	0.941
dtr	07	clip	valid	0.318	4.253	2.062	0.924

Cutting_df

lr/dtr	#	tag	part	MAE	MSE	RMSE	R2
dtr	00	cut	train	5.026	49.197	7.014	0
dtr	00	cut	valid	5.084	50.731	7.122	0
dtr	01	cut	train	2.405	20.088	4.482	0.597
dtr	01	cut	valid	2.844	27.532	5.247	0.457
dtr	02	cut	train	2.177	17.583	4.193	0.642
dtr	02	cut	valid	2.901	29.413	5.423	0.421
dtr	03	cut	train	1.387	1.495	1.222	1
dtr	03	cut	valid	0.071	0.061	0.248	0.998
dtr	04	cut	train	1.517	1.691	1.301	1
dtr	04	cut	valid	0.075	0.065	0.255	0.998
dtr	05	cut	train	2.978	28.166	5.307	0.427
dtr	05	cut	valid	3.007	28.867	5.372	0.431
dtr	06	cut	train	1.584	10.907	3.302	0.778
dtr	06	cut	valid	2.902	32.634	5.712	0.356
dtr	07	cut	train	1.546	1.771	1.331	1
dtr	07	cut	valid	0.078	0.071	0.266	0.998

dtr_test

FEATURES:

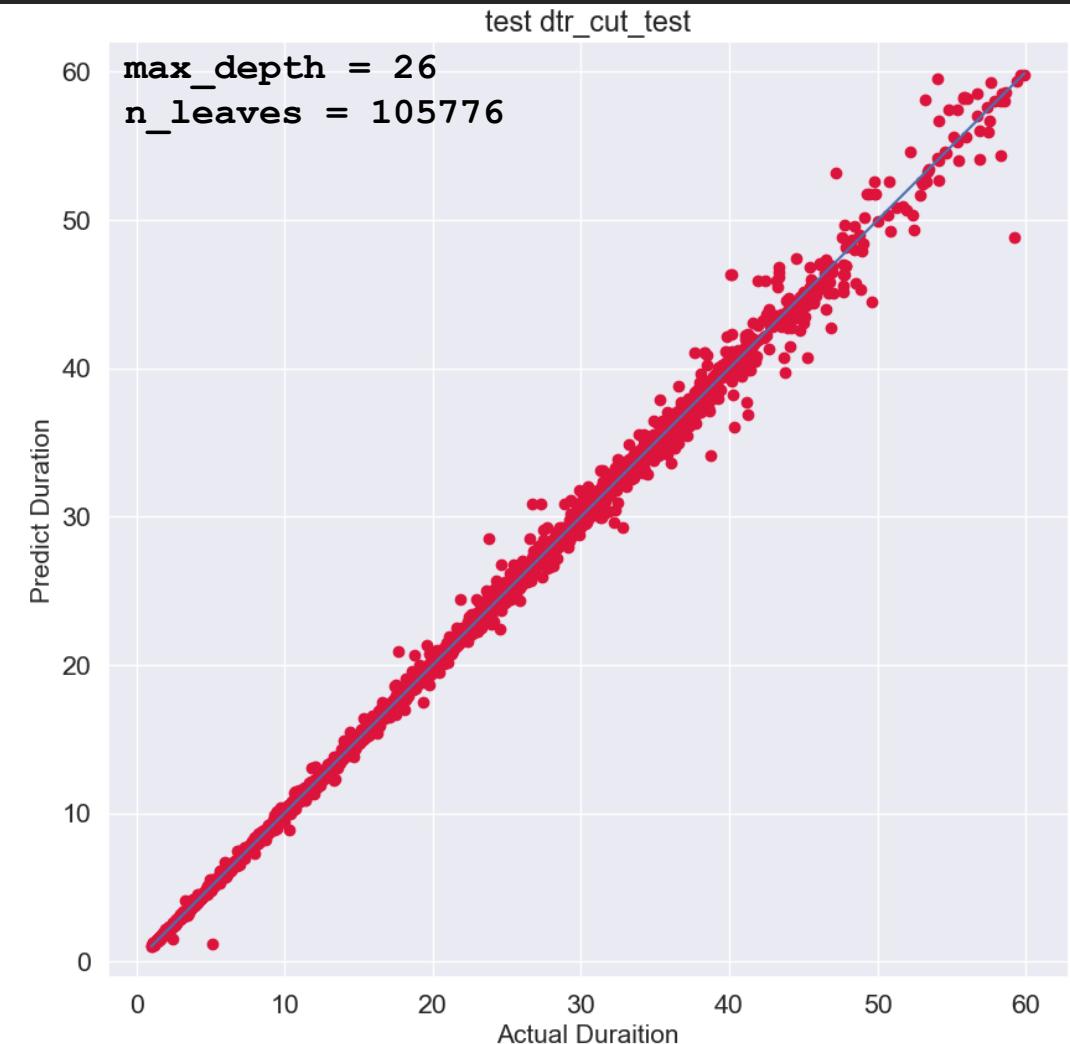
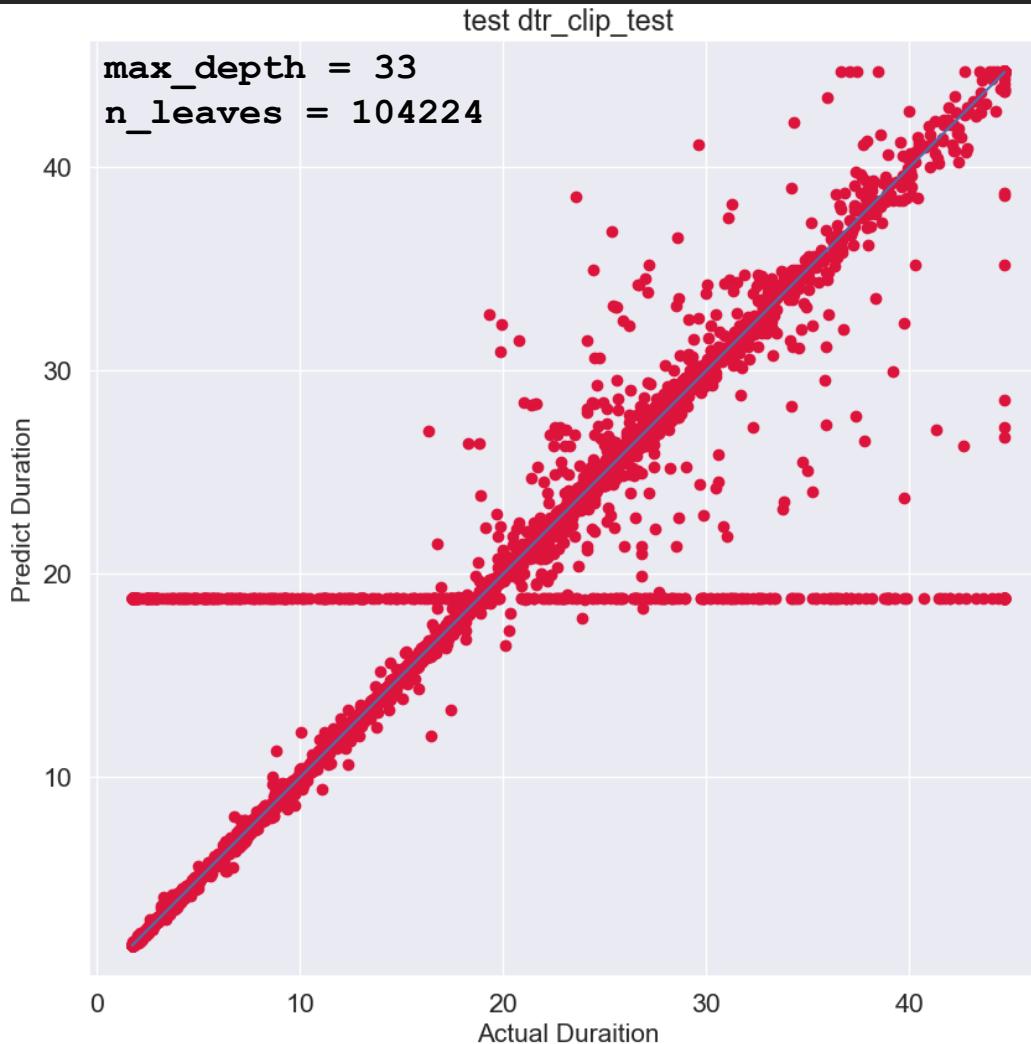
['trip_distance', 'diff_alt', 'avg_trip_speed']

HP:

max_dep = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0

Train: MAE = 0.253, MSE = 4.334, RMSE = 2.081, R^2 = 0.923
Test: MAE = 0.336, MSE = 4.471, RMSE = 2.114, R^2 = 0.921

Train: MAE = $1.7e^{-17}$, MSE = $1.9e^{-32}$, RMSE = $1.3e^{-16}$, R^2 = 1
Test: MAE = 0.064, MSE = 0.056, RMSE = 0.237, R^2 = 0.998



Decision-Tree Conclusions

The best model we created through the Decision-Tree method was **lr_03** which based on the following features:
['trip_distance', 'diff_alt', `avg_trip_speed`]

Make sense ??? Business insight ???

We didn't find any business insight when we create the model based on the whole data of the user when it finish to use.
In addition, we based the model on features when those calculated by the target itself.

Genius but instructive!

So, we made another session to think how we can use the current code to get some insights.

lr/dtr	#	tag	part	MAE	MSE	RMSE	R2
dtr	test	clip	train	0.253	4.334	2.081	0.923
dtr	test	clip	test	0.336	4.471	2.114	0.921
dtr	test	cut	train	1.75E-17	1.94E-32	1.39E-16	1
dtr	test	cut	test	0.064	0.056	0.237	0.998

*The RMSE is minutes.

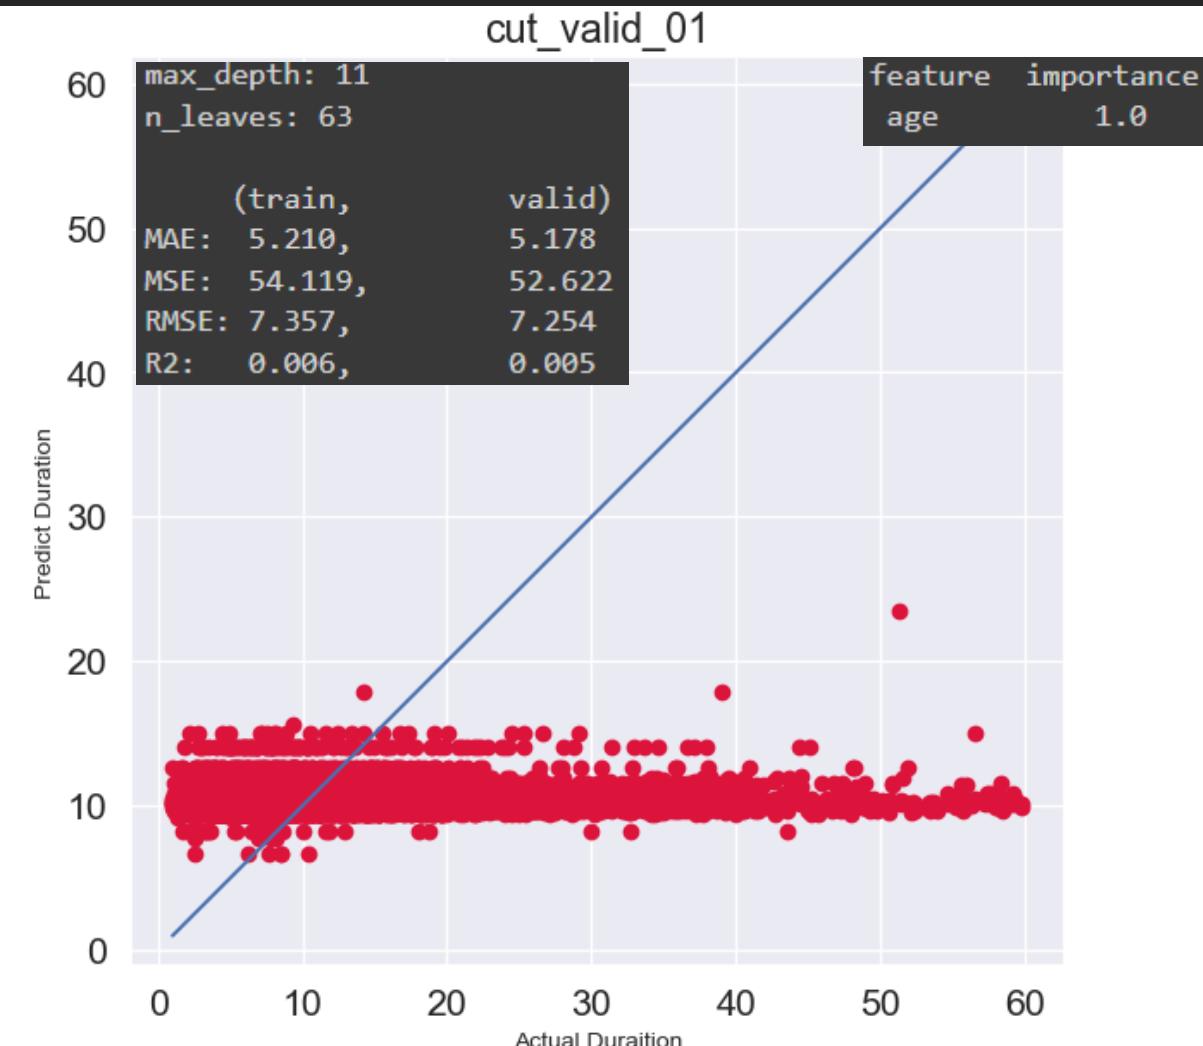
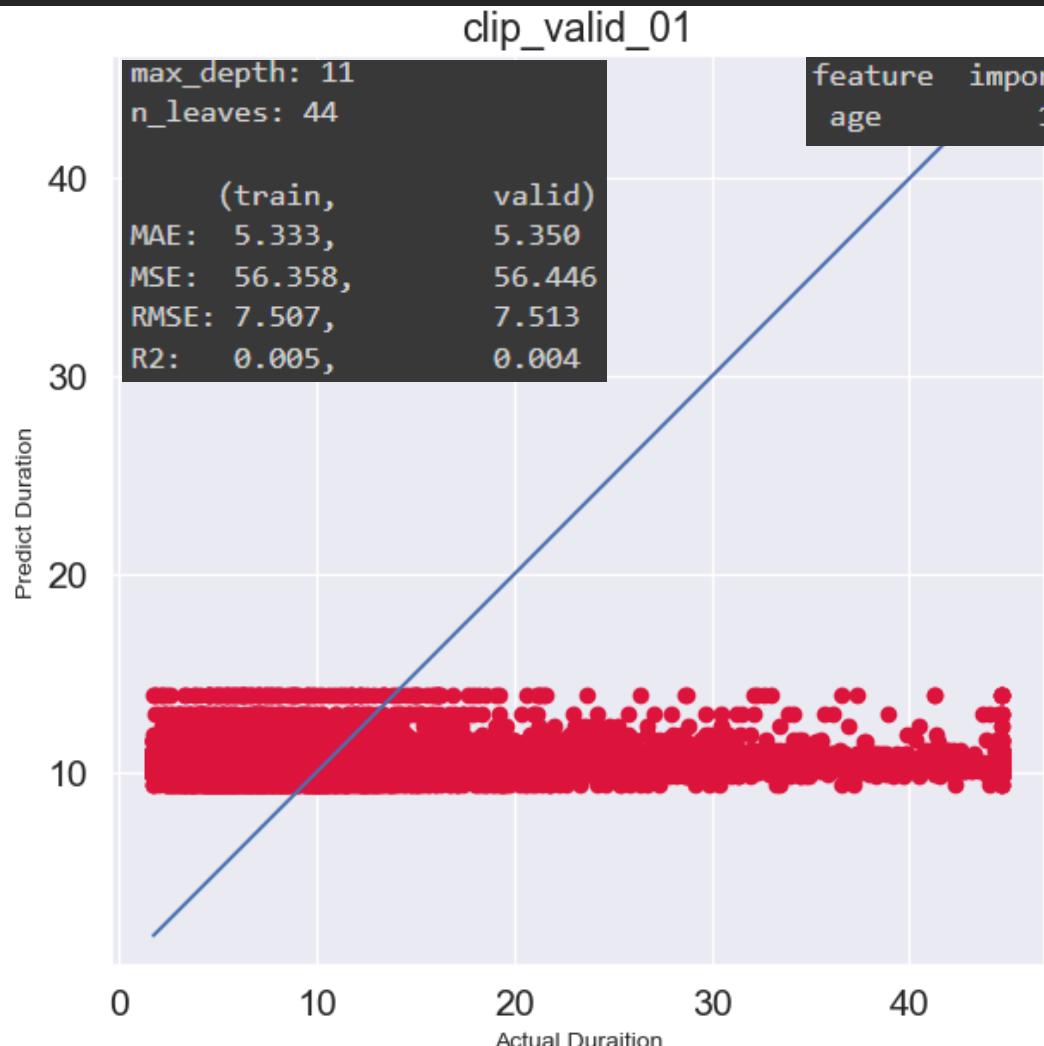
Practically, the features we can use as an estimators for a model which can contribute an insight are the features were we know about the user at the beginning of using:

- Age
- User type (Customer/ Subscriber)
- Gender (Male/ Female)
- Payment method (Direct/Indirect)

Ndtr_01

FEATURES:
['age']

HP:
max_depth = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0



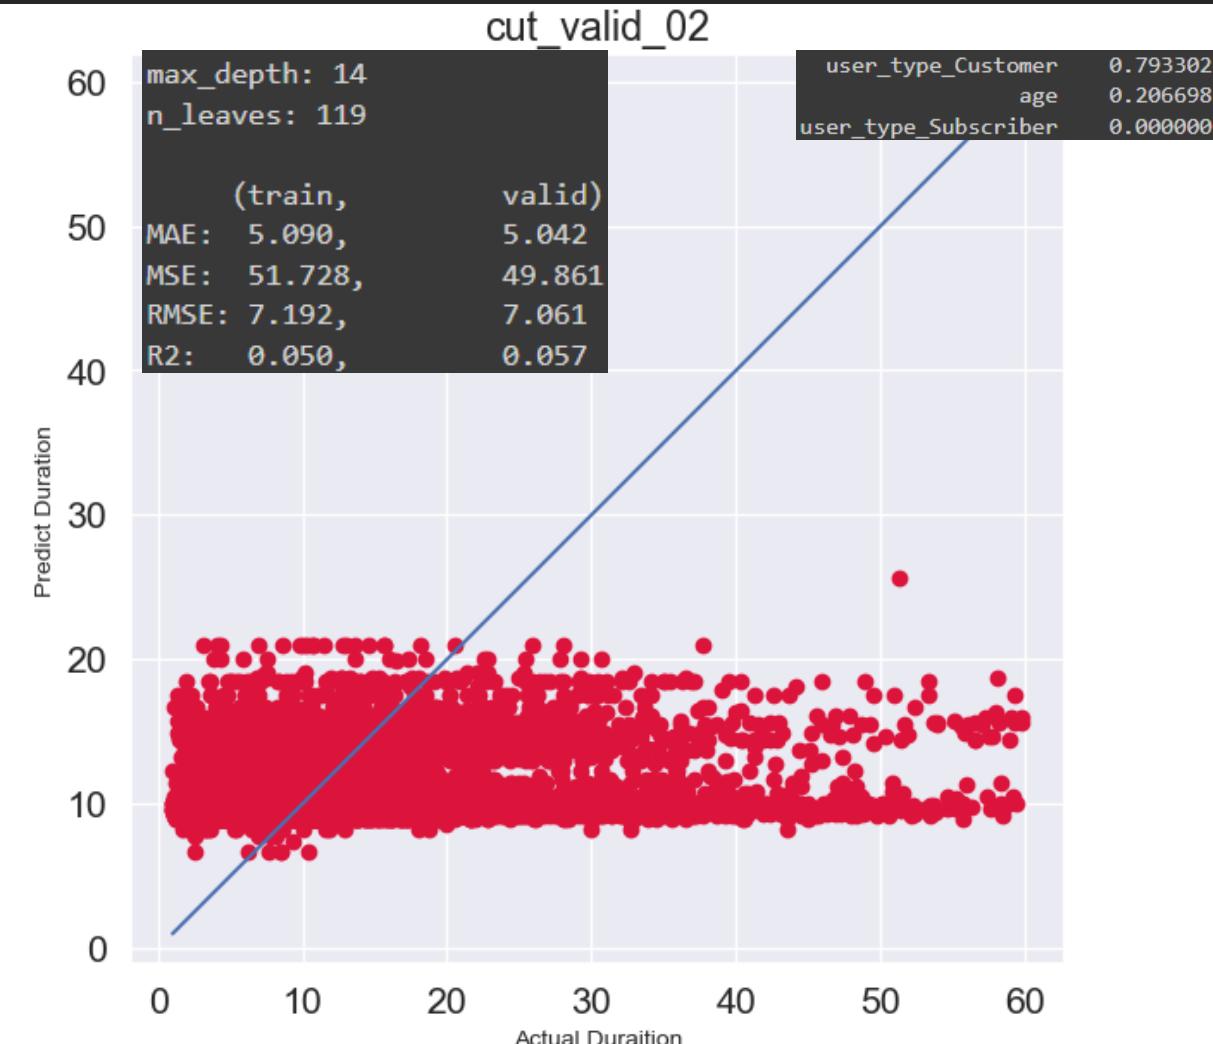
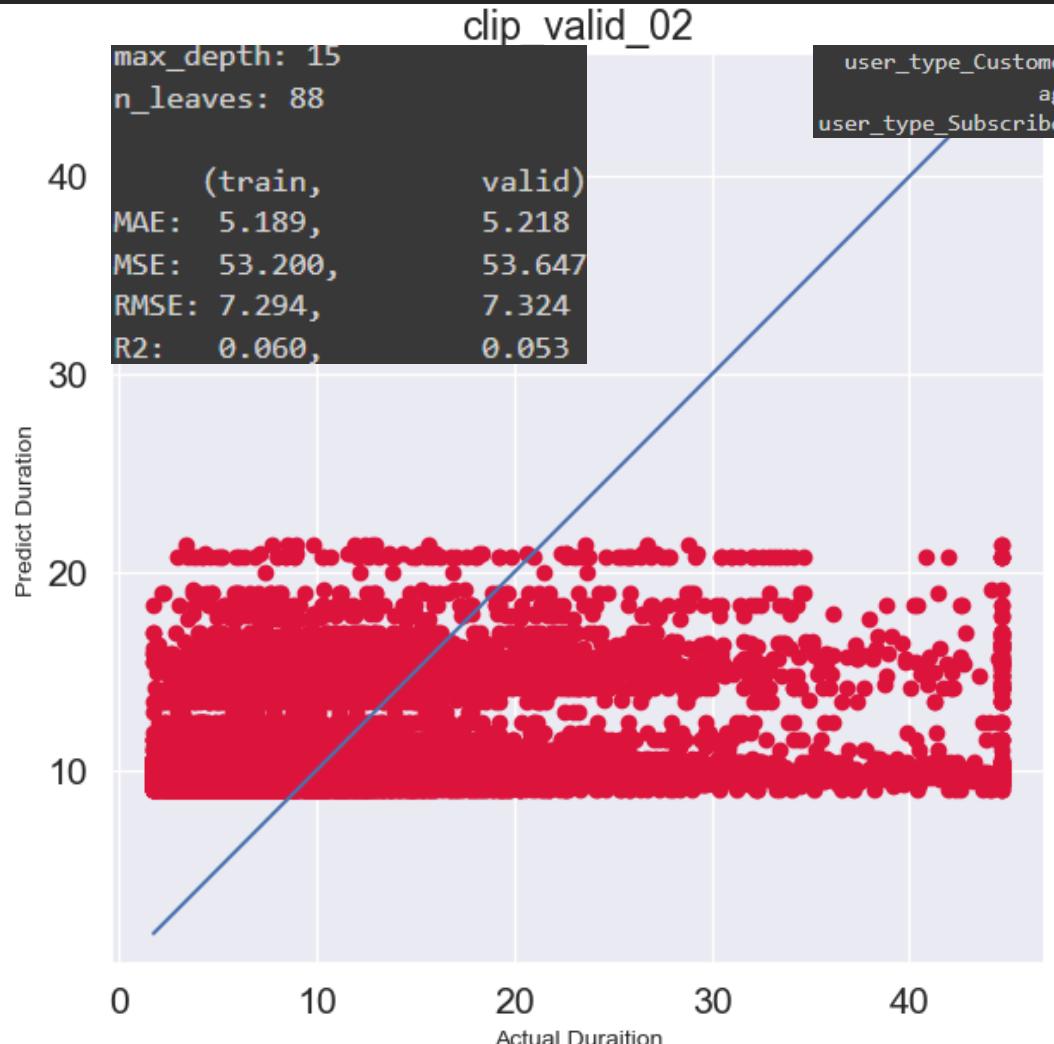
Ndtr_02

FEATURES:

['age', 'user_type_Customer', 'user_type_Subscriber']

HP:

max_depth = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0

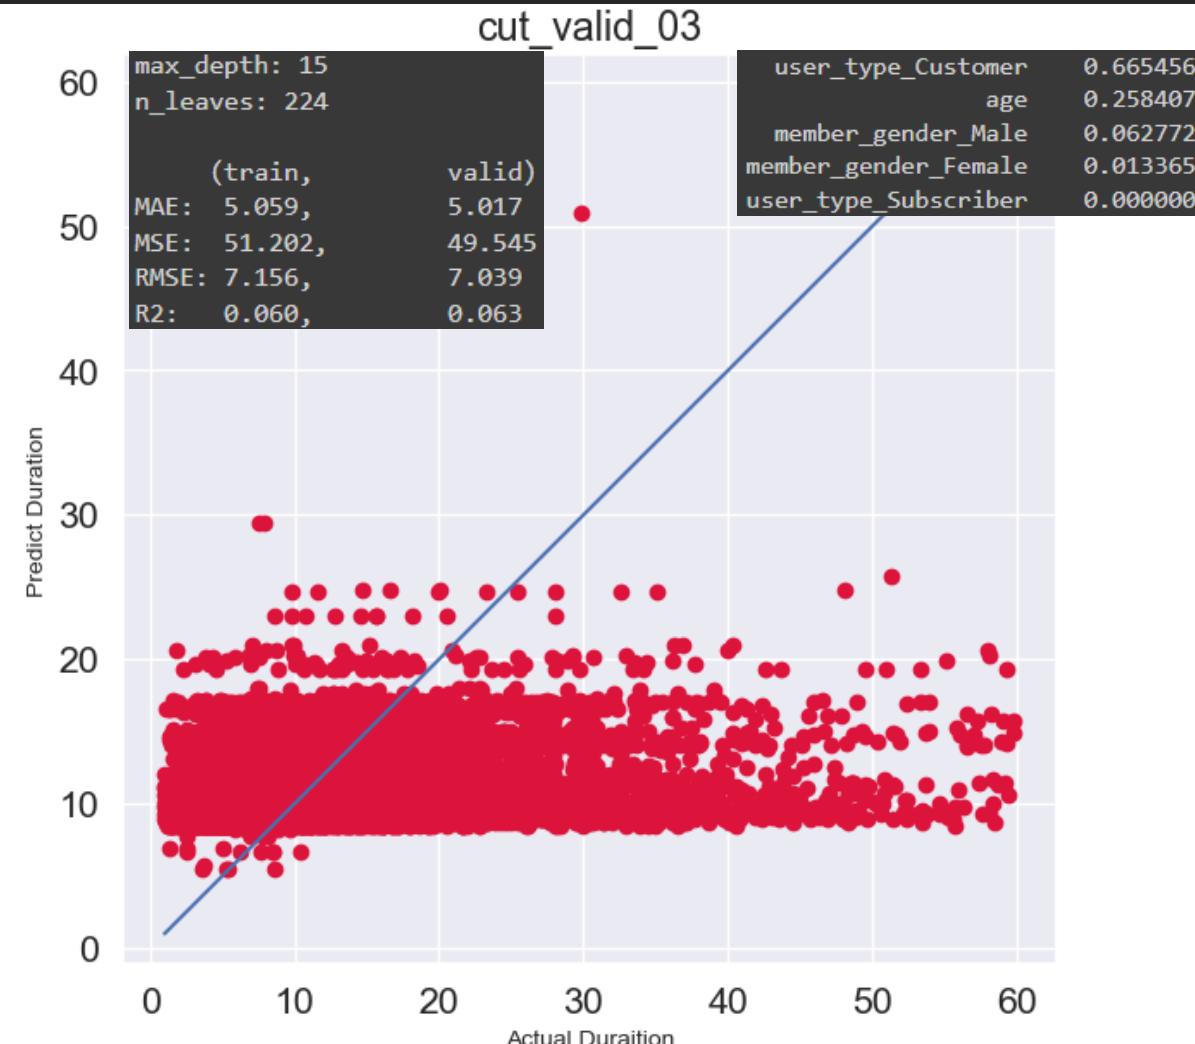
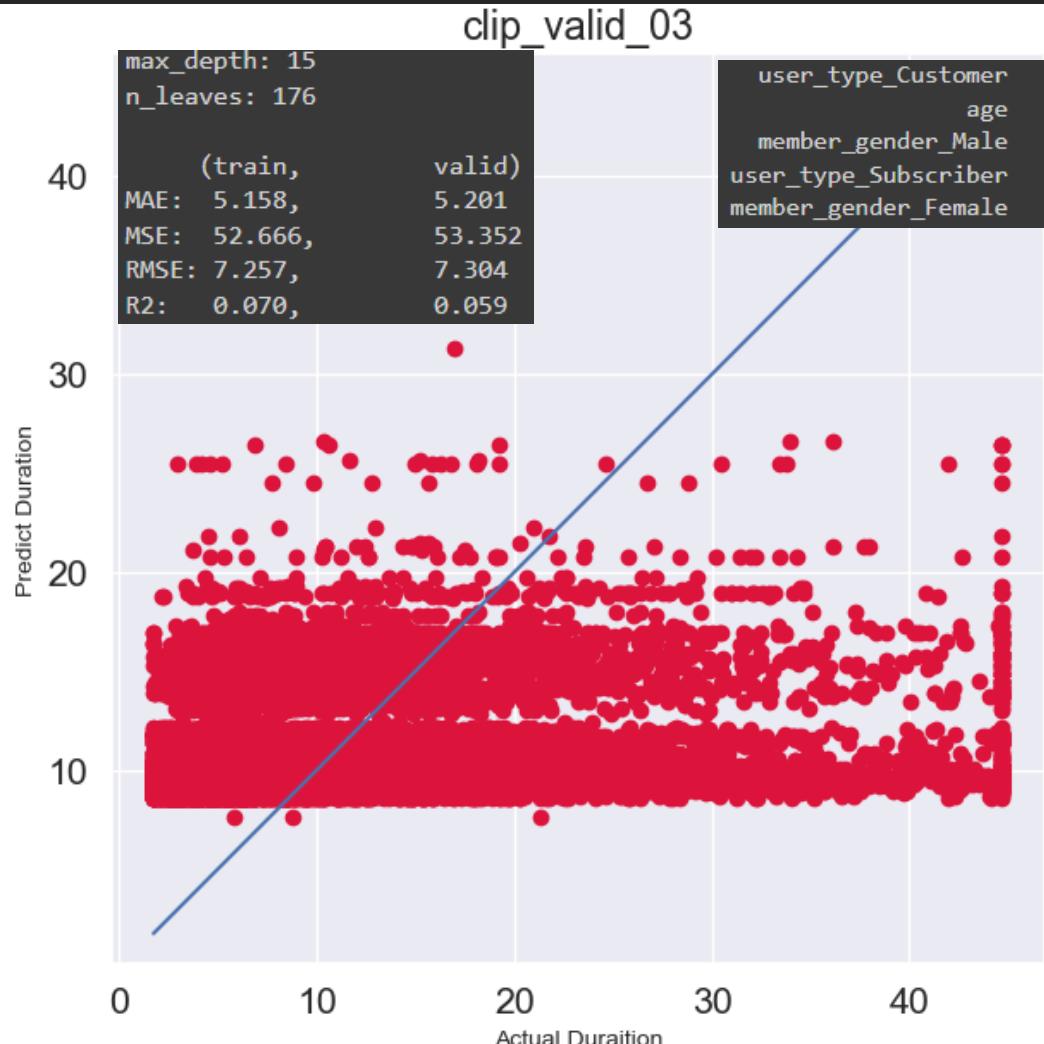


Ndtr_03

FEATURES:

['age', 'user_type_Customer', 'user_type_Subscriber',
'member_gender_Female', 'member_gender_Male']

HP:
max_depth = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0



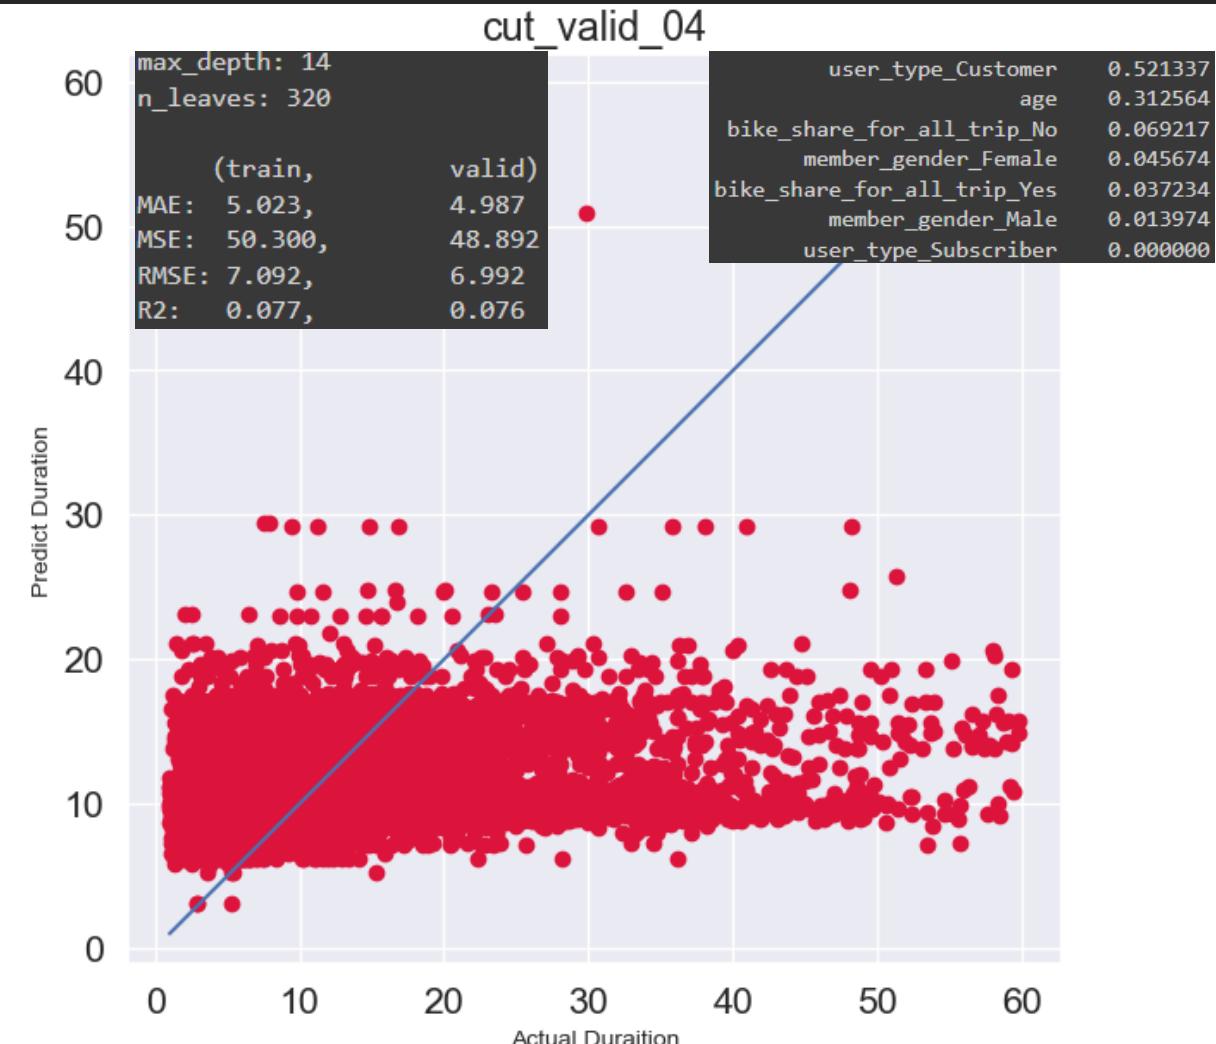
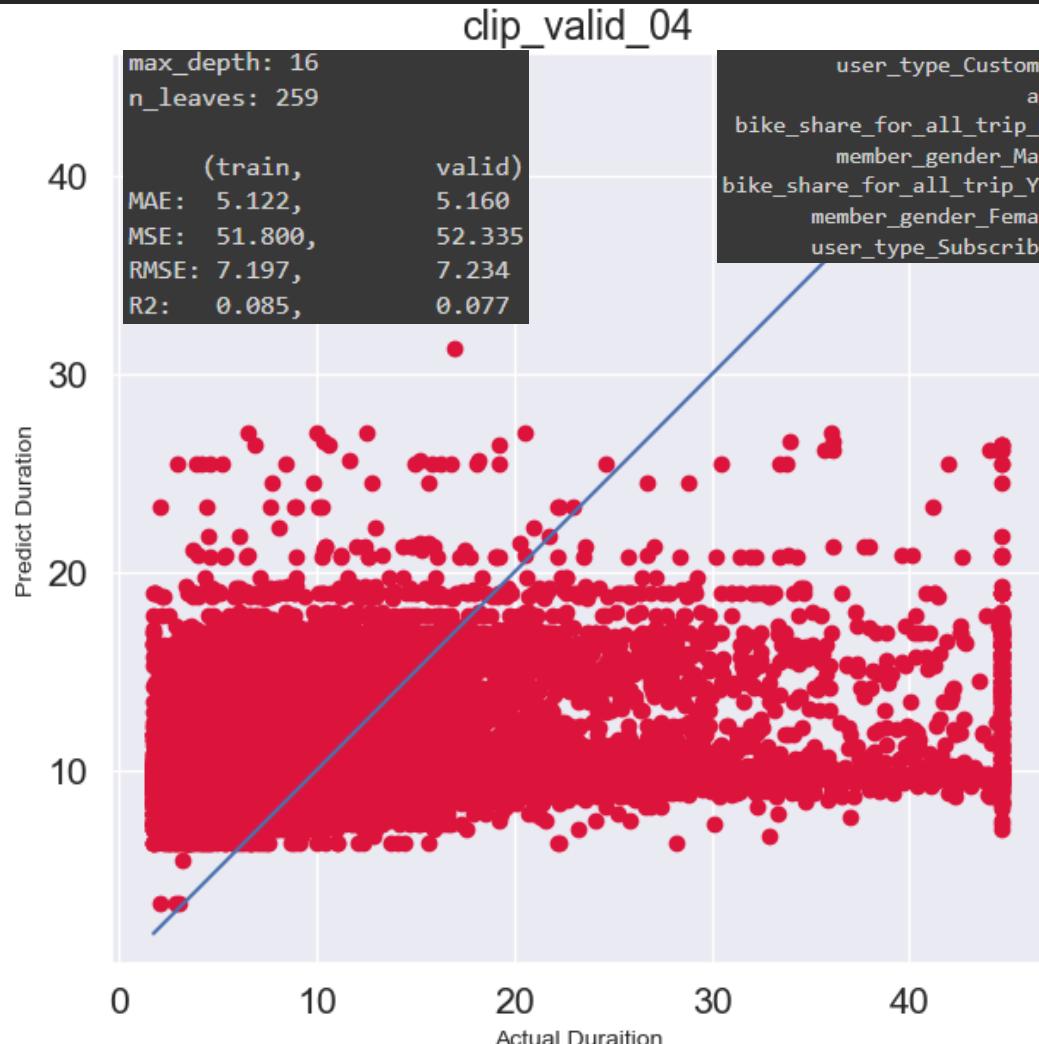
Ndtr_04

FEATURES:

['age', 'user_type_Customer', 'user_type_Subscriber',
'member_gender_Female', 'member_gender_Male']

HP:

max_depth = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0



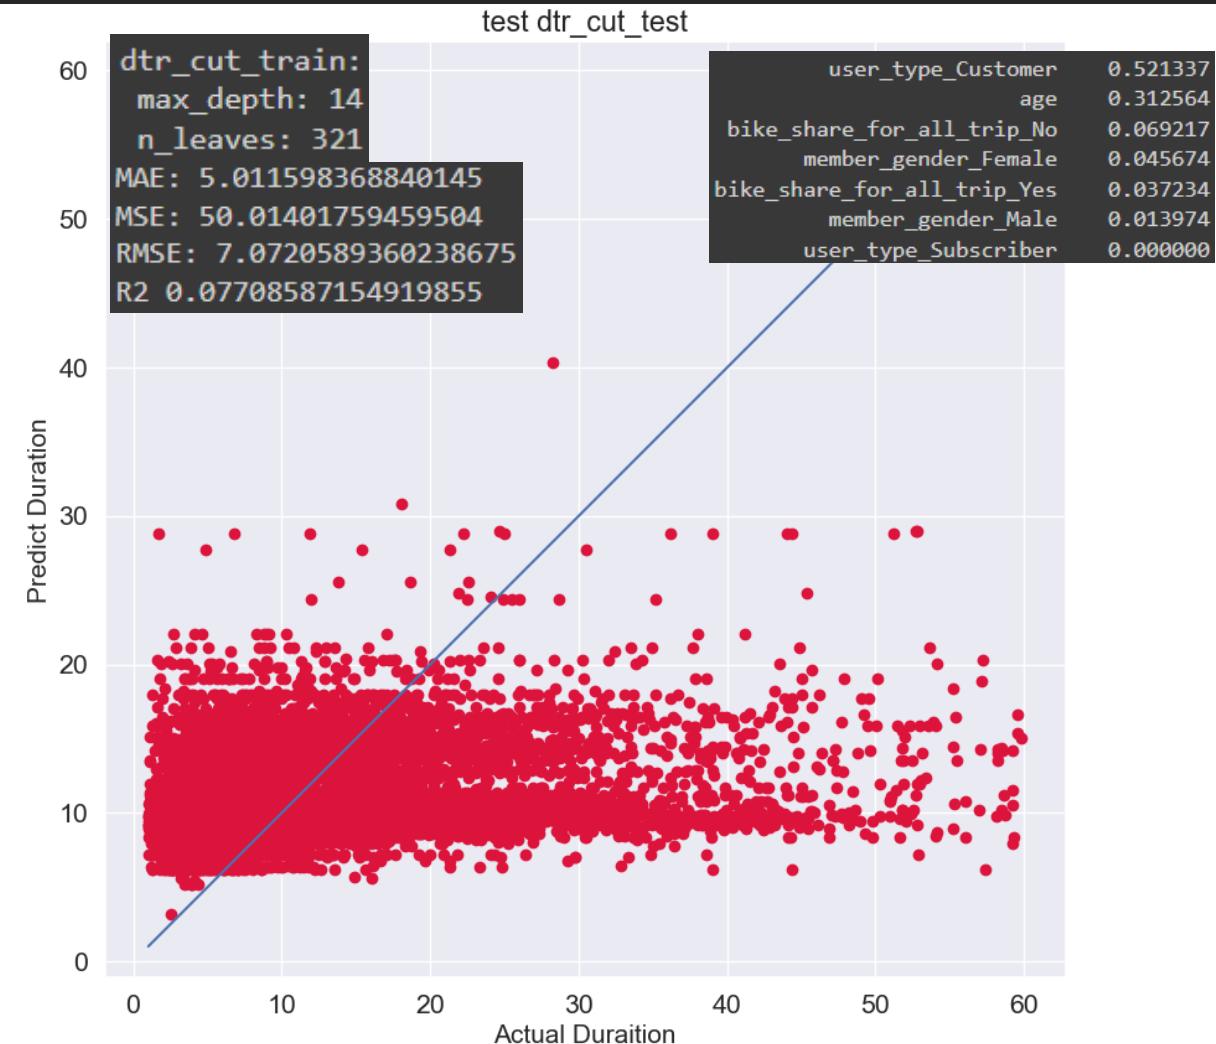
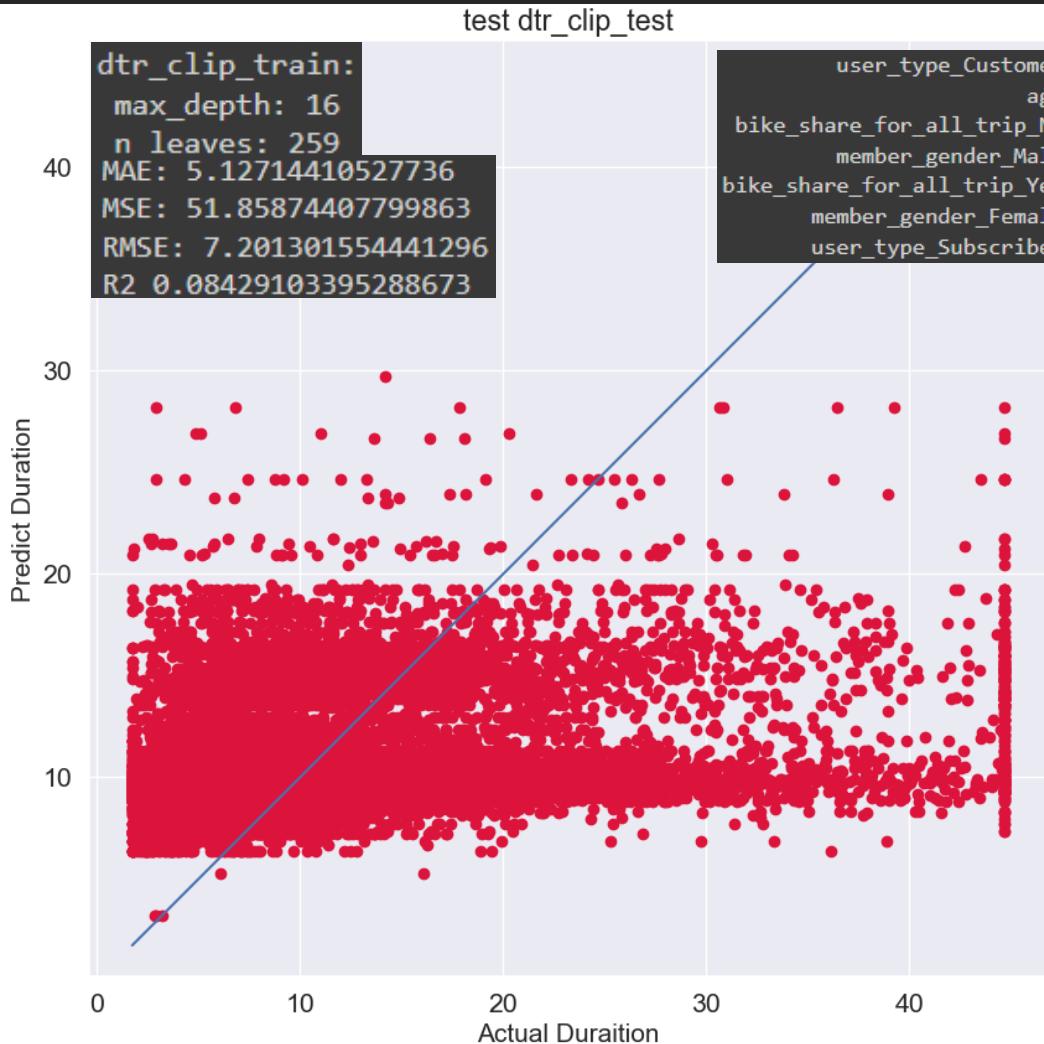
Ndtr_test

FEATURES:

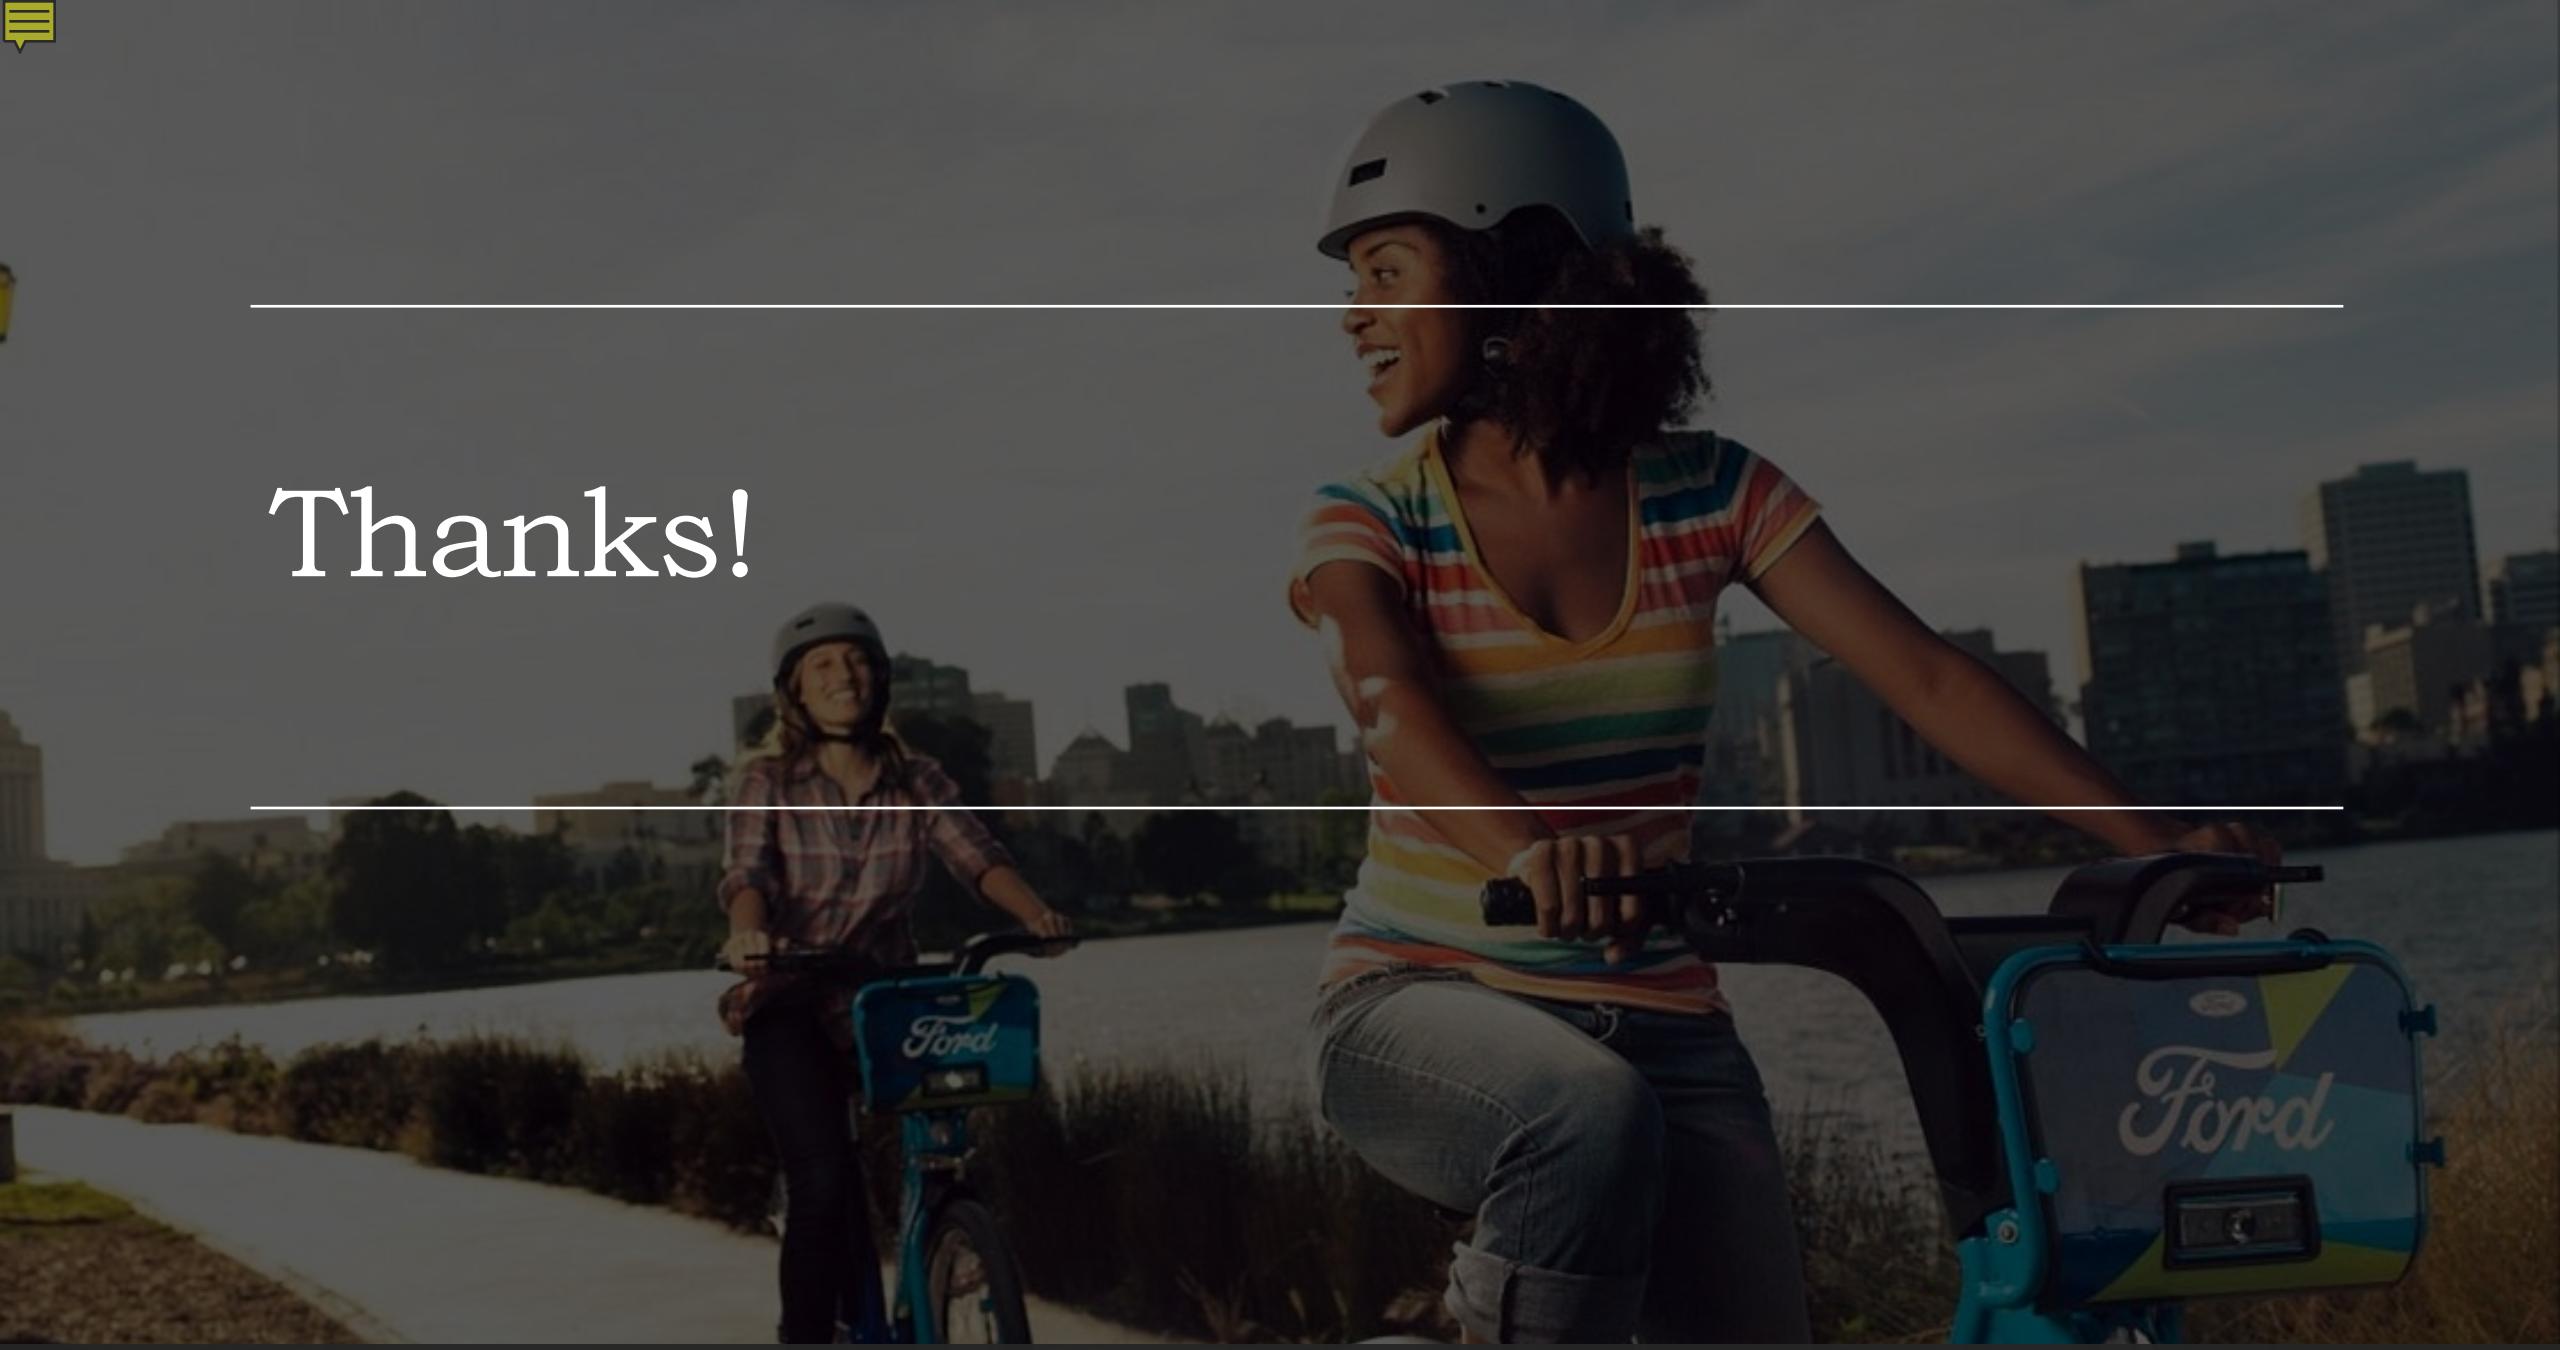
['age', 'user_type_Customer', 'user_type_Subscriber',
 'member_gender_Female', 'member_gender_Male']

HP:

max_dep = None
min_samp_split = 2
min_samp_leaf = 1
ccp = 0.0

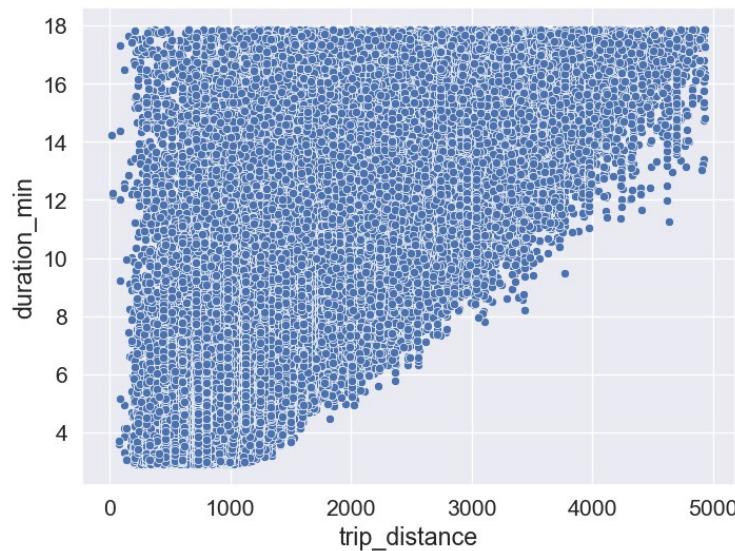
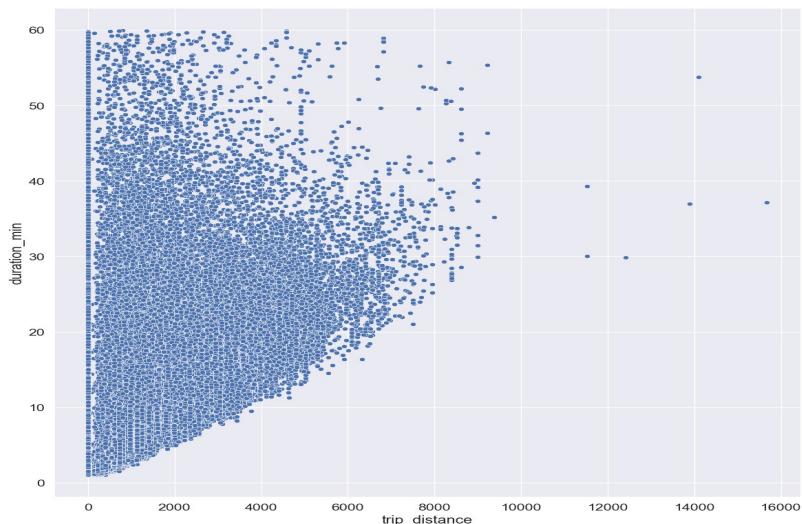


Thanks!



Statistical tuning data

```
def func_statistics_cleaning_numerical_values(  
    data, x_col, y_col,  
    x_k=3, max_x_bound=1, min_x_bound=1,  
    y_k=3, max_y_bound=1, min_y_bound=1  
):
```



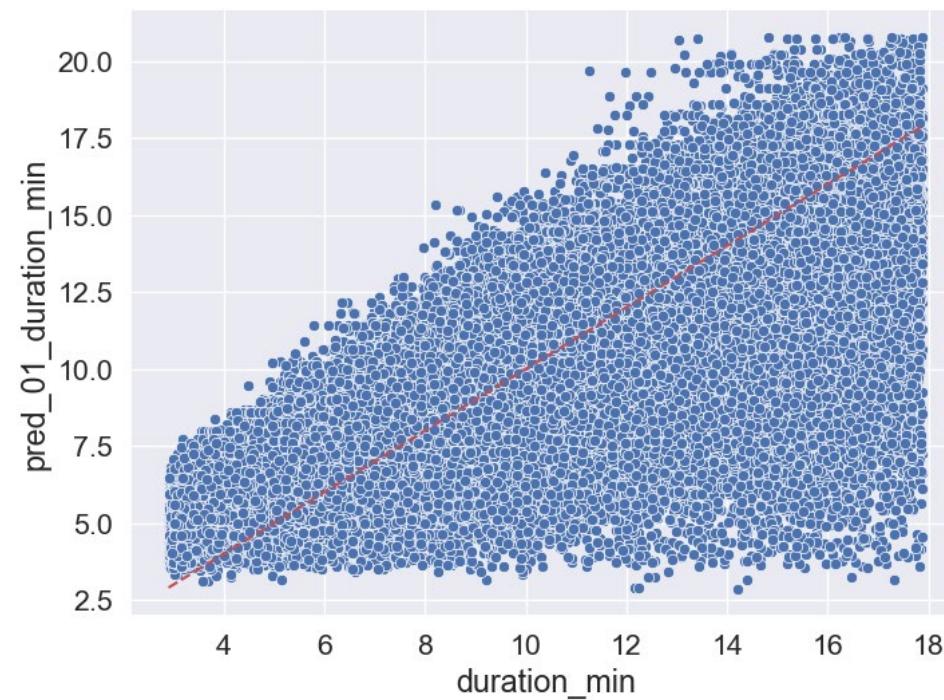
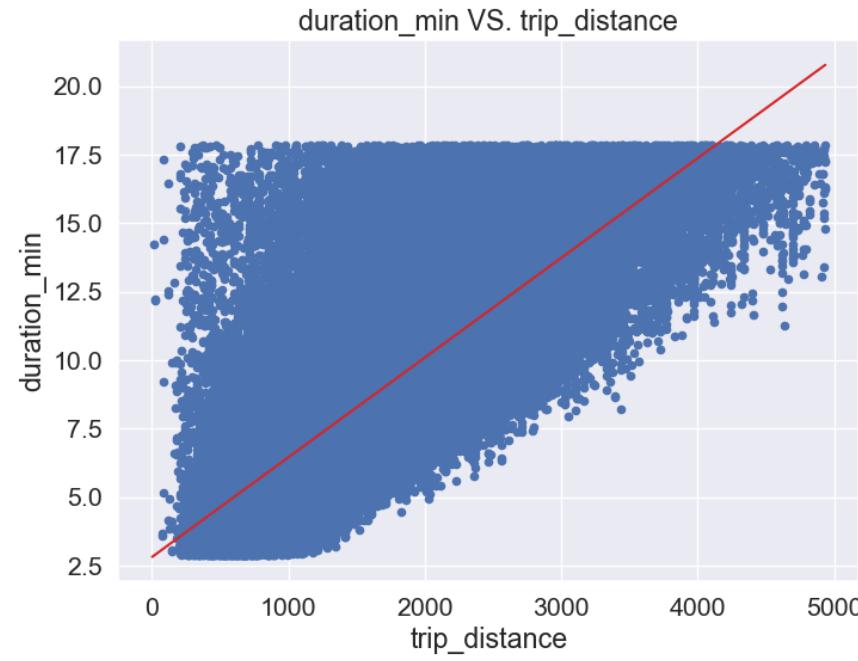
Raw data: (181434, 14)
Splitting data: (148053, 14)
Train data shape is : (103637, 14)
Validation data shape is: (22208, 14)
Test data shape is : (22208, 14)
Devaluation of 18.39%

lr_01

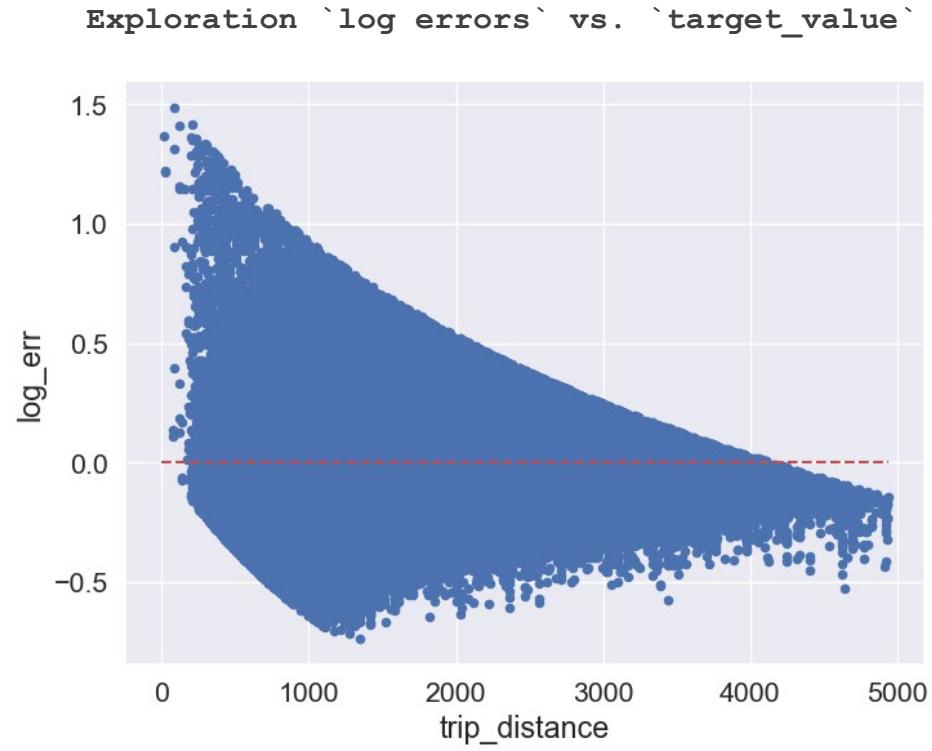
duration_min vs. trip_distance

TARGET = 'duration_min'
FEATURES = 'trip_distance'

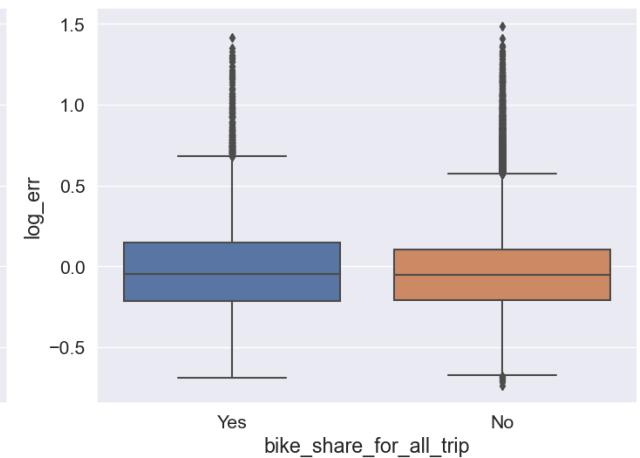
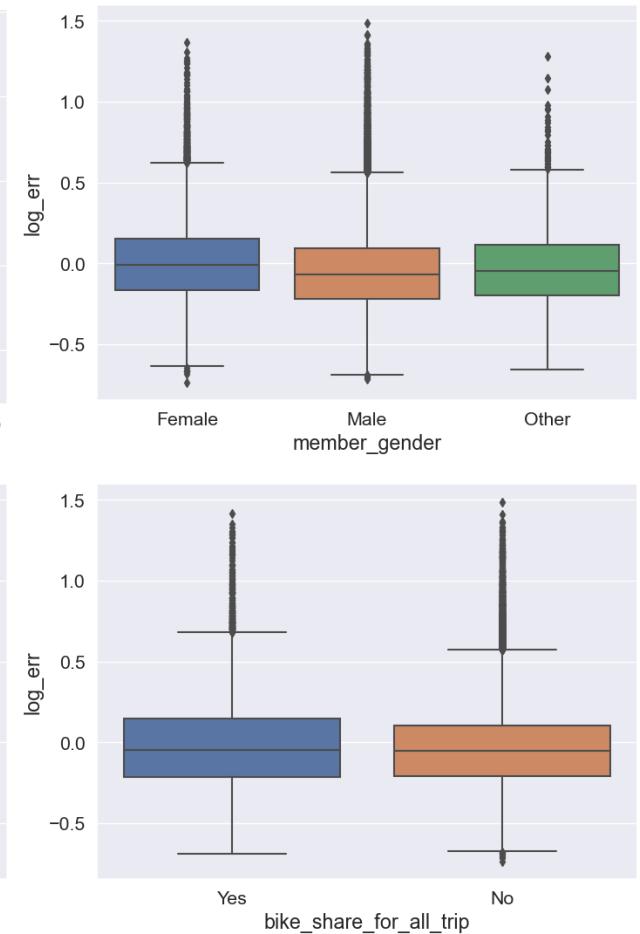
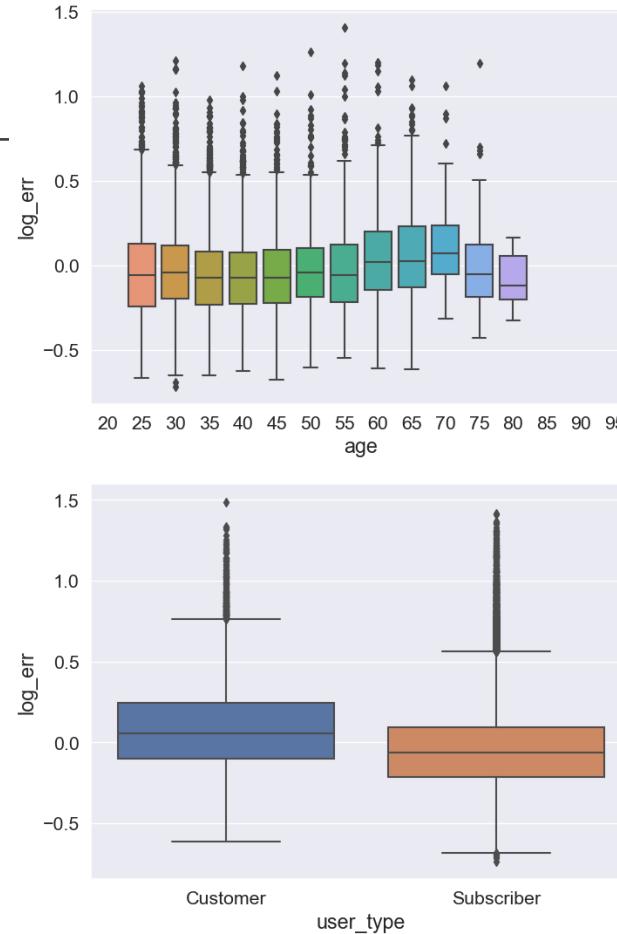
The empirical equation is
 $y = 0.0036x + 2.8227$



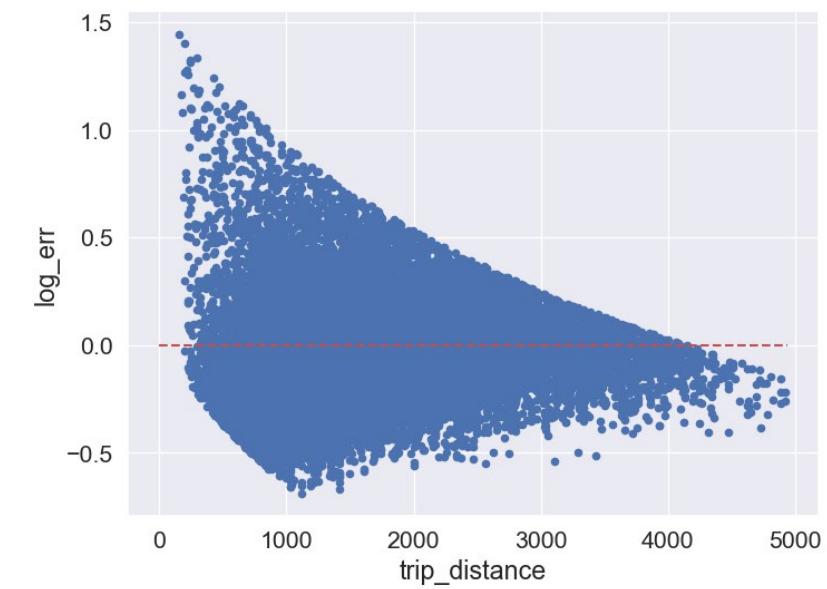
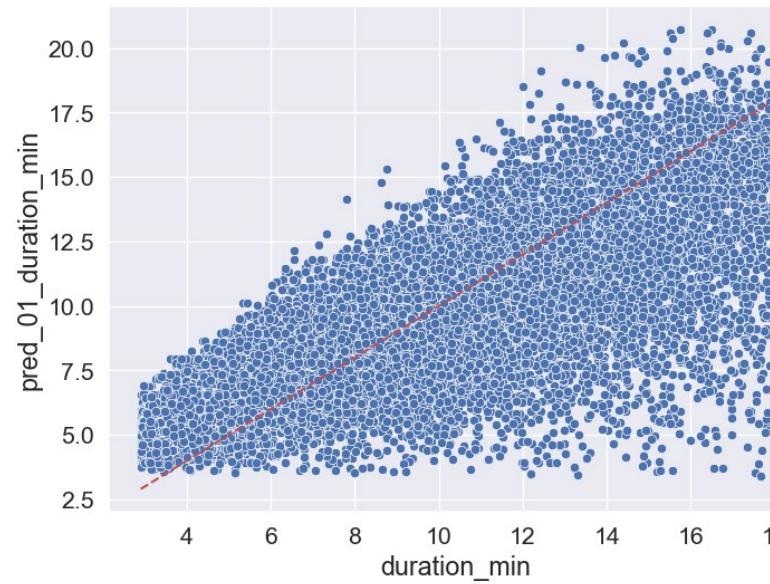
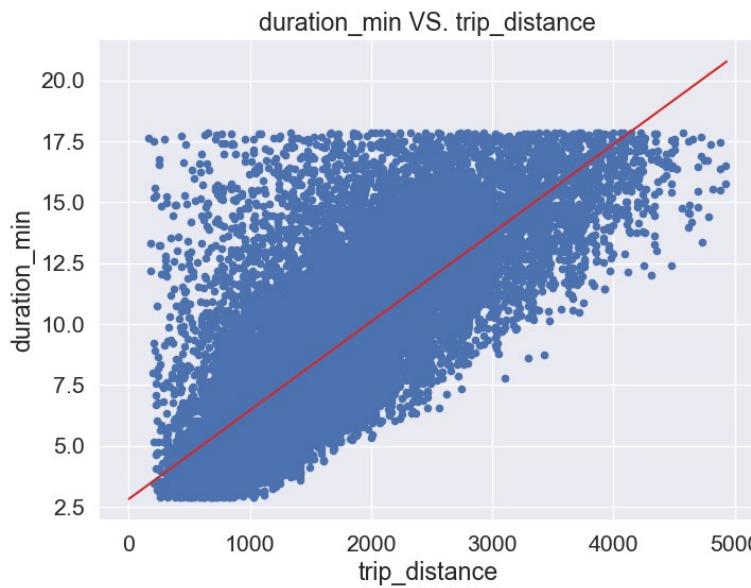
Residual



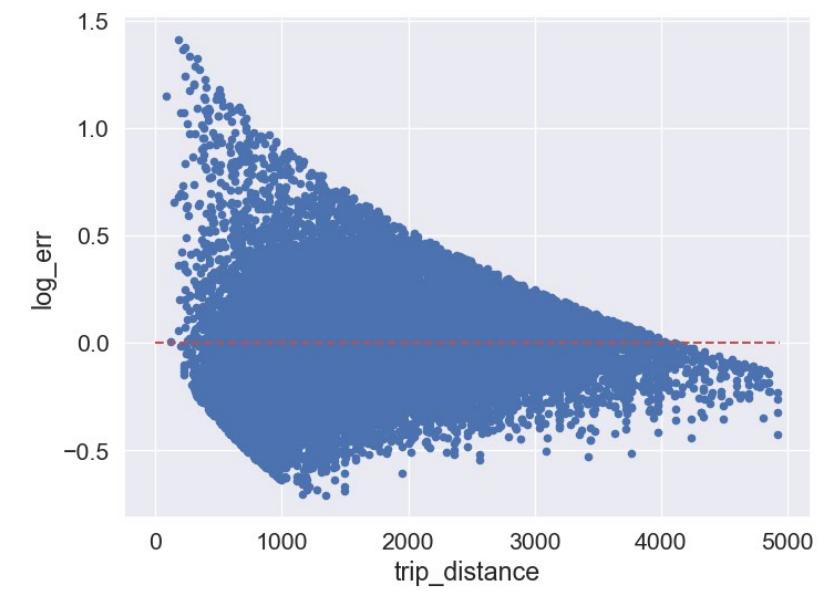
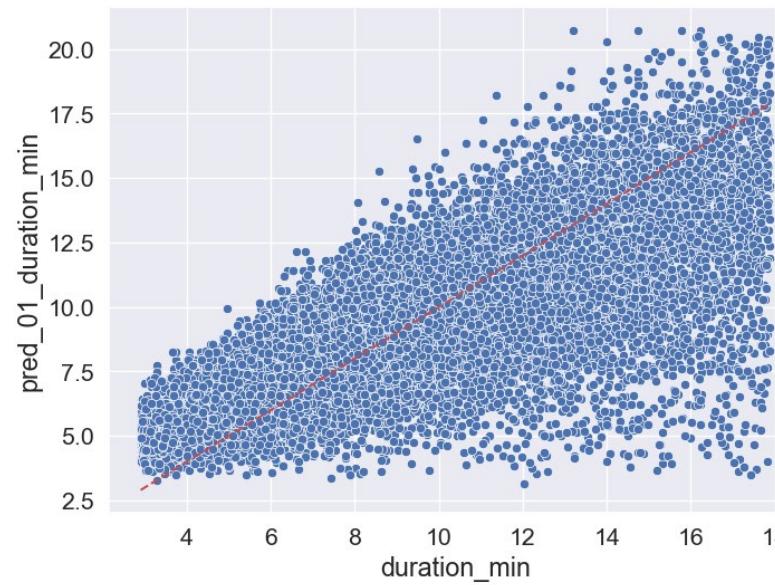
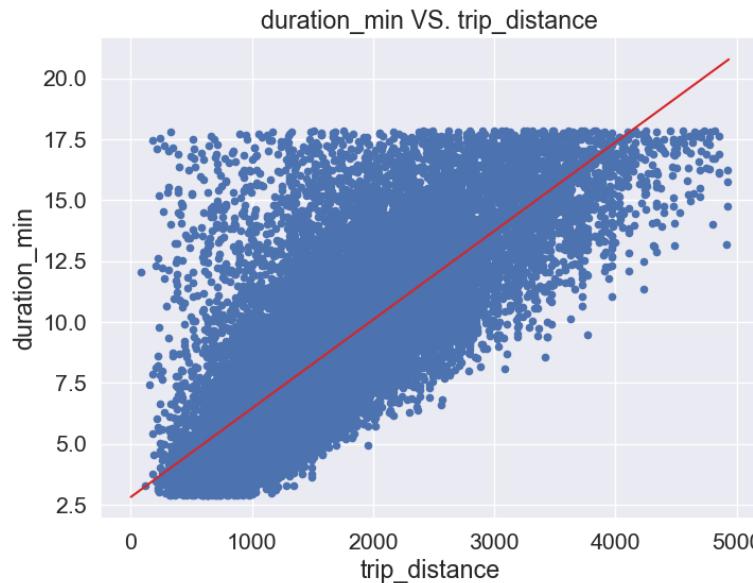
Exploration additional relationships for potential improvement of tuning



Validate the model



Testing the model



Evaluation

RMSE score

As lower *RMSE score* as better the prediction results;
For this project, the *RMSE score* value represent a minutes.

R² score

Coefficient of Determination.

The best possible *R² score* is **1.0** and it can be negative (because the model can be arbitrarily worse).

Generally, when the **true Y** is non-constant, a constant model that always predicts **the Average of Y** disregarding the input features would get a *R² score* of 0.

lr_01	fdf_train	fdf_valid	fdf_test
<i>RMSE score</i>	2.3723961060413967	2.365876472027405	2.37110246288299
<i>R² score</i>	0.6110326559694819	0.6137572392225006	0.6136474098447284

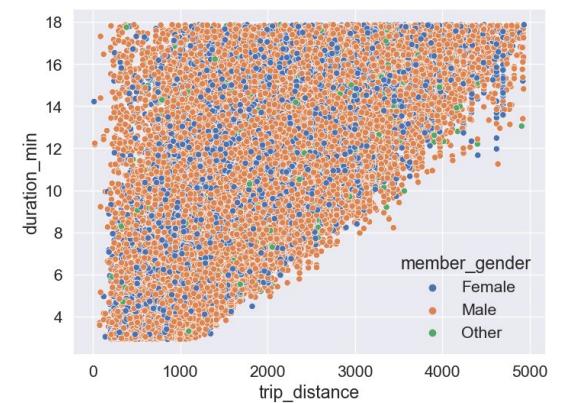
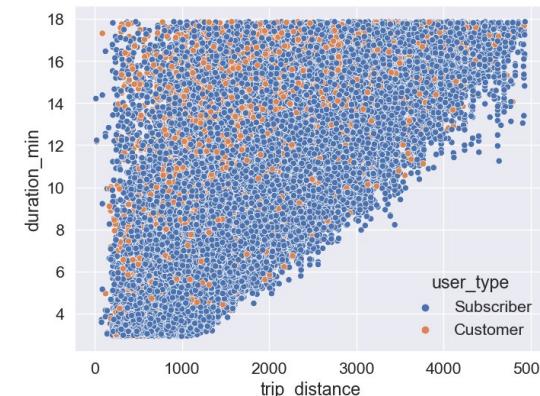
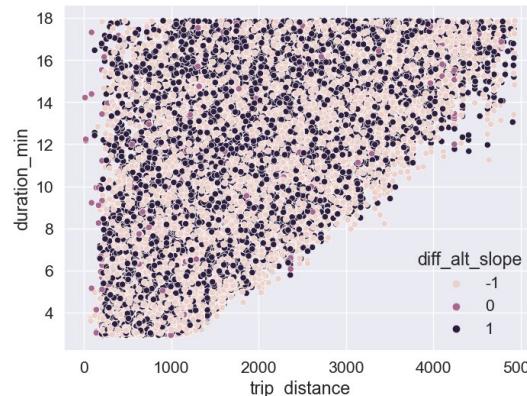
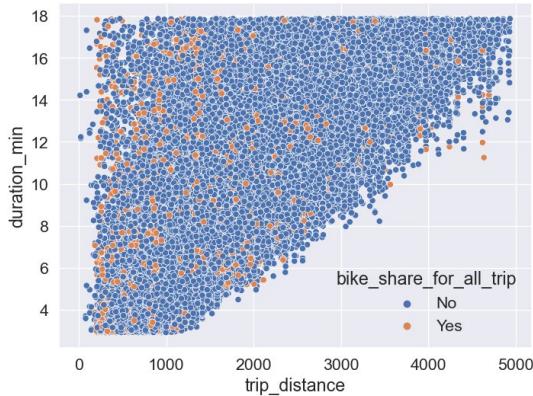
lr_02

duration_min vs. trip_distance, member_gender, user_type, bike_share_for_all_trip, diff_alt_slope, 'avg_trip_speed'

```
Enumerate categorical columns: ['gender_numeric', 'numeric_type', 'numeric_all_trip']
```

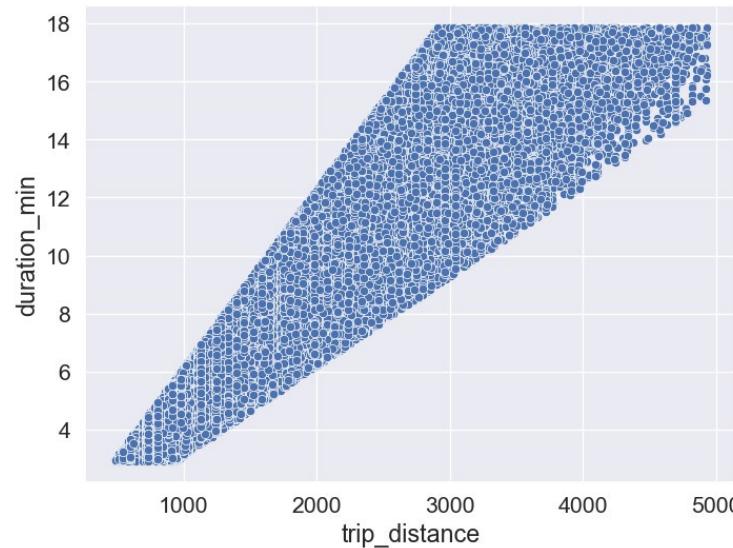
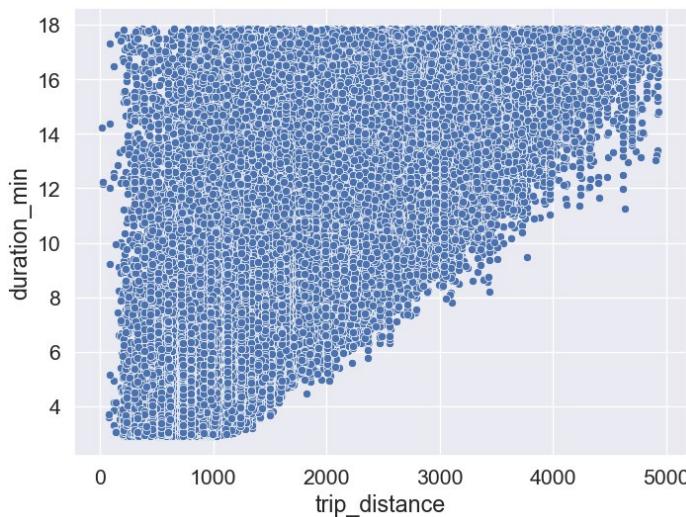
```
TARGET = 'duration_min'
```

```
FEATURES = ['trip_distance', 'gender_numeric', 'numeric_type', 'numeric_all_trip', 'diff_alt_slope', 'avg_trip_speed']
```



Circumstantial cleaning

By cleaning the observations below the average speed minus 1 standard deviation threshold we clean the users which probably take a tour and using Ford-Go-Bike service as a “tour vehicle” and not as a transportation tool



Raw data: (148053, 17)
Splitting data: (101277, 17)
Train data shape is : (70893, 17)
Validation data shape is: (15192, 17)
Test data shape is : (15192, 17)

Another devaluation of 31.59%
Total devaluation of 44.17%

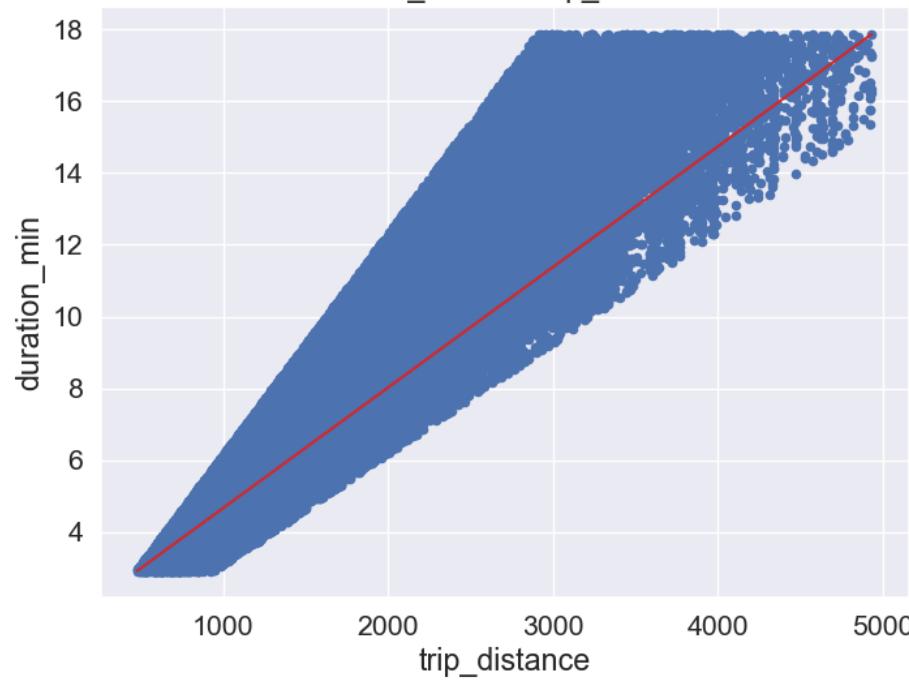
Predict

duration_min vs. trip_distance

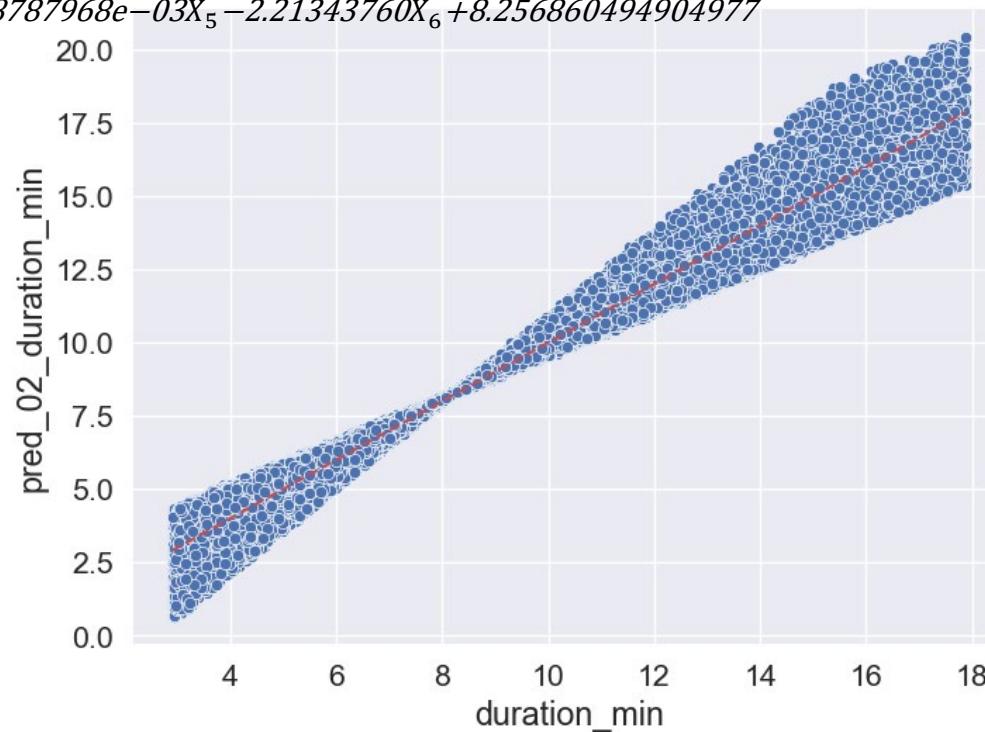
The empirical equation is

$$y = 4.54056103e-03X_1 +$$

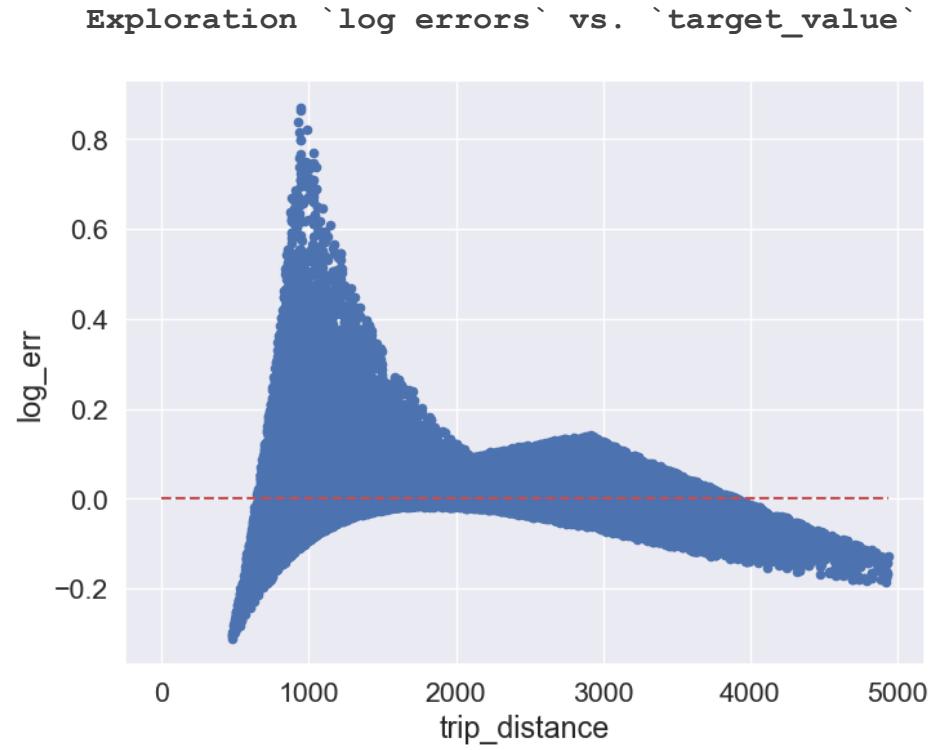
duration_min VS. trip_distance



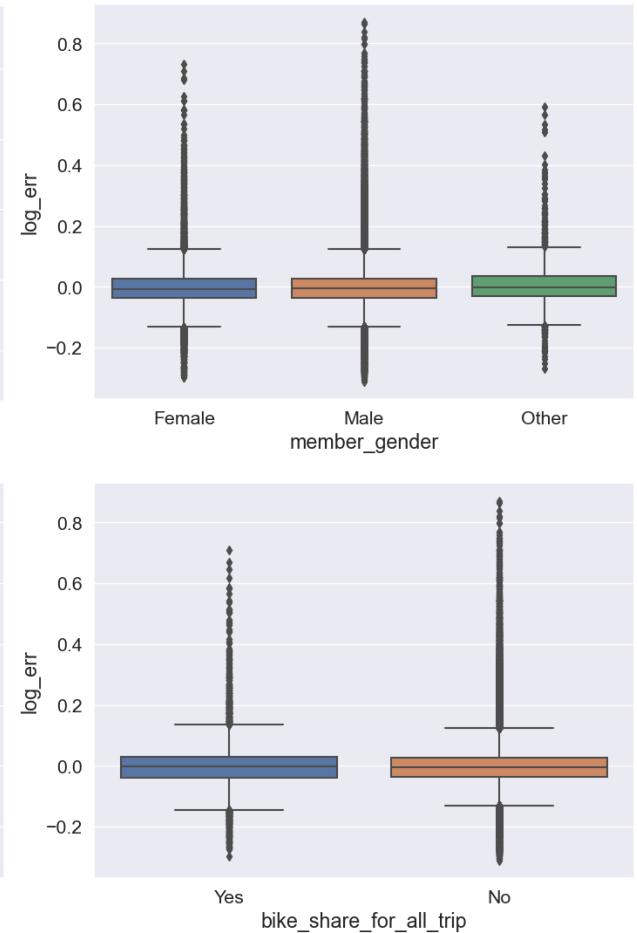
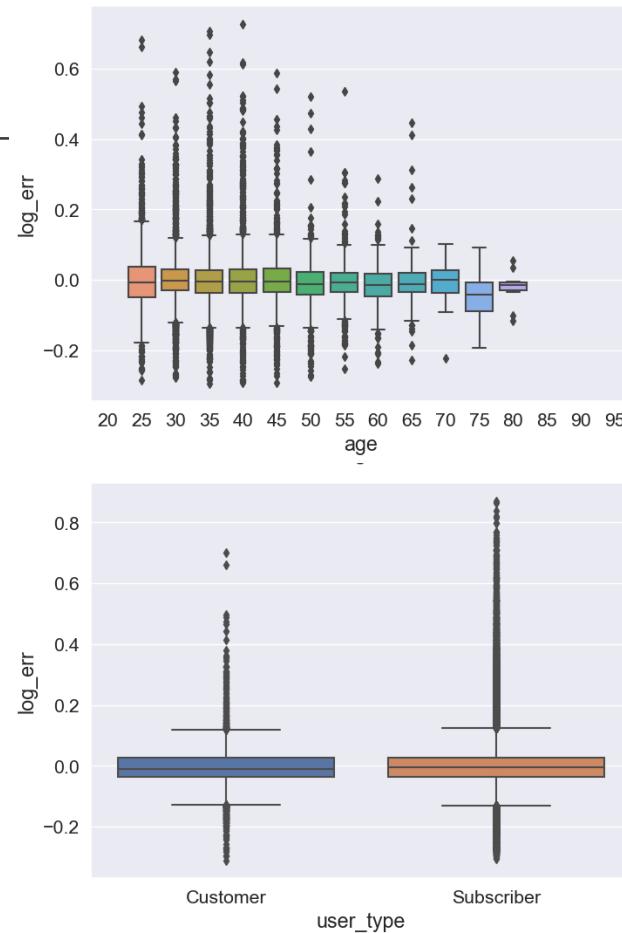
$$-1.38787968e-03X_5 - 2.21343760X_6 + 8.256860494904977$$



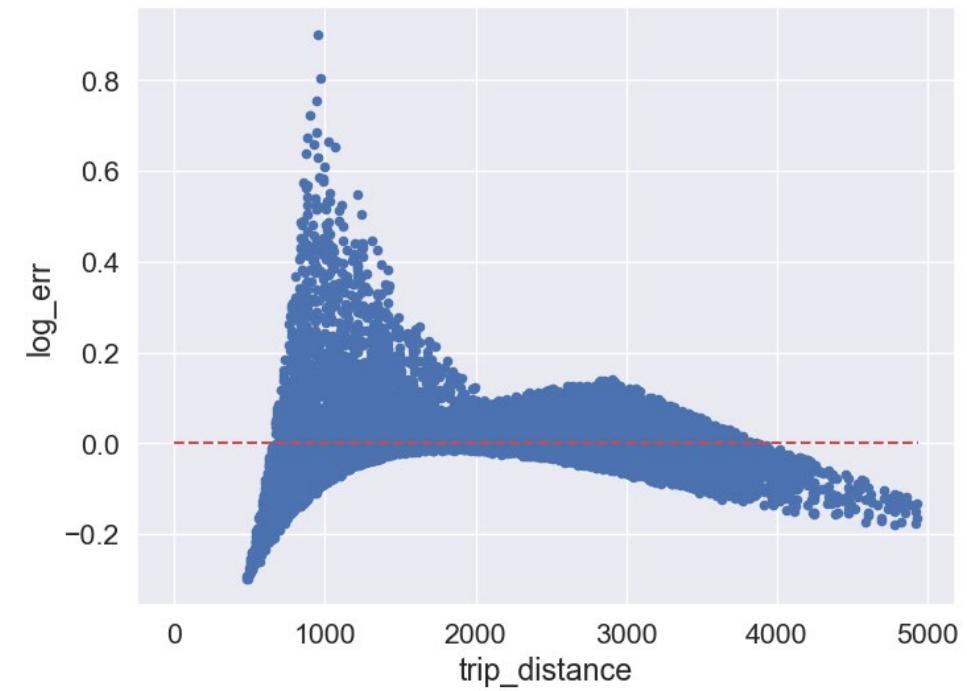
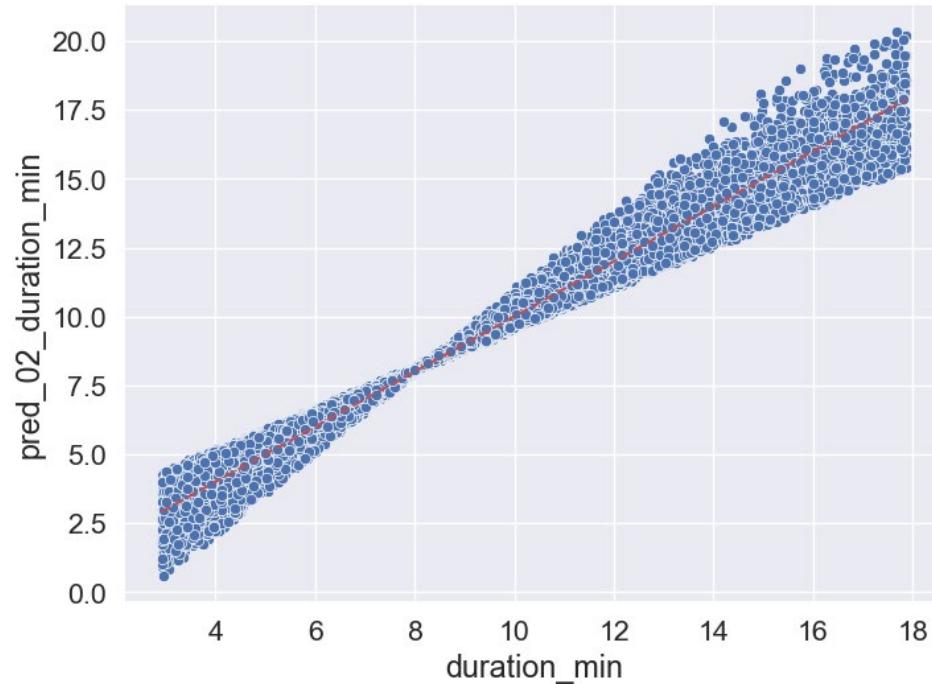
Residual



Exploration additional relationships for potential improvement of tuning



Validate the model



Testing the model

